

Neural Networks with TensorFlow

Overview

Neural networks are representation based machine learning algorithms

Neural networks are made up of building blocks called neurons


Each neuron is made up of a linear function and an activation function

Performance is very sensitive to details such as proper choice of activation

Overfitting in neural networks is mitigated using techniques such as dropout


Understanding Machine Learning

Whales: Fish or Mammals?



Mammals

Members of the infraorder Cetacea



Fish

Look like fish, swim like fish, move with fish

Whales: Fish or Mammals?



ML-based Classifier

Training

Feed in a large corpus of data classified correctly

Prediction

Use it to classify new instances which it has not seen before

Training the ML-based Classifier



Corpus



ML-based Classifier



Classification



Feedback - loss
function or cost
function

Improves model parameters

ML-based Binary Classifier

Breathes like a mammal
Gives birth like a mammal



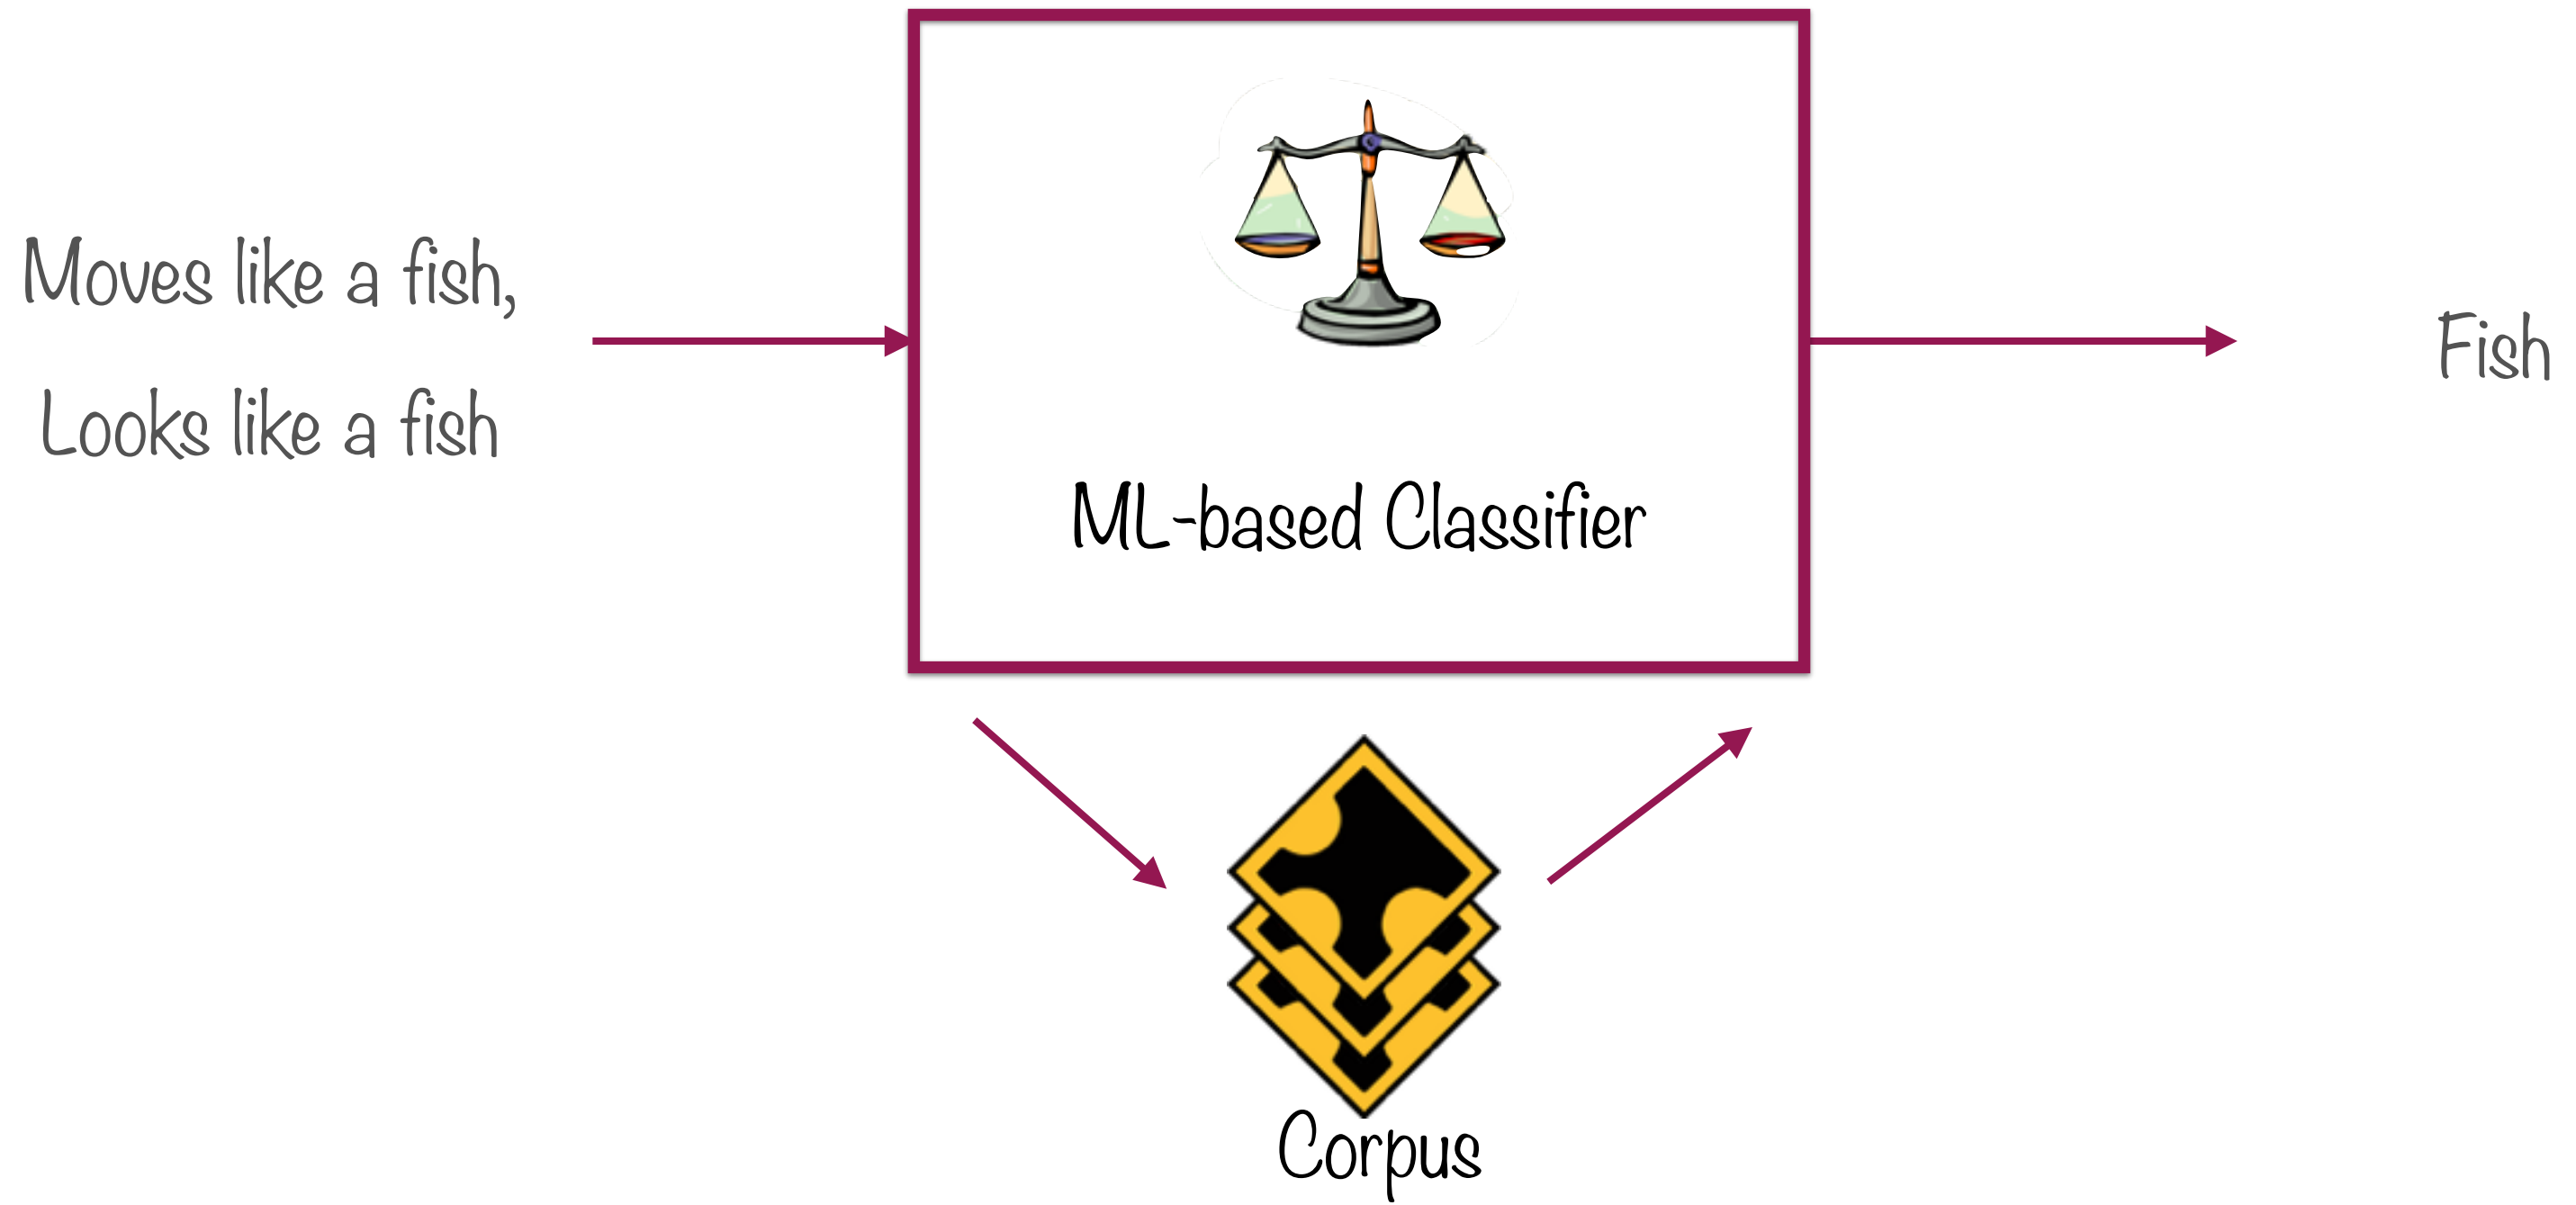
ML-based Classifier

Mammal




Corpus


“Traditional” ML-based Binary Classifier



ML-based Binary Classifier



Corpus




Classification Algorithm



ML-based Classifier

ML-based Binary Classifier



Corpus



Naive Bayes, Support
Vector Machines,
Decision Trees



ML-based Classifier

ML-based Binary Classifier

Breathes like a mammal
Gives birth like a mammal

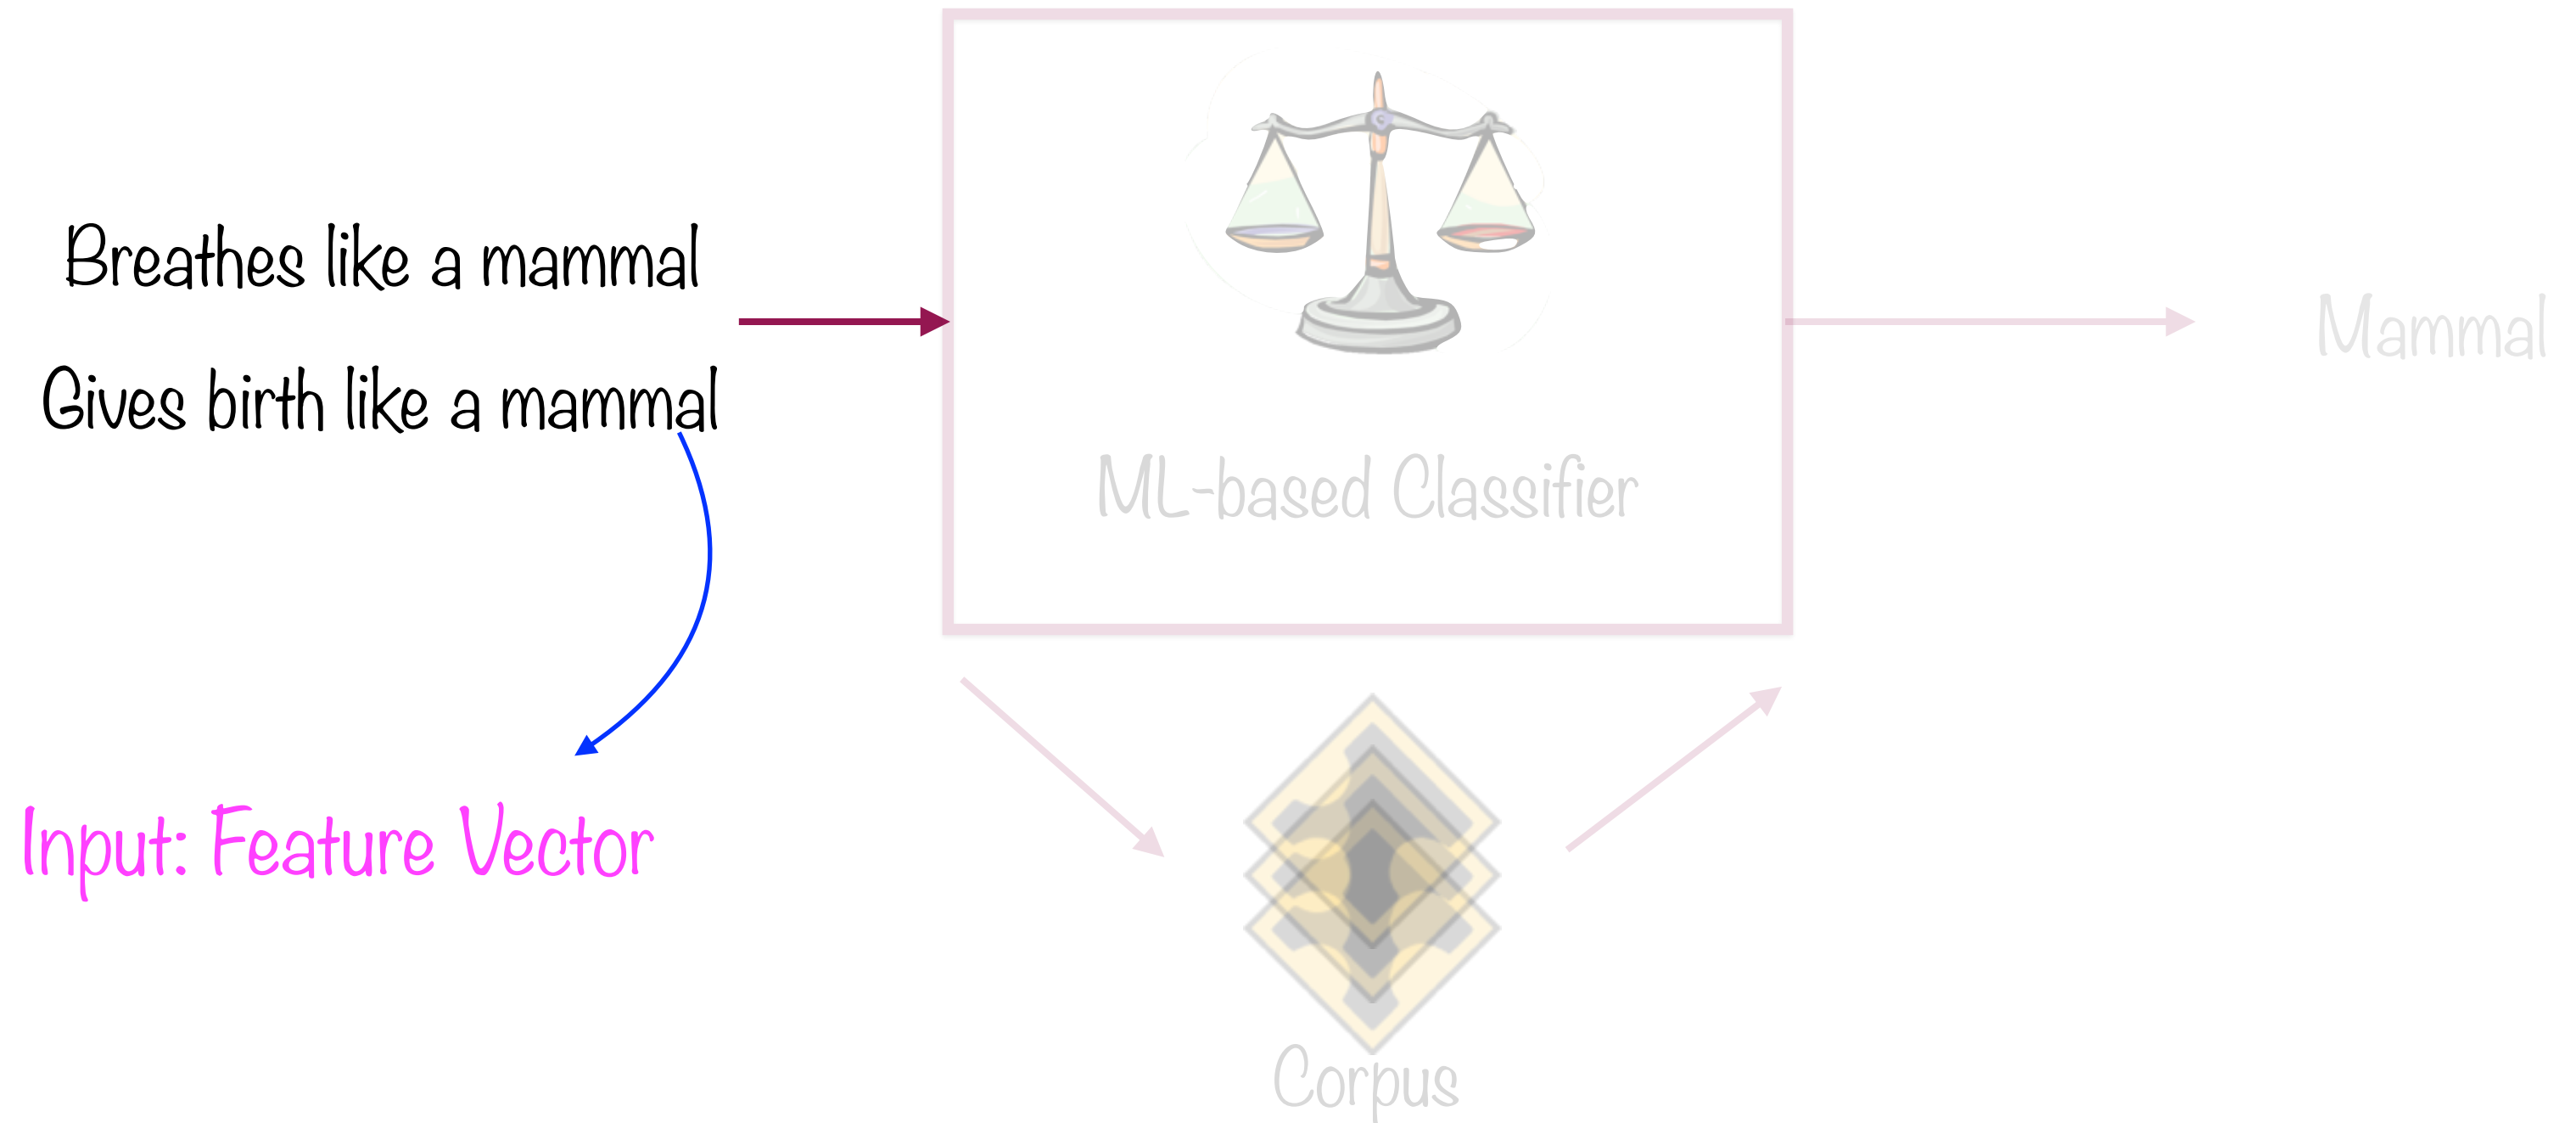


Mammal

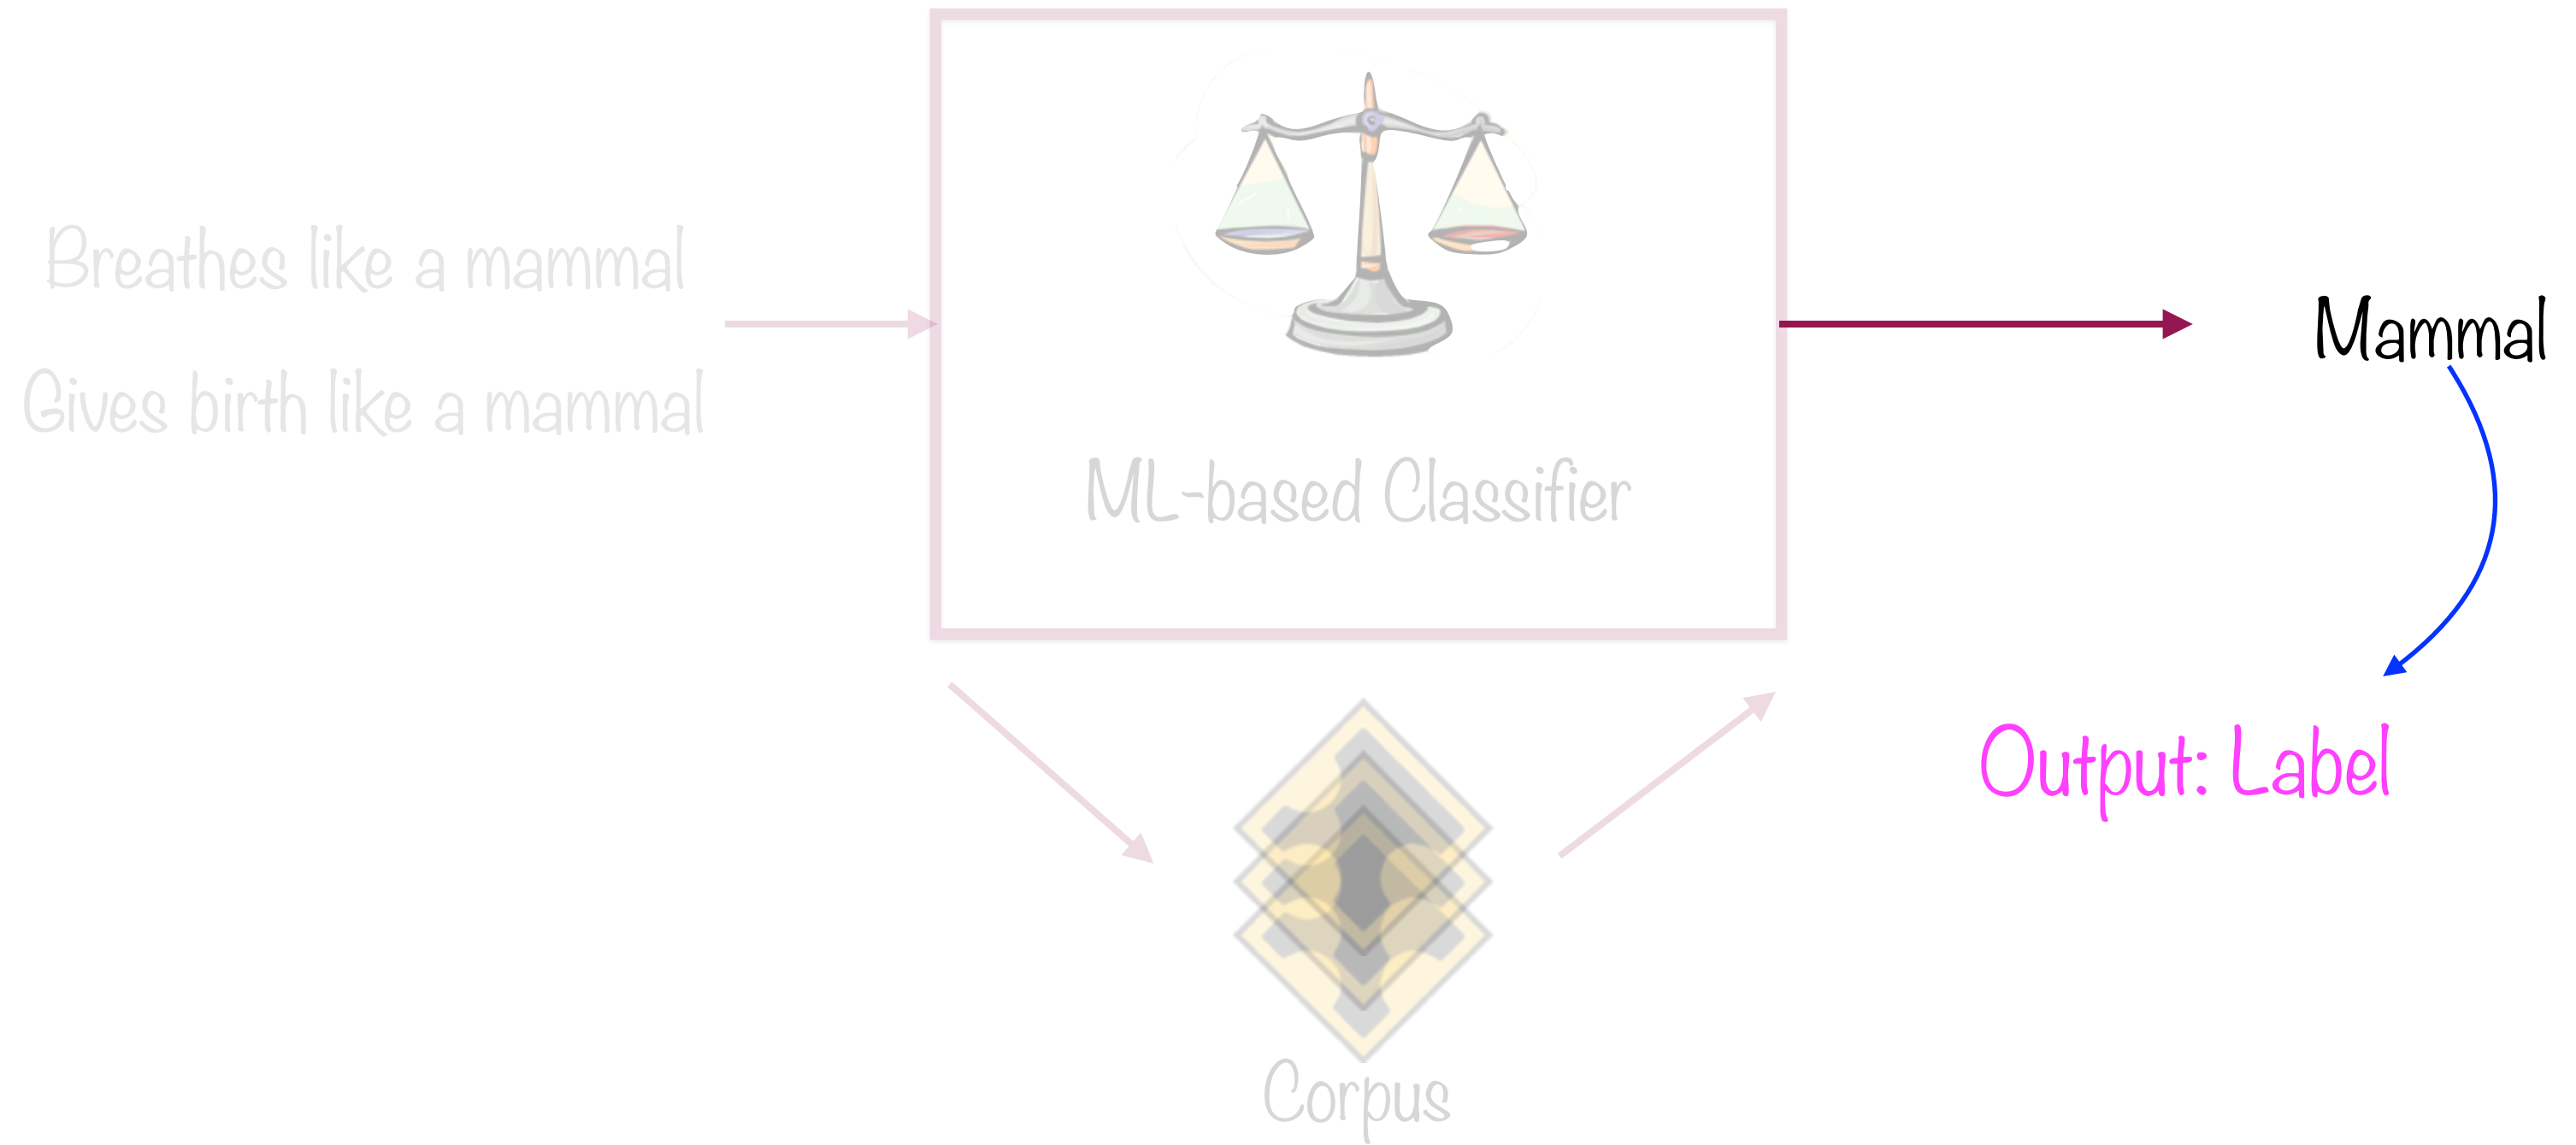


Corpus

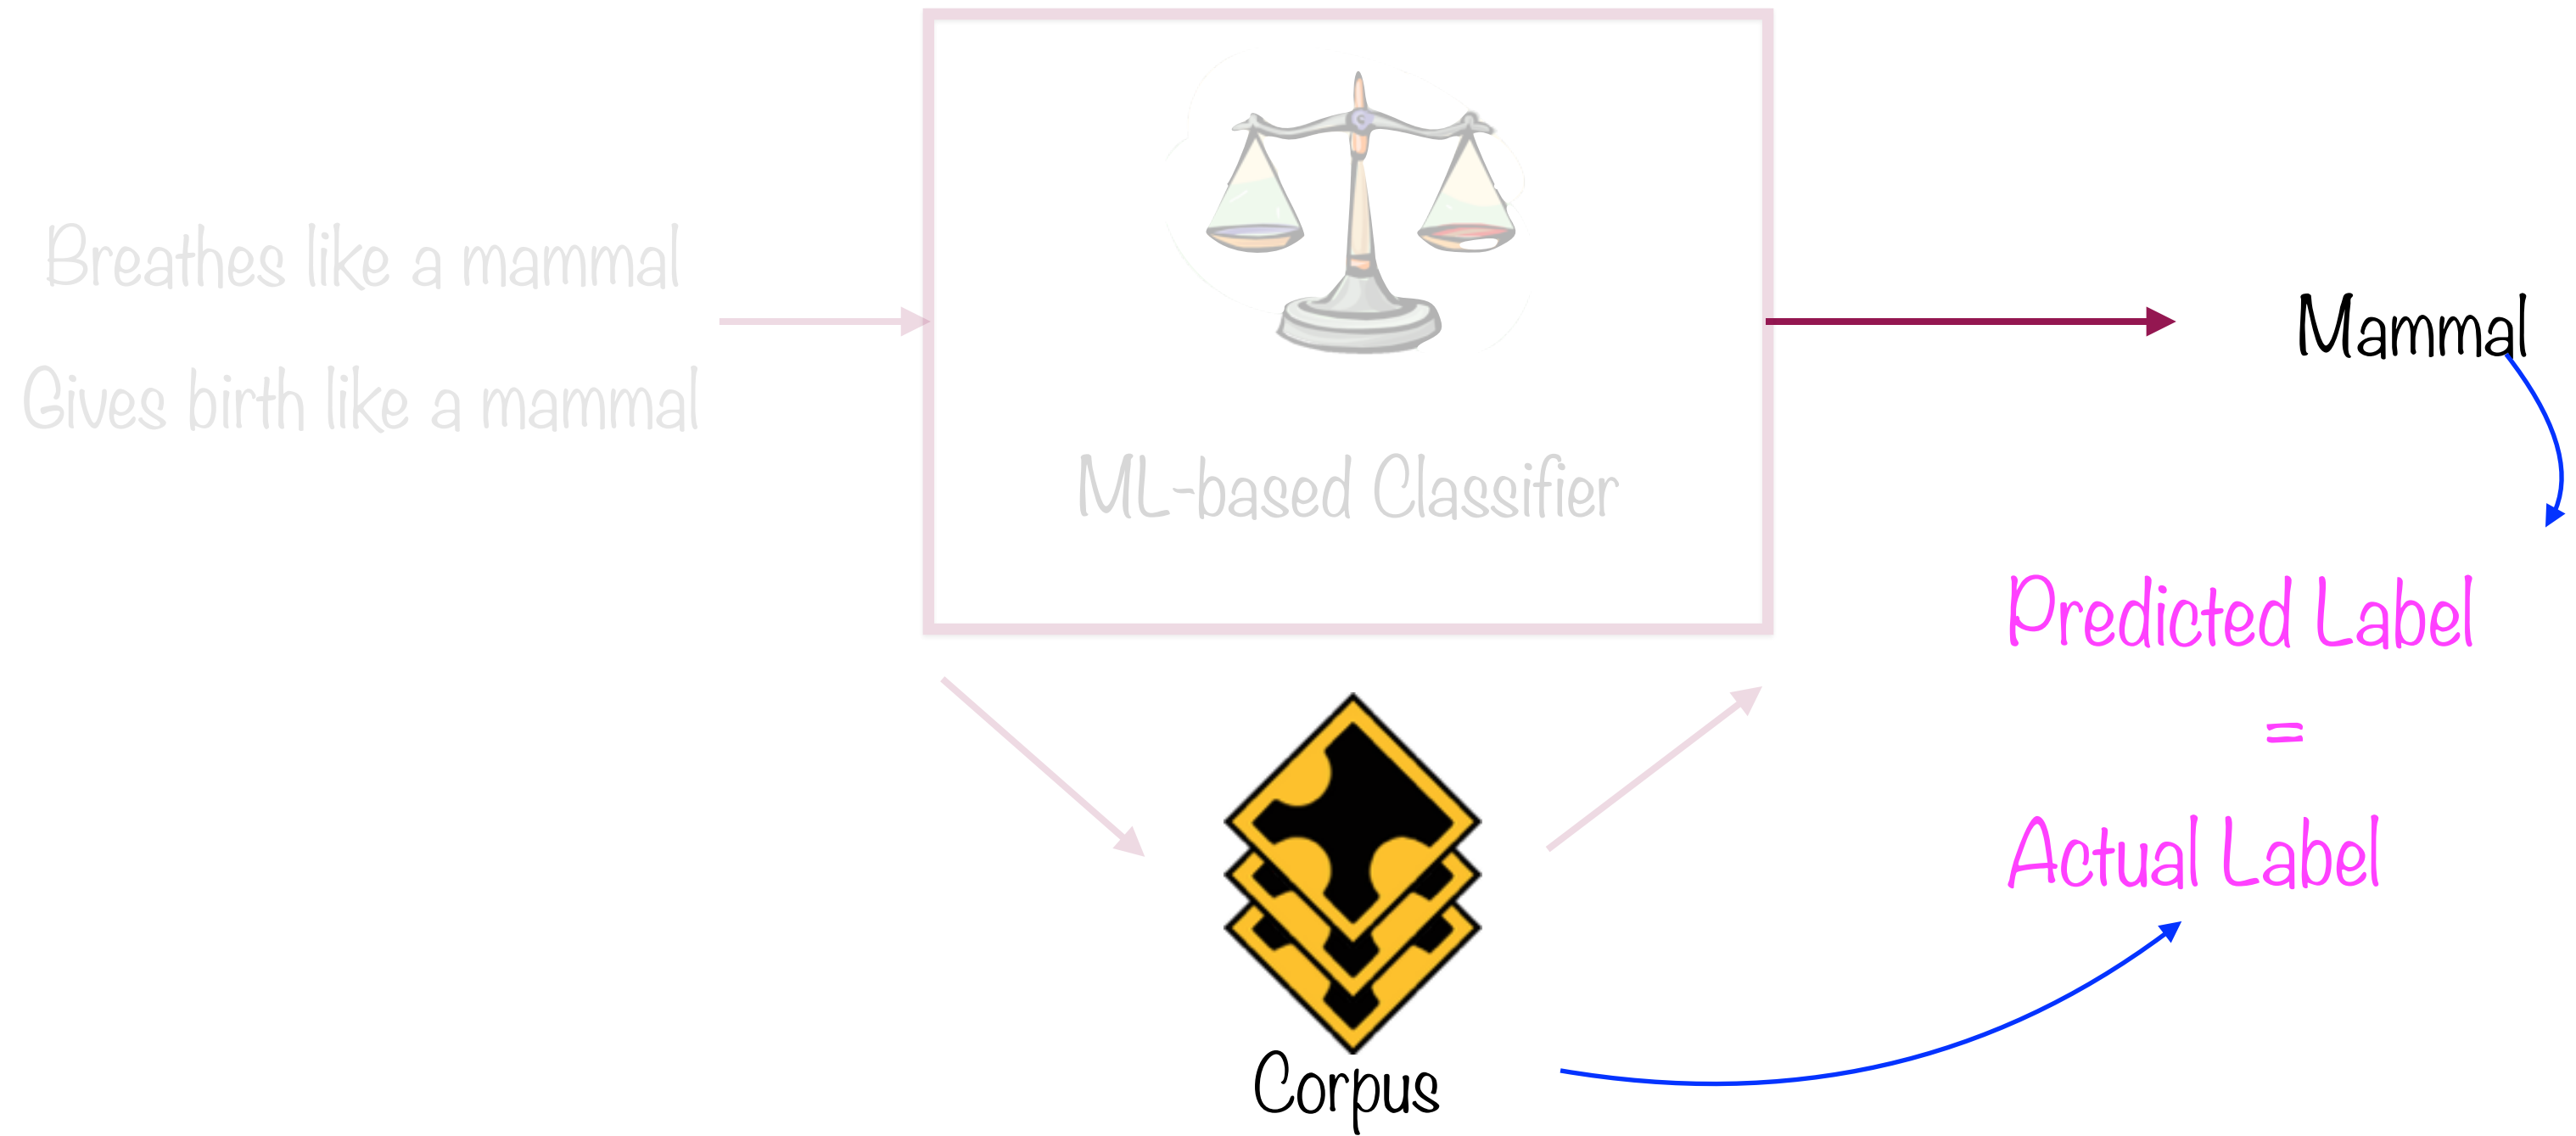
ML-based Binary Classifier



ML-based Binary Classifier



ML-based Binary Classifier




ML-based Binary Classifier

Moves like a fish,
Looks like a fish



Fish



ML-based Binary Classifier

Moves like a fish,
Looks like a fish

Input: Feature Vector



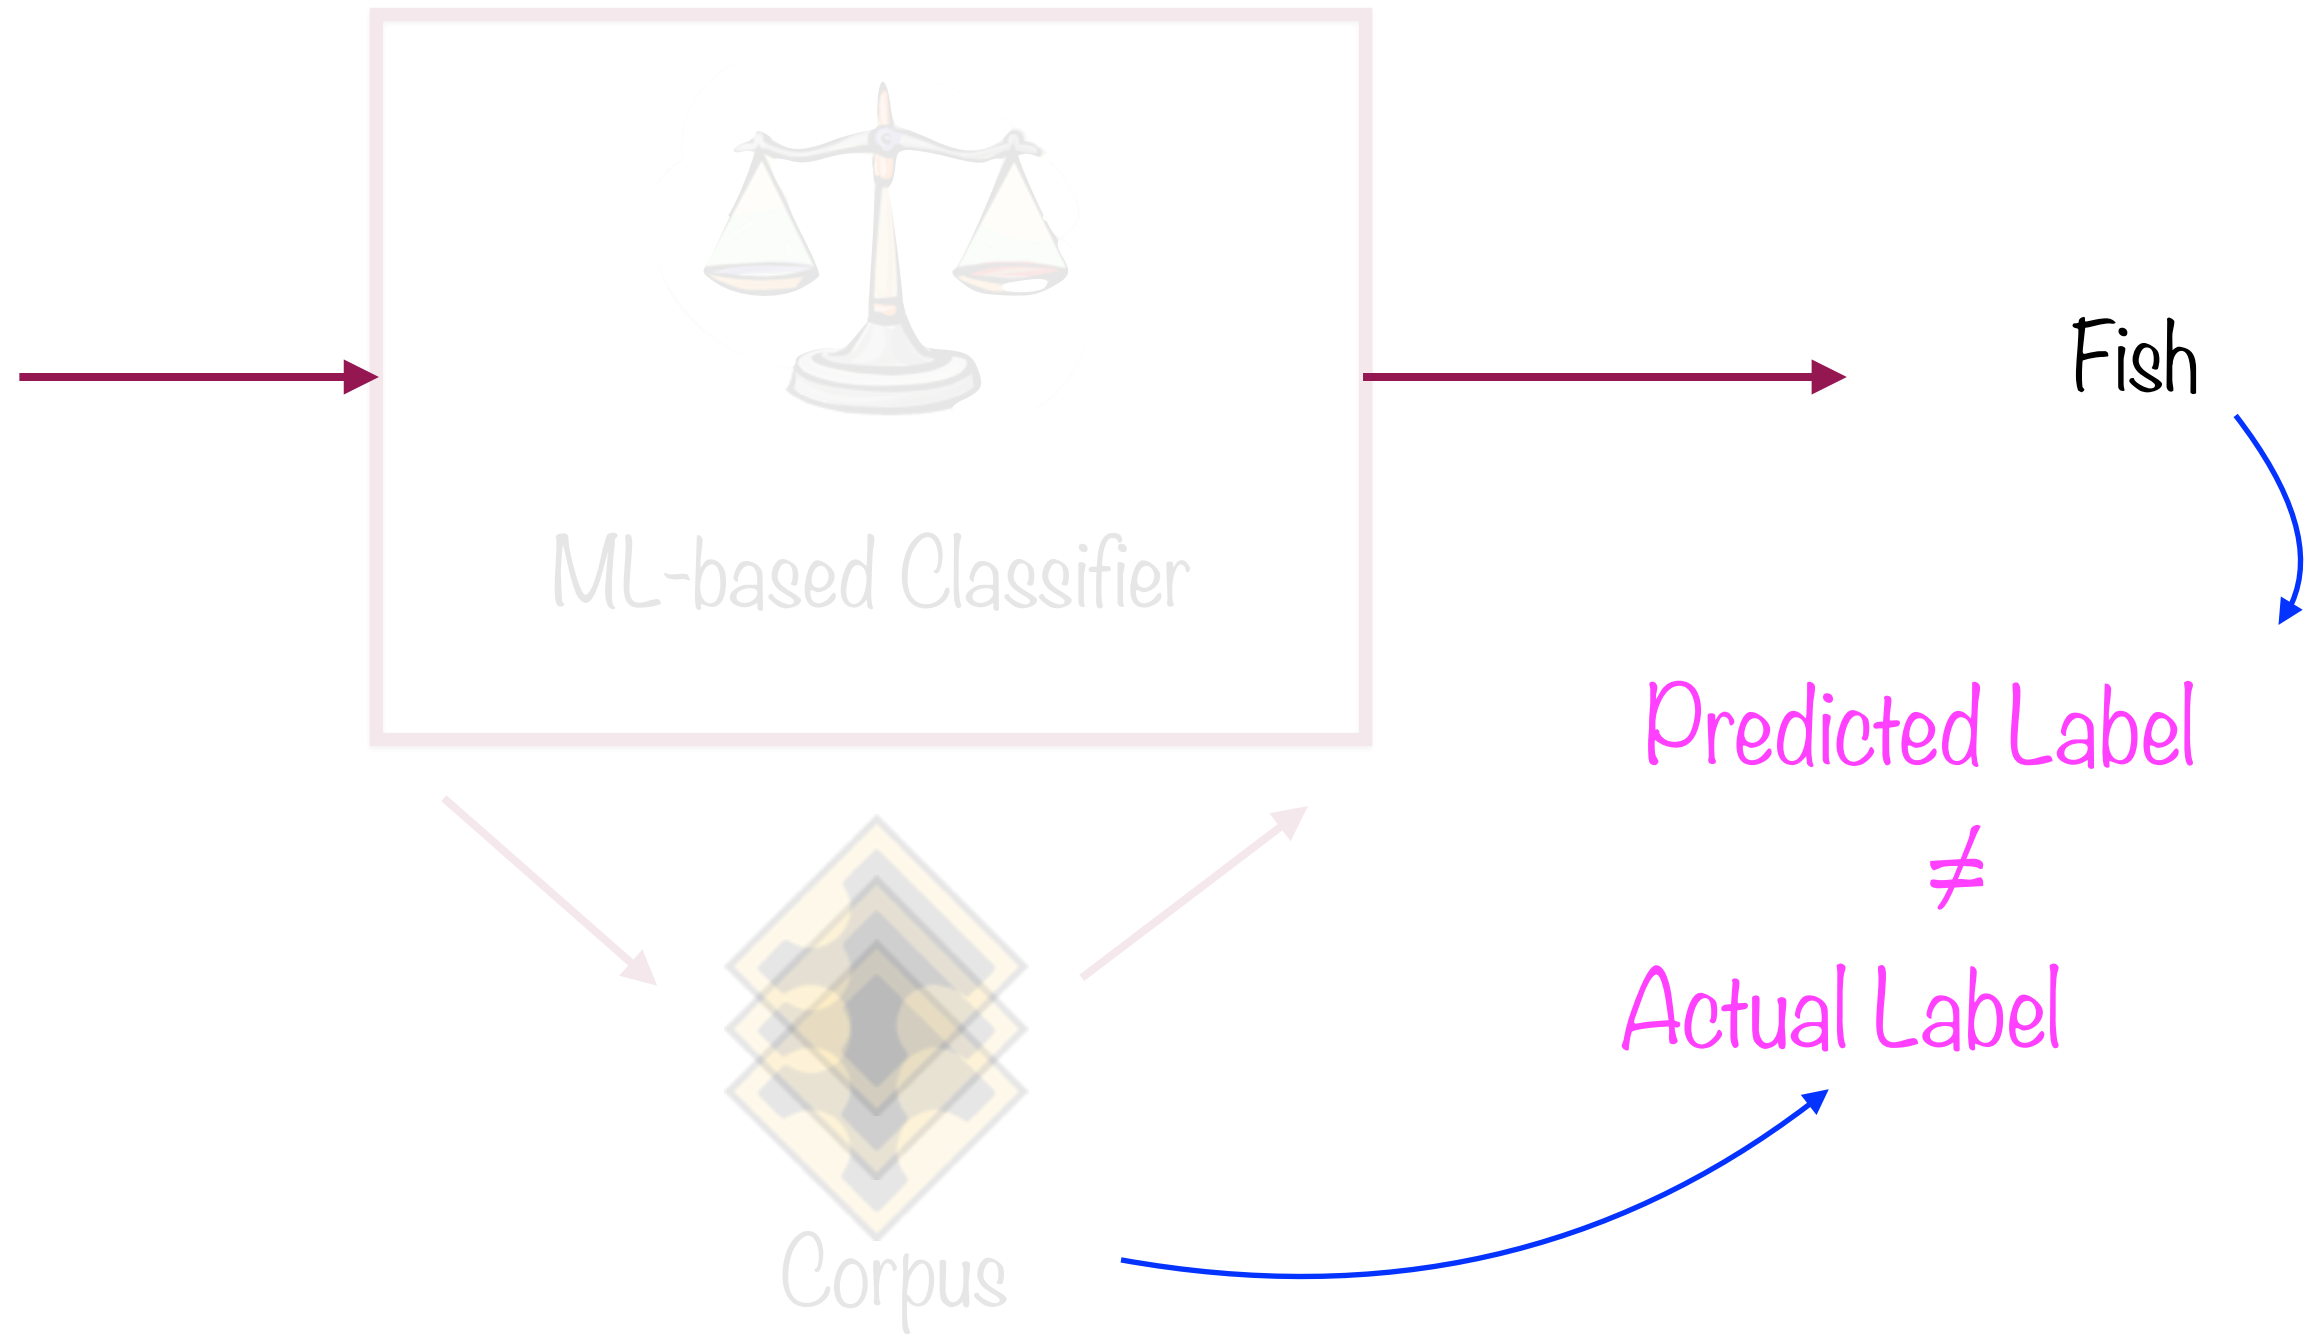
Fish



Corpus



ML-based Binary Classifier

Moves like a fish,
Looks like a fish



Understanding Deep Learning

“Traditional” ML-based Binary Classifier



Corpus




Classification Algorithm



ML-based Classifier

“Traditional” ML-based Binary Classifier



Corpus



Feature Selection by
Experts




Classification
Algorithm



ML-based Classifier

“Traditional” ML-based Binary Classifier



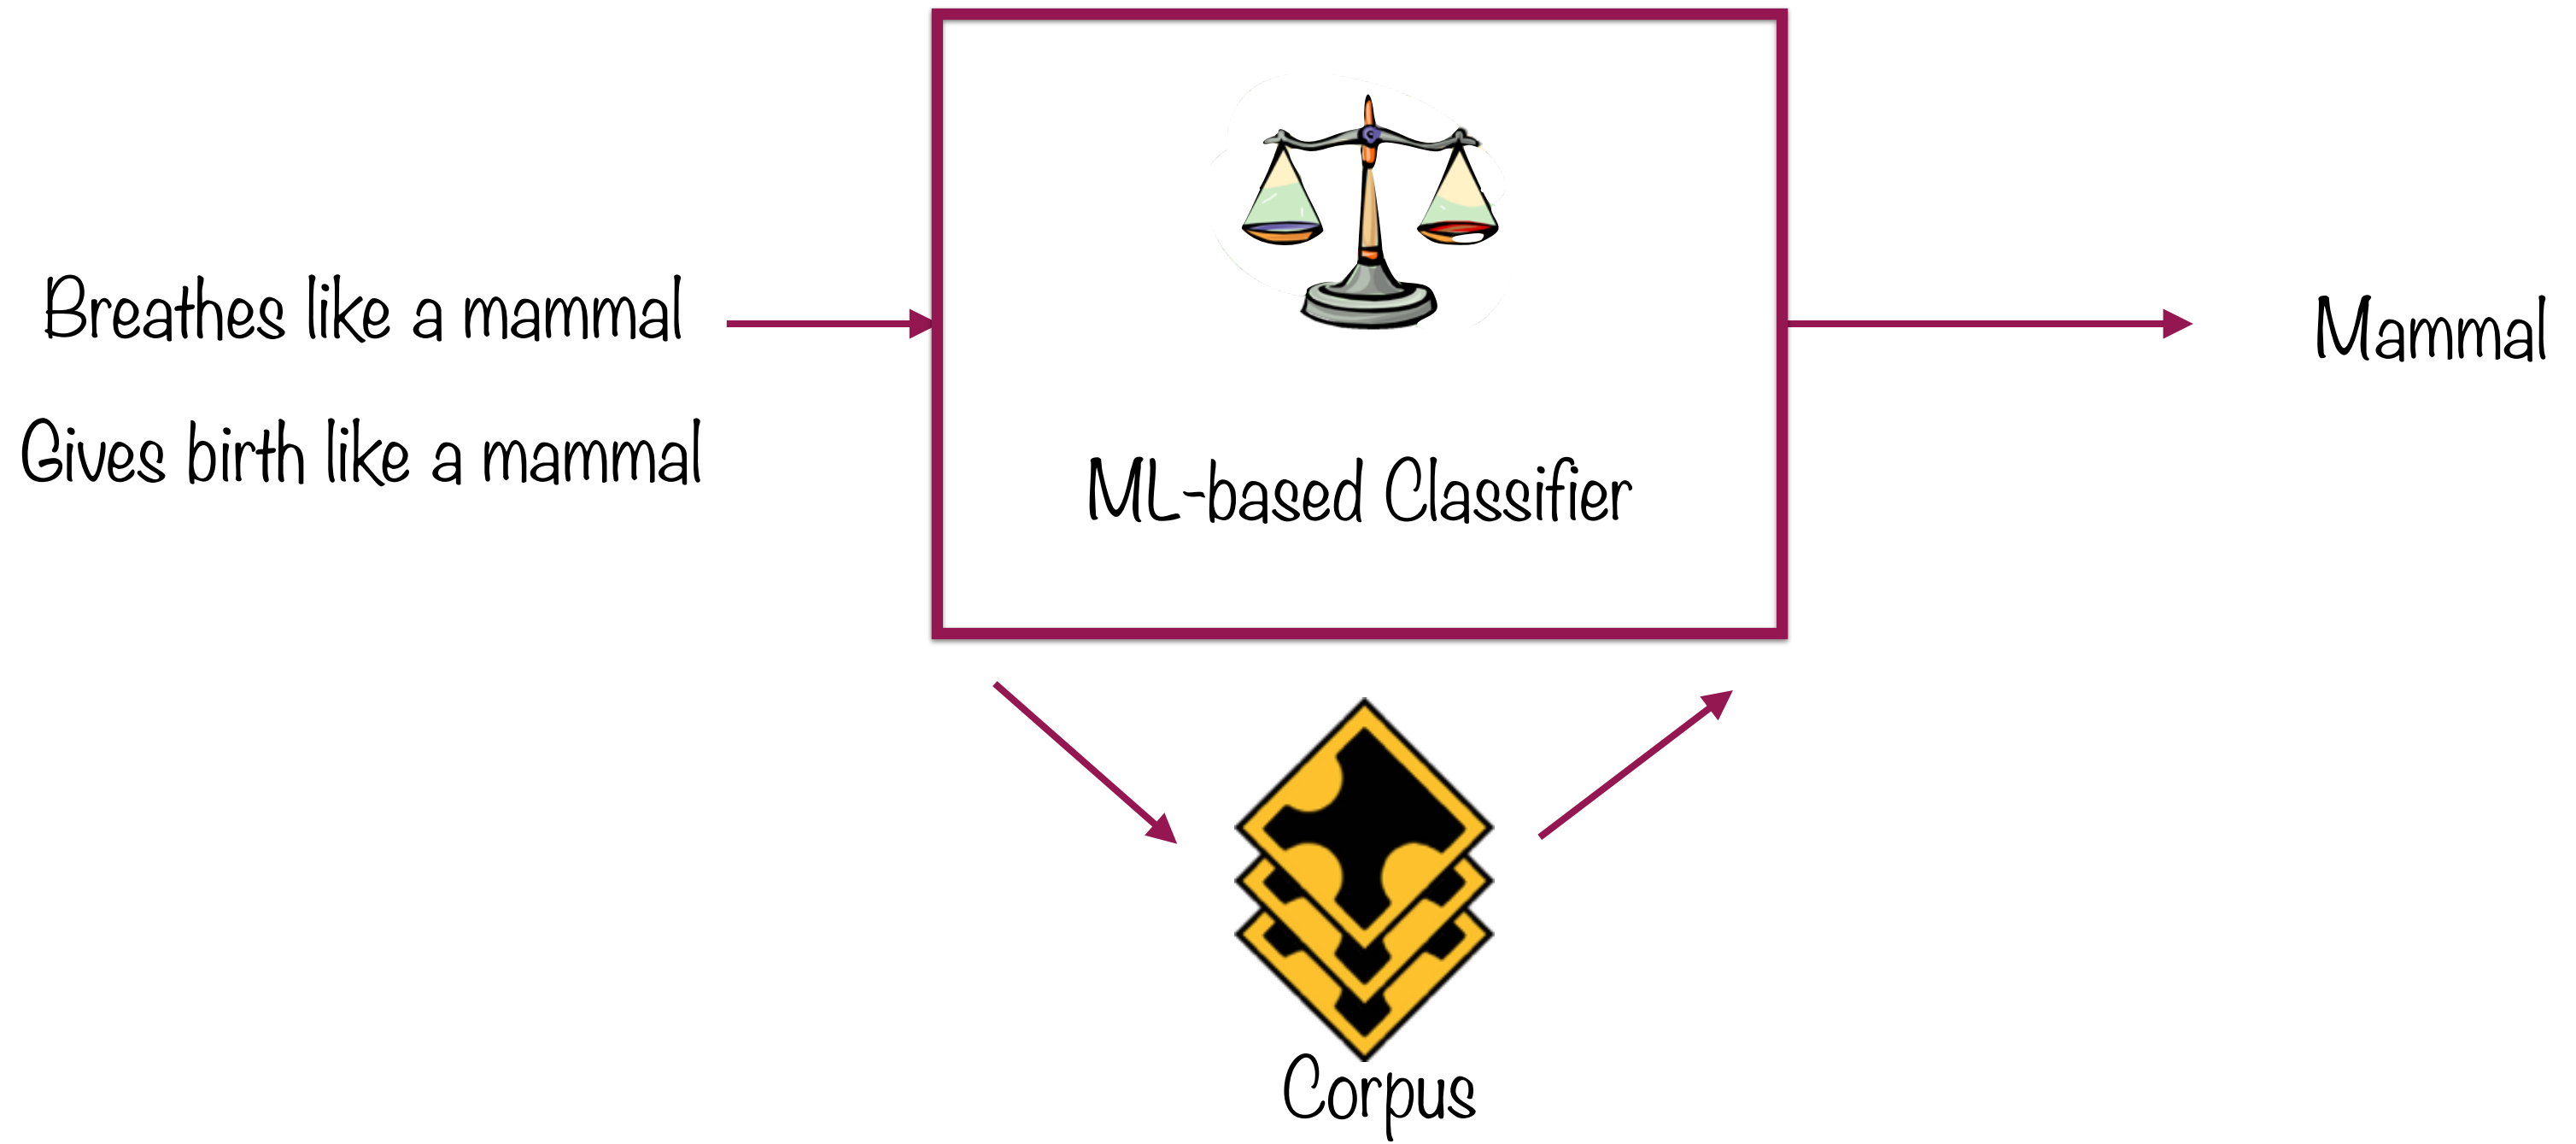
Corpus

Feature Selection by
Experts

Classification
Algorithm

ML-based Classifier

“Traditional” ML-based Binary Classifier



“Traditional” ML-based systems still rely on experts to decide what features to pay attention to

“Traditional” ML-based Binary Classifier



Corpus



Feature Selection by
Experts

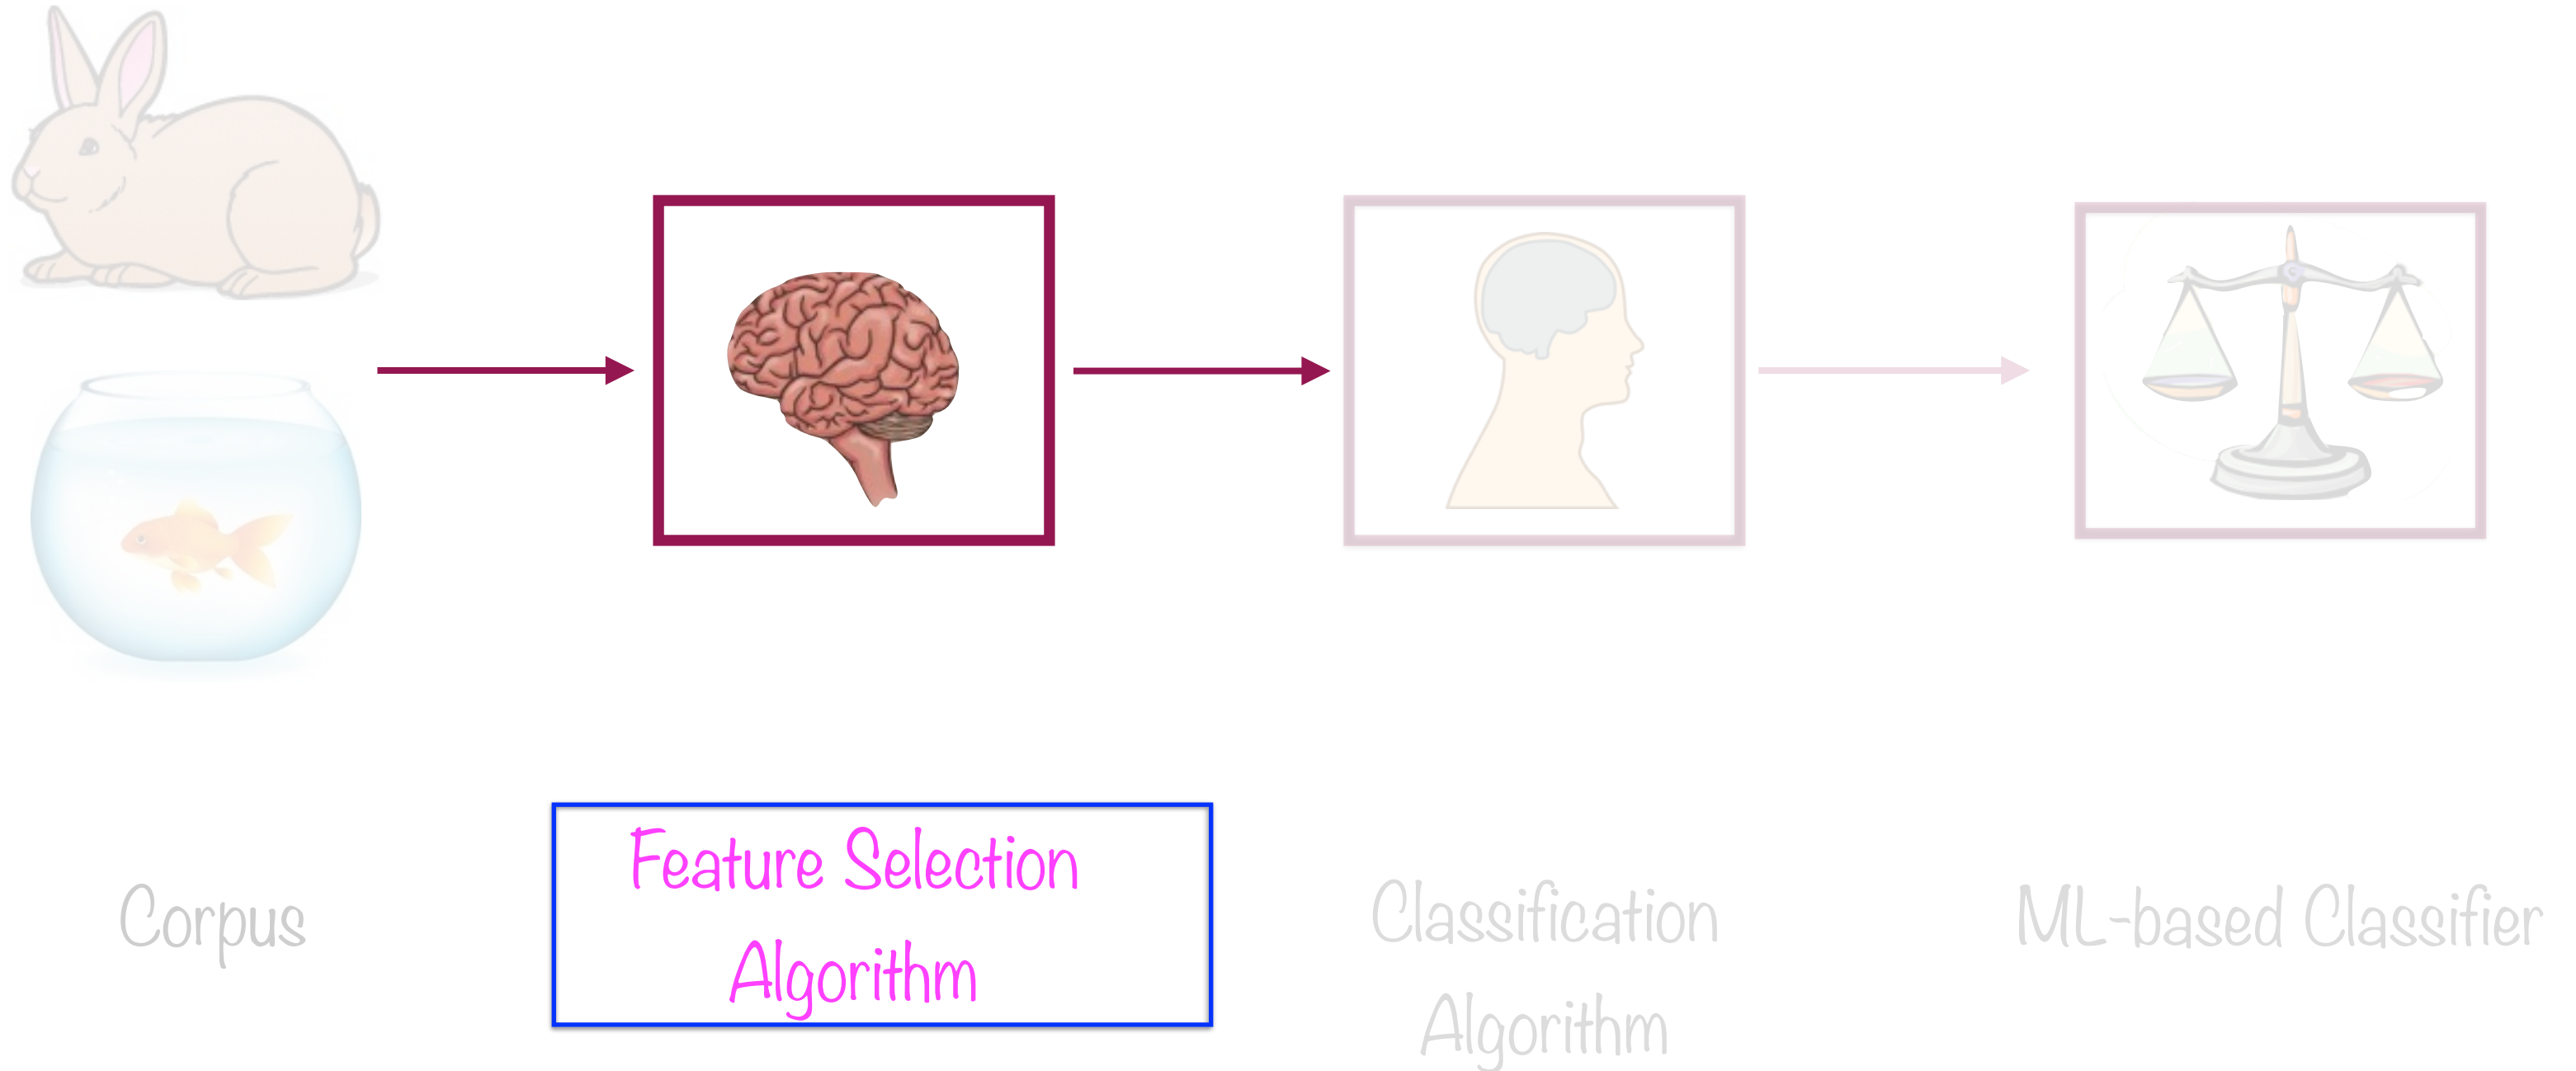


Classification
Algorithm




ML-based Classifier


“Representation” ML-based Binary Classifier



“Representation” ML-based Binary Classifier



Corpus



Feature Selection
Algorithm



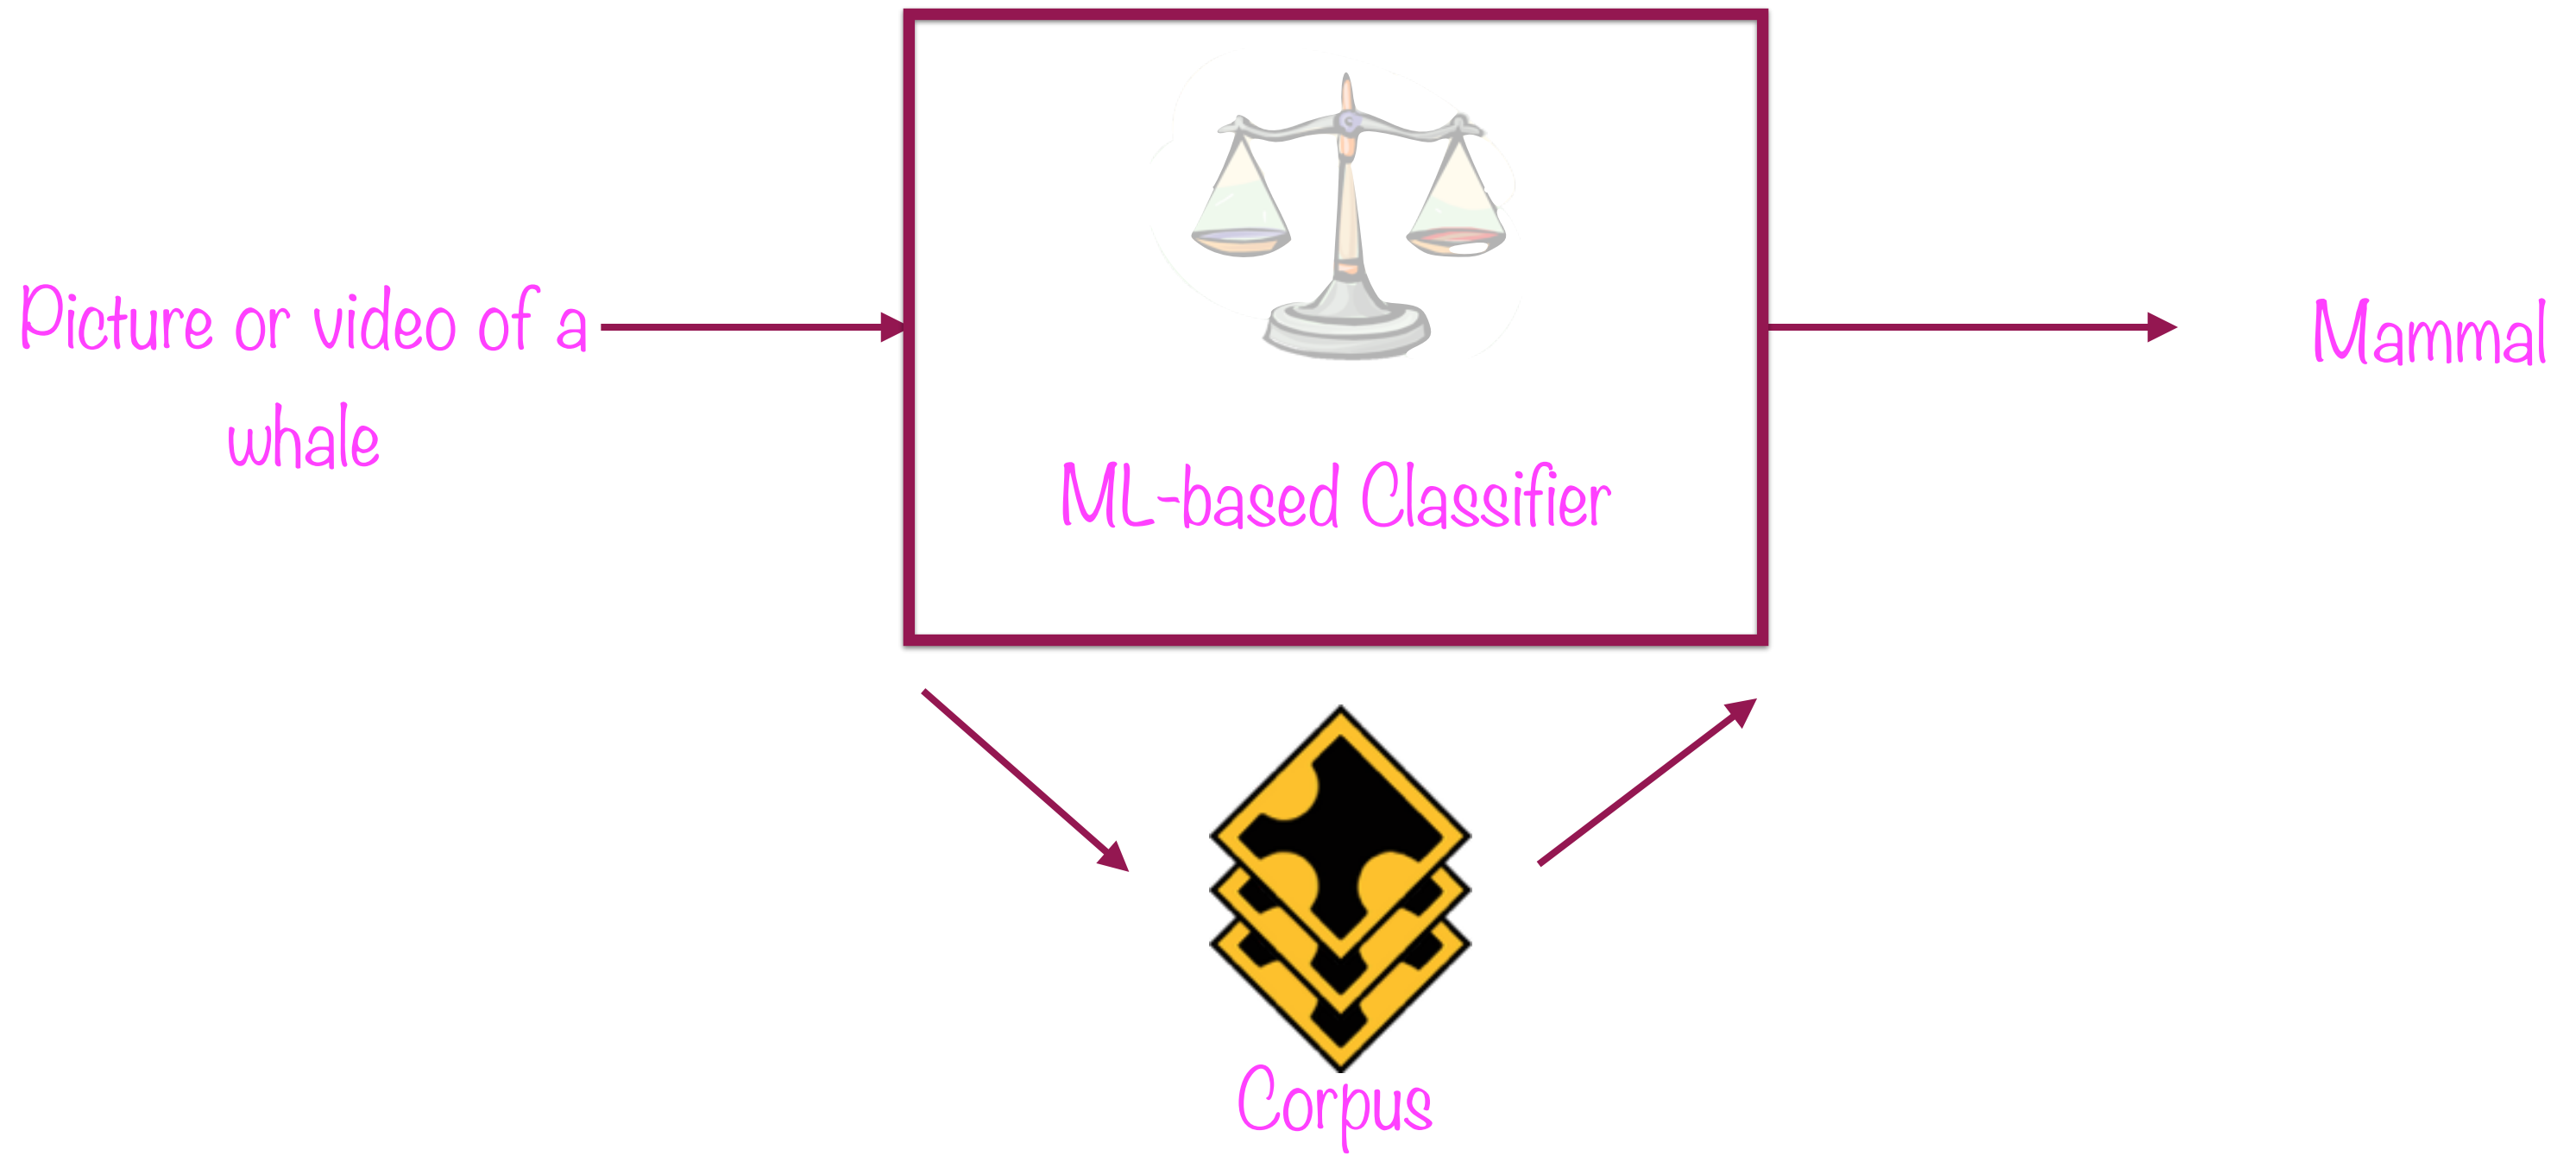
Classification
Algorithm



ML-based Classifier

“Representation” ML-based systems figure out by themselves what features to pay attention to

“Representation” ML-based Binary Classifier



**“Deep Learning” systems are one type of
representation systems**

Deep Learning and Neural Networks

Deep Learning

Algorithms that learn what features matter

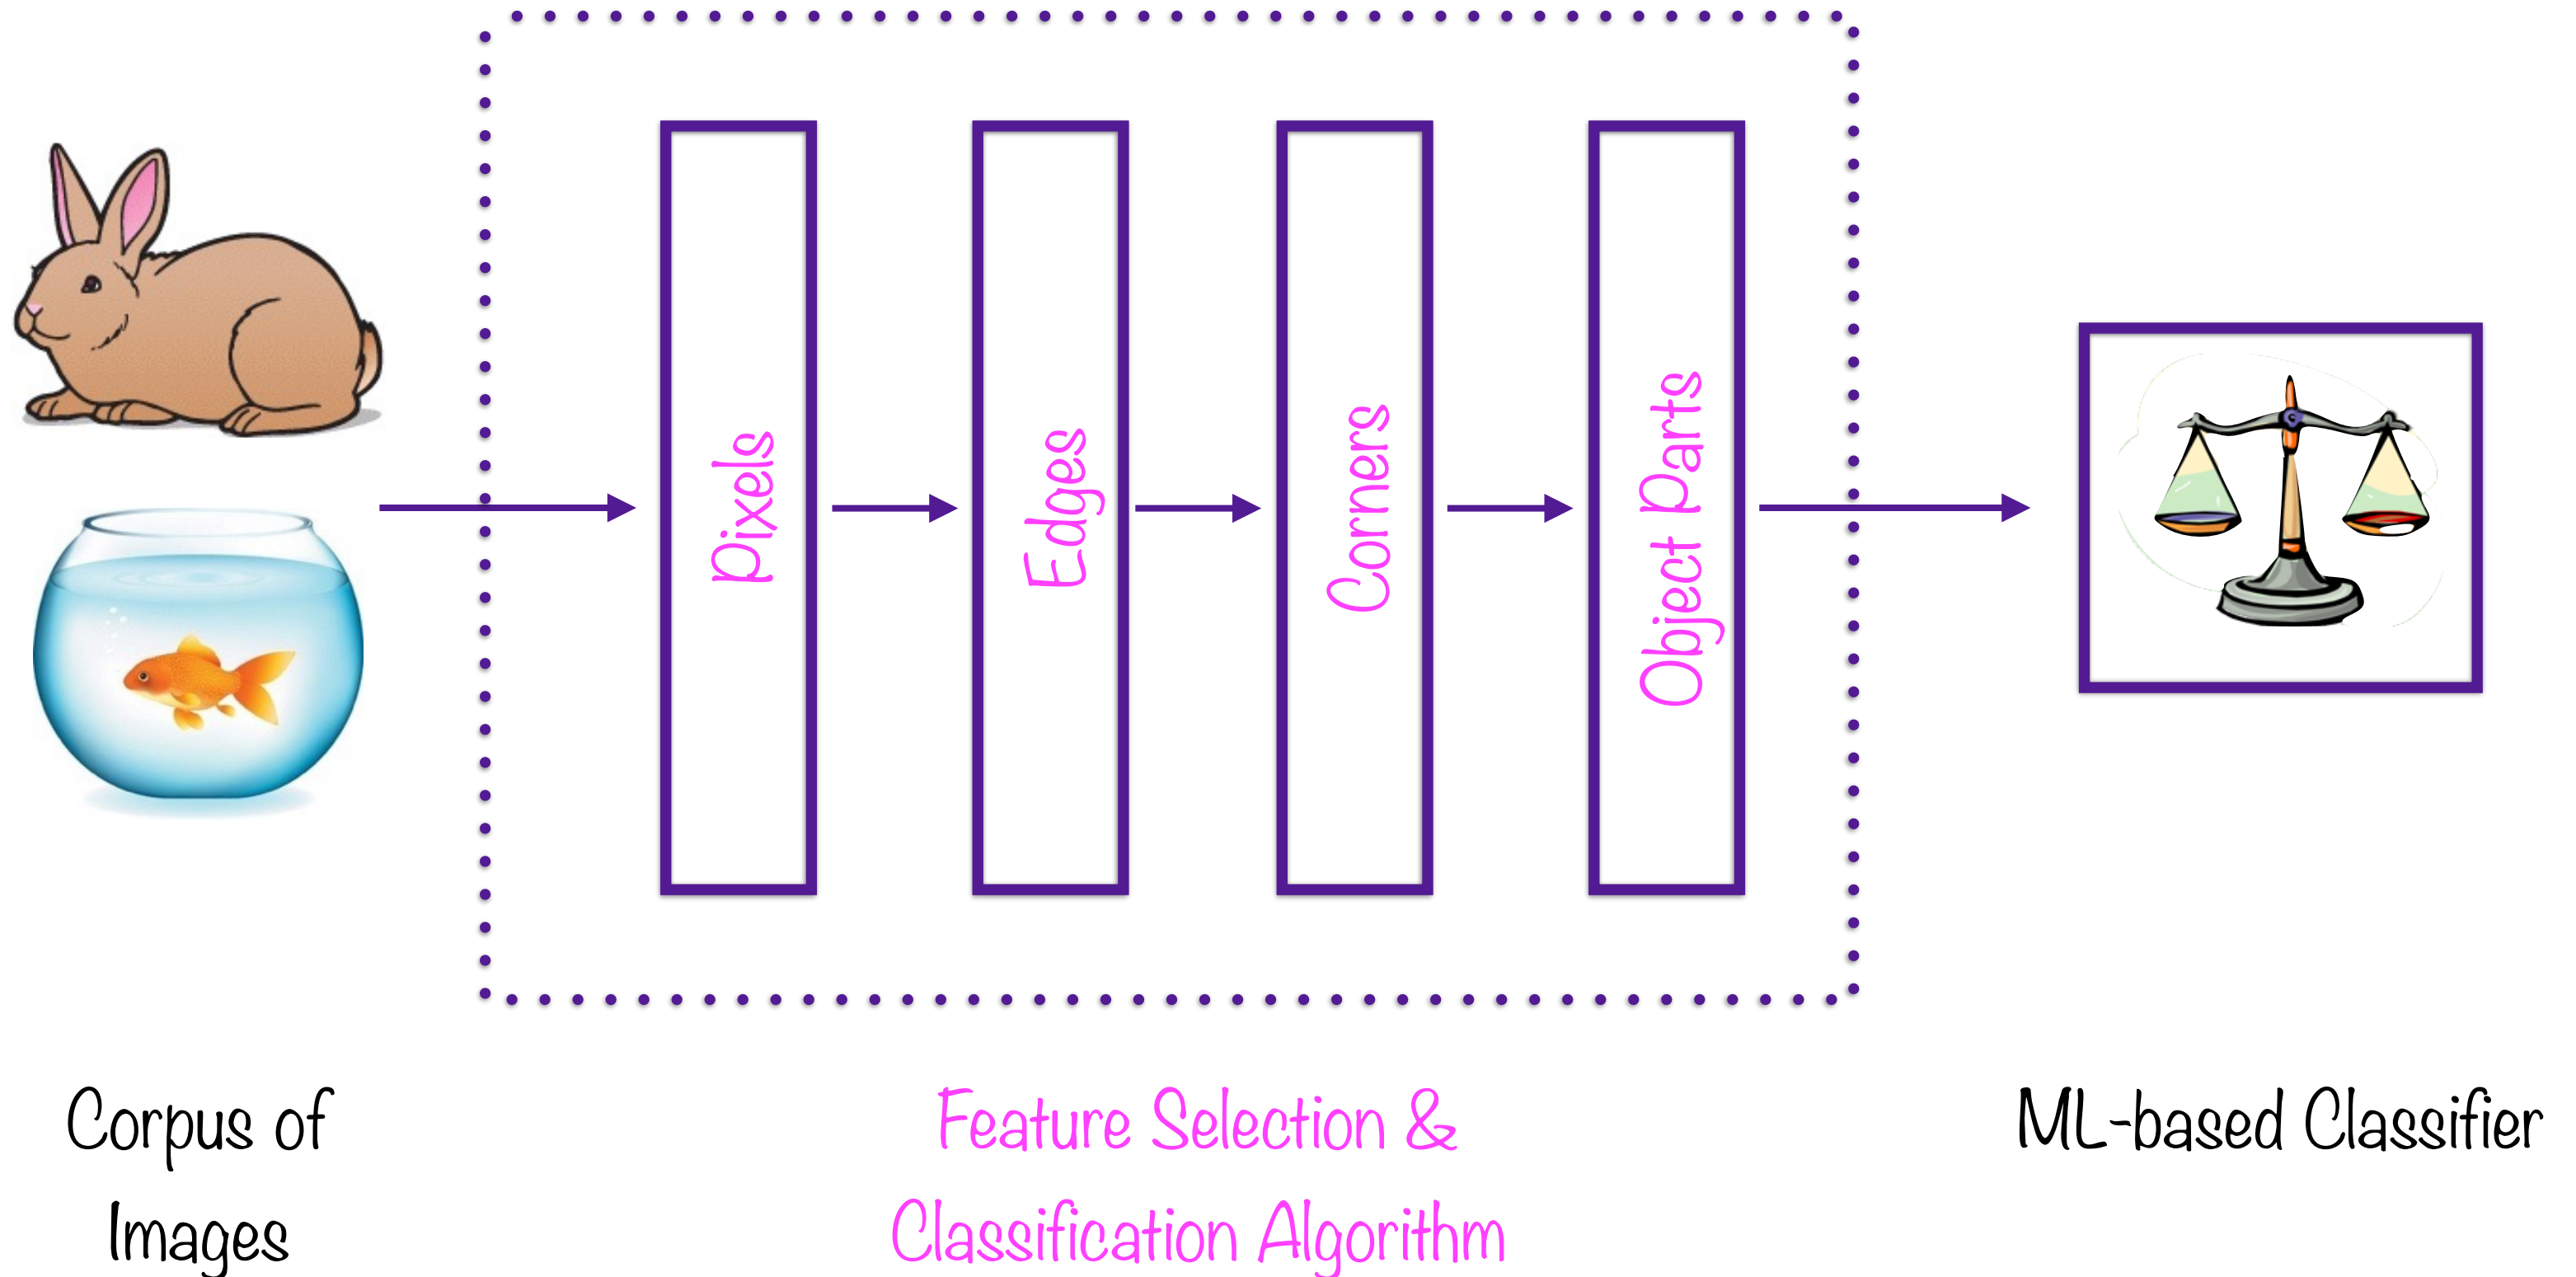
Neural Networks

The most common class of deep learning algorithms


Neurons

Simple building blocks that actually “learn”


"Deep Learning"-based Binary Classifier



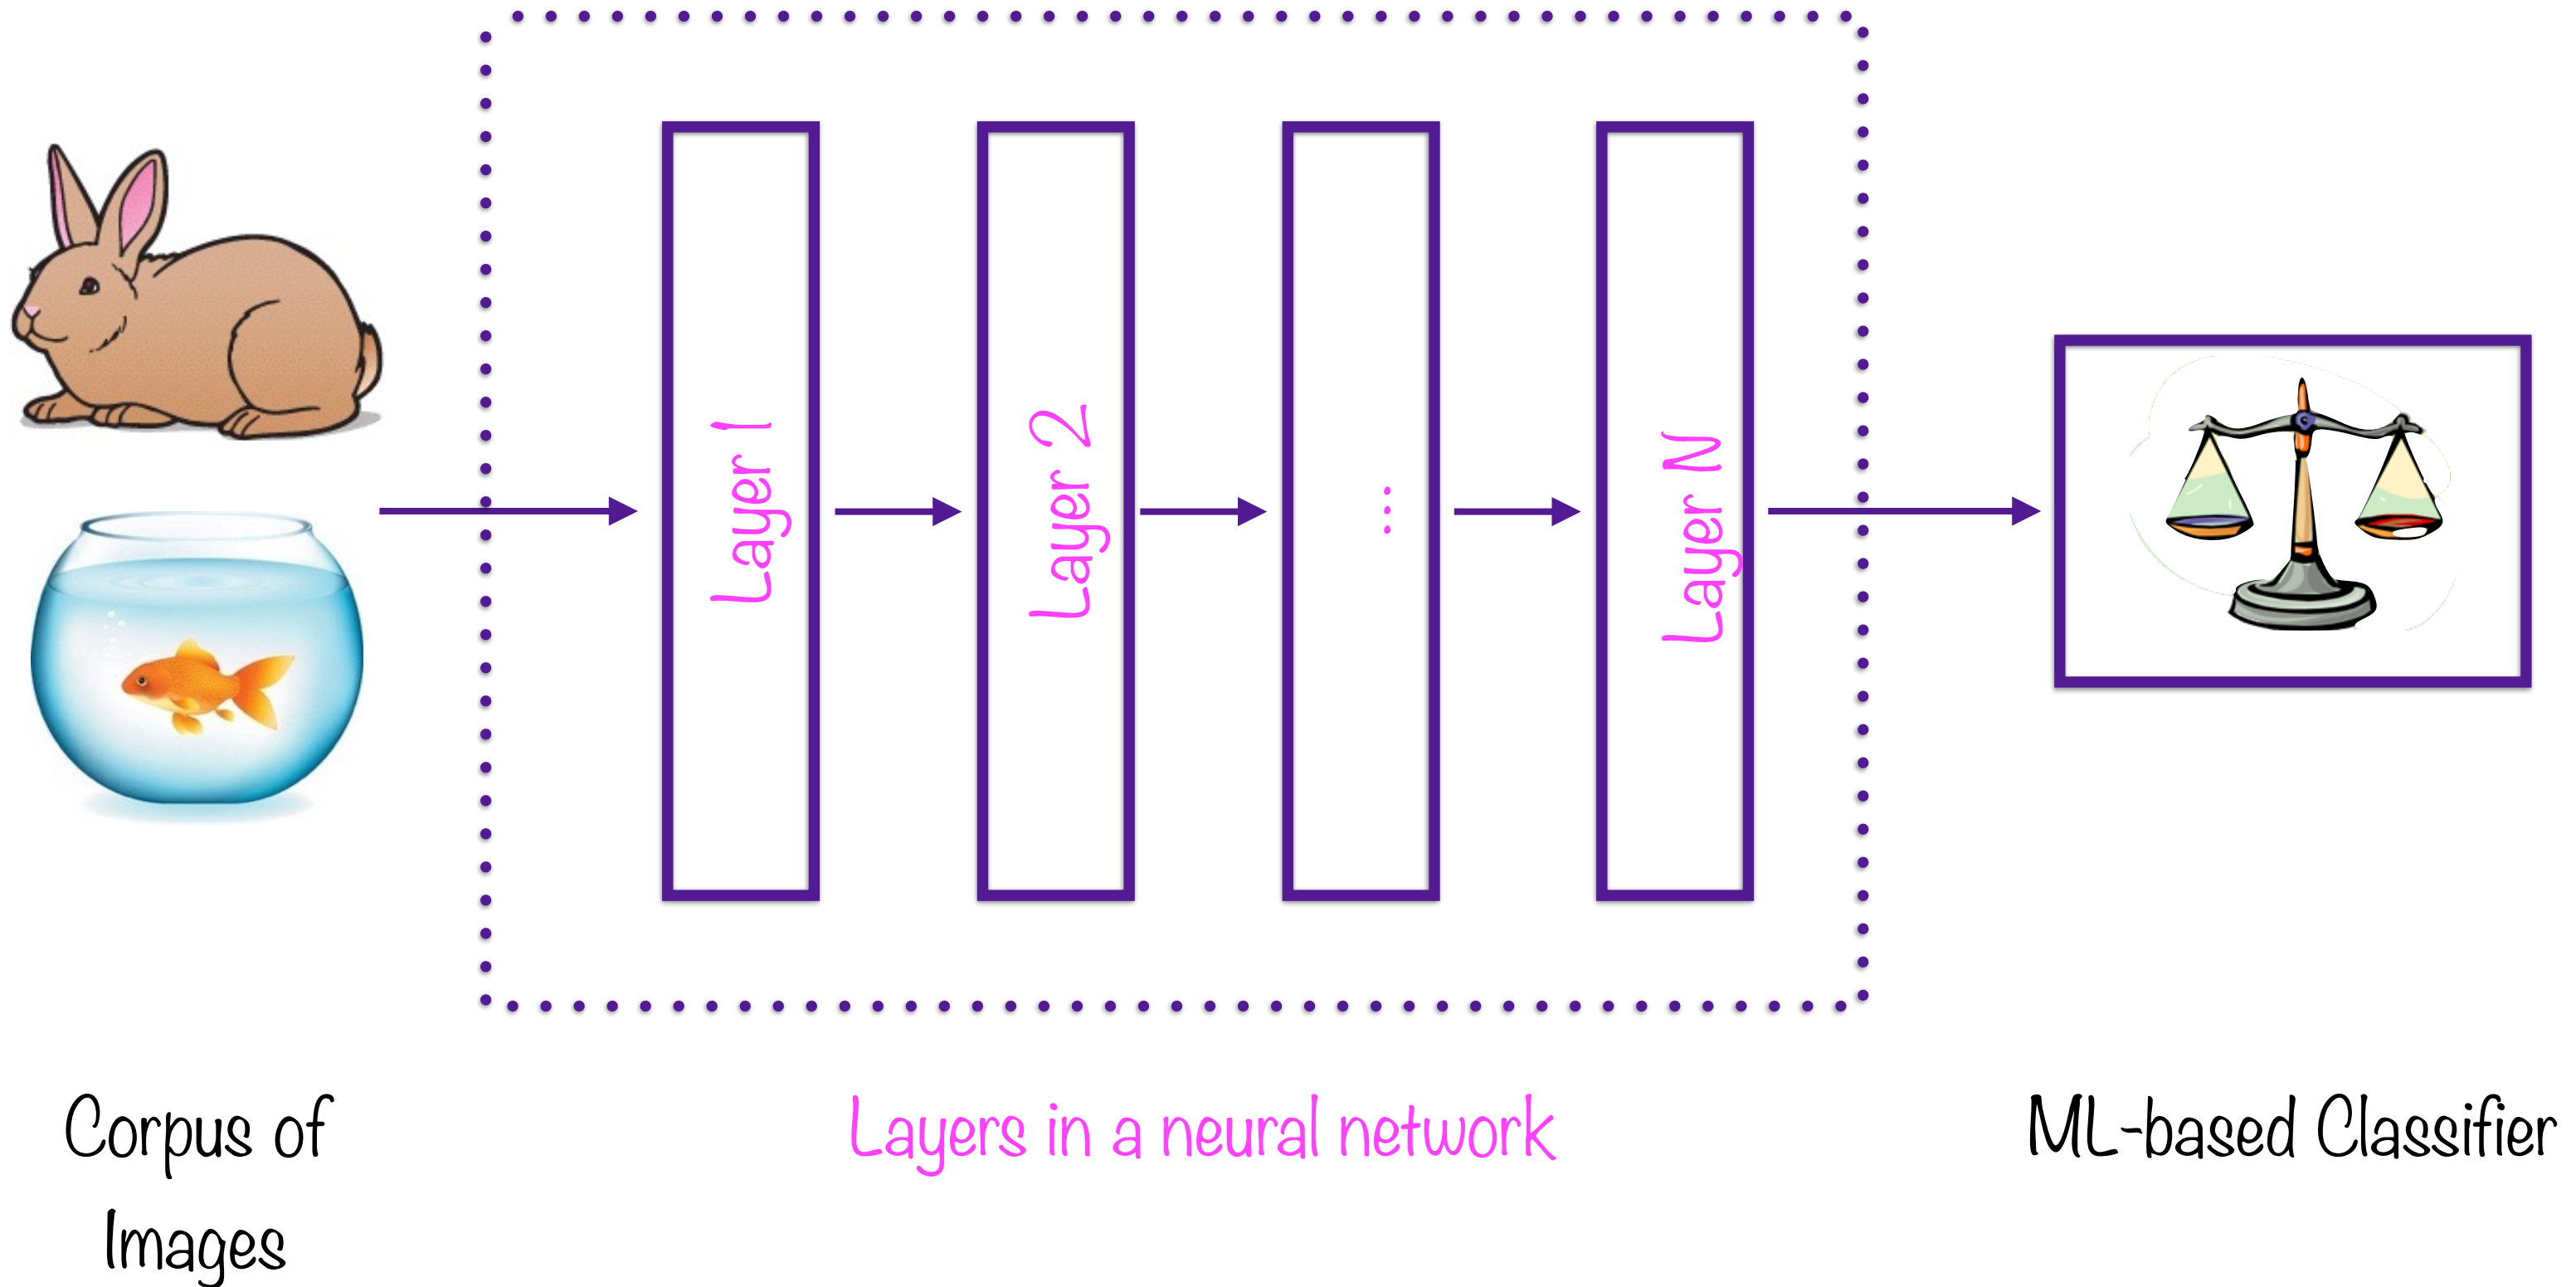
“Deep Learning”-based Binary Classifier



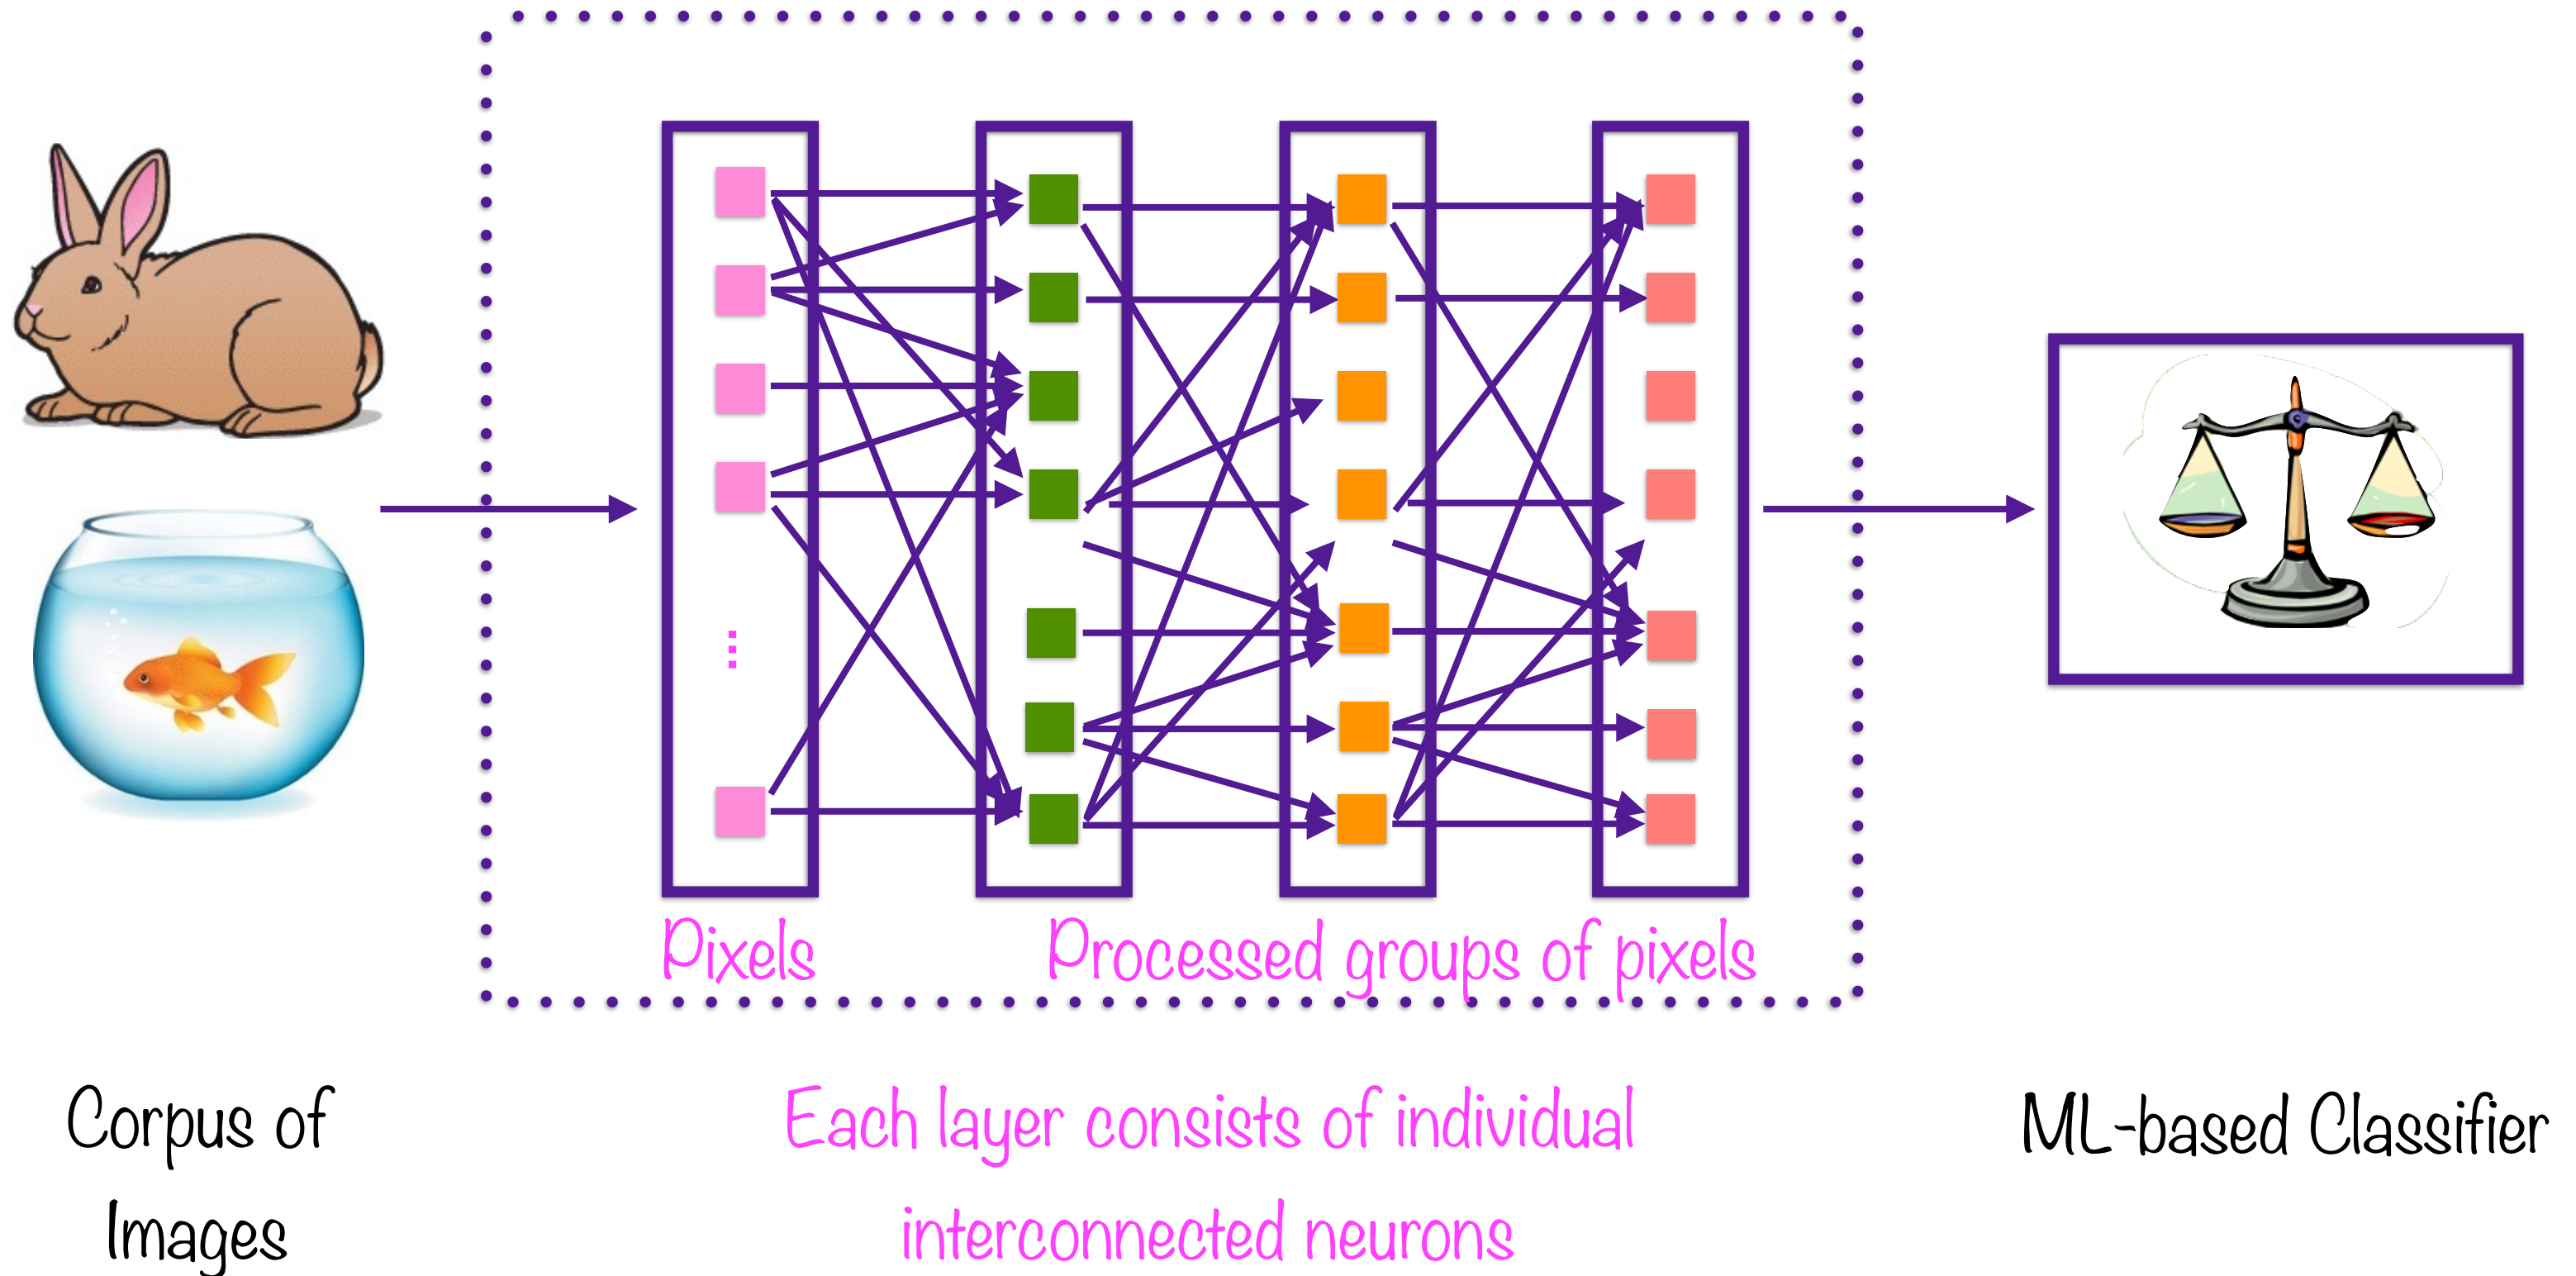
“Deep Learning”-based Binary Classifier



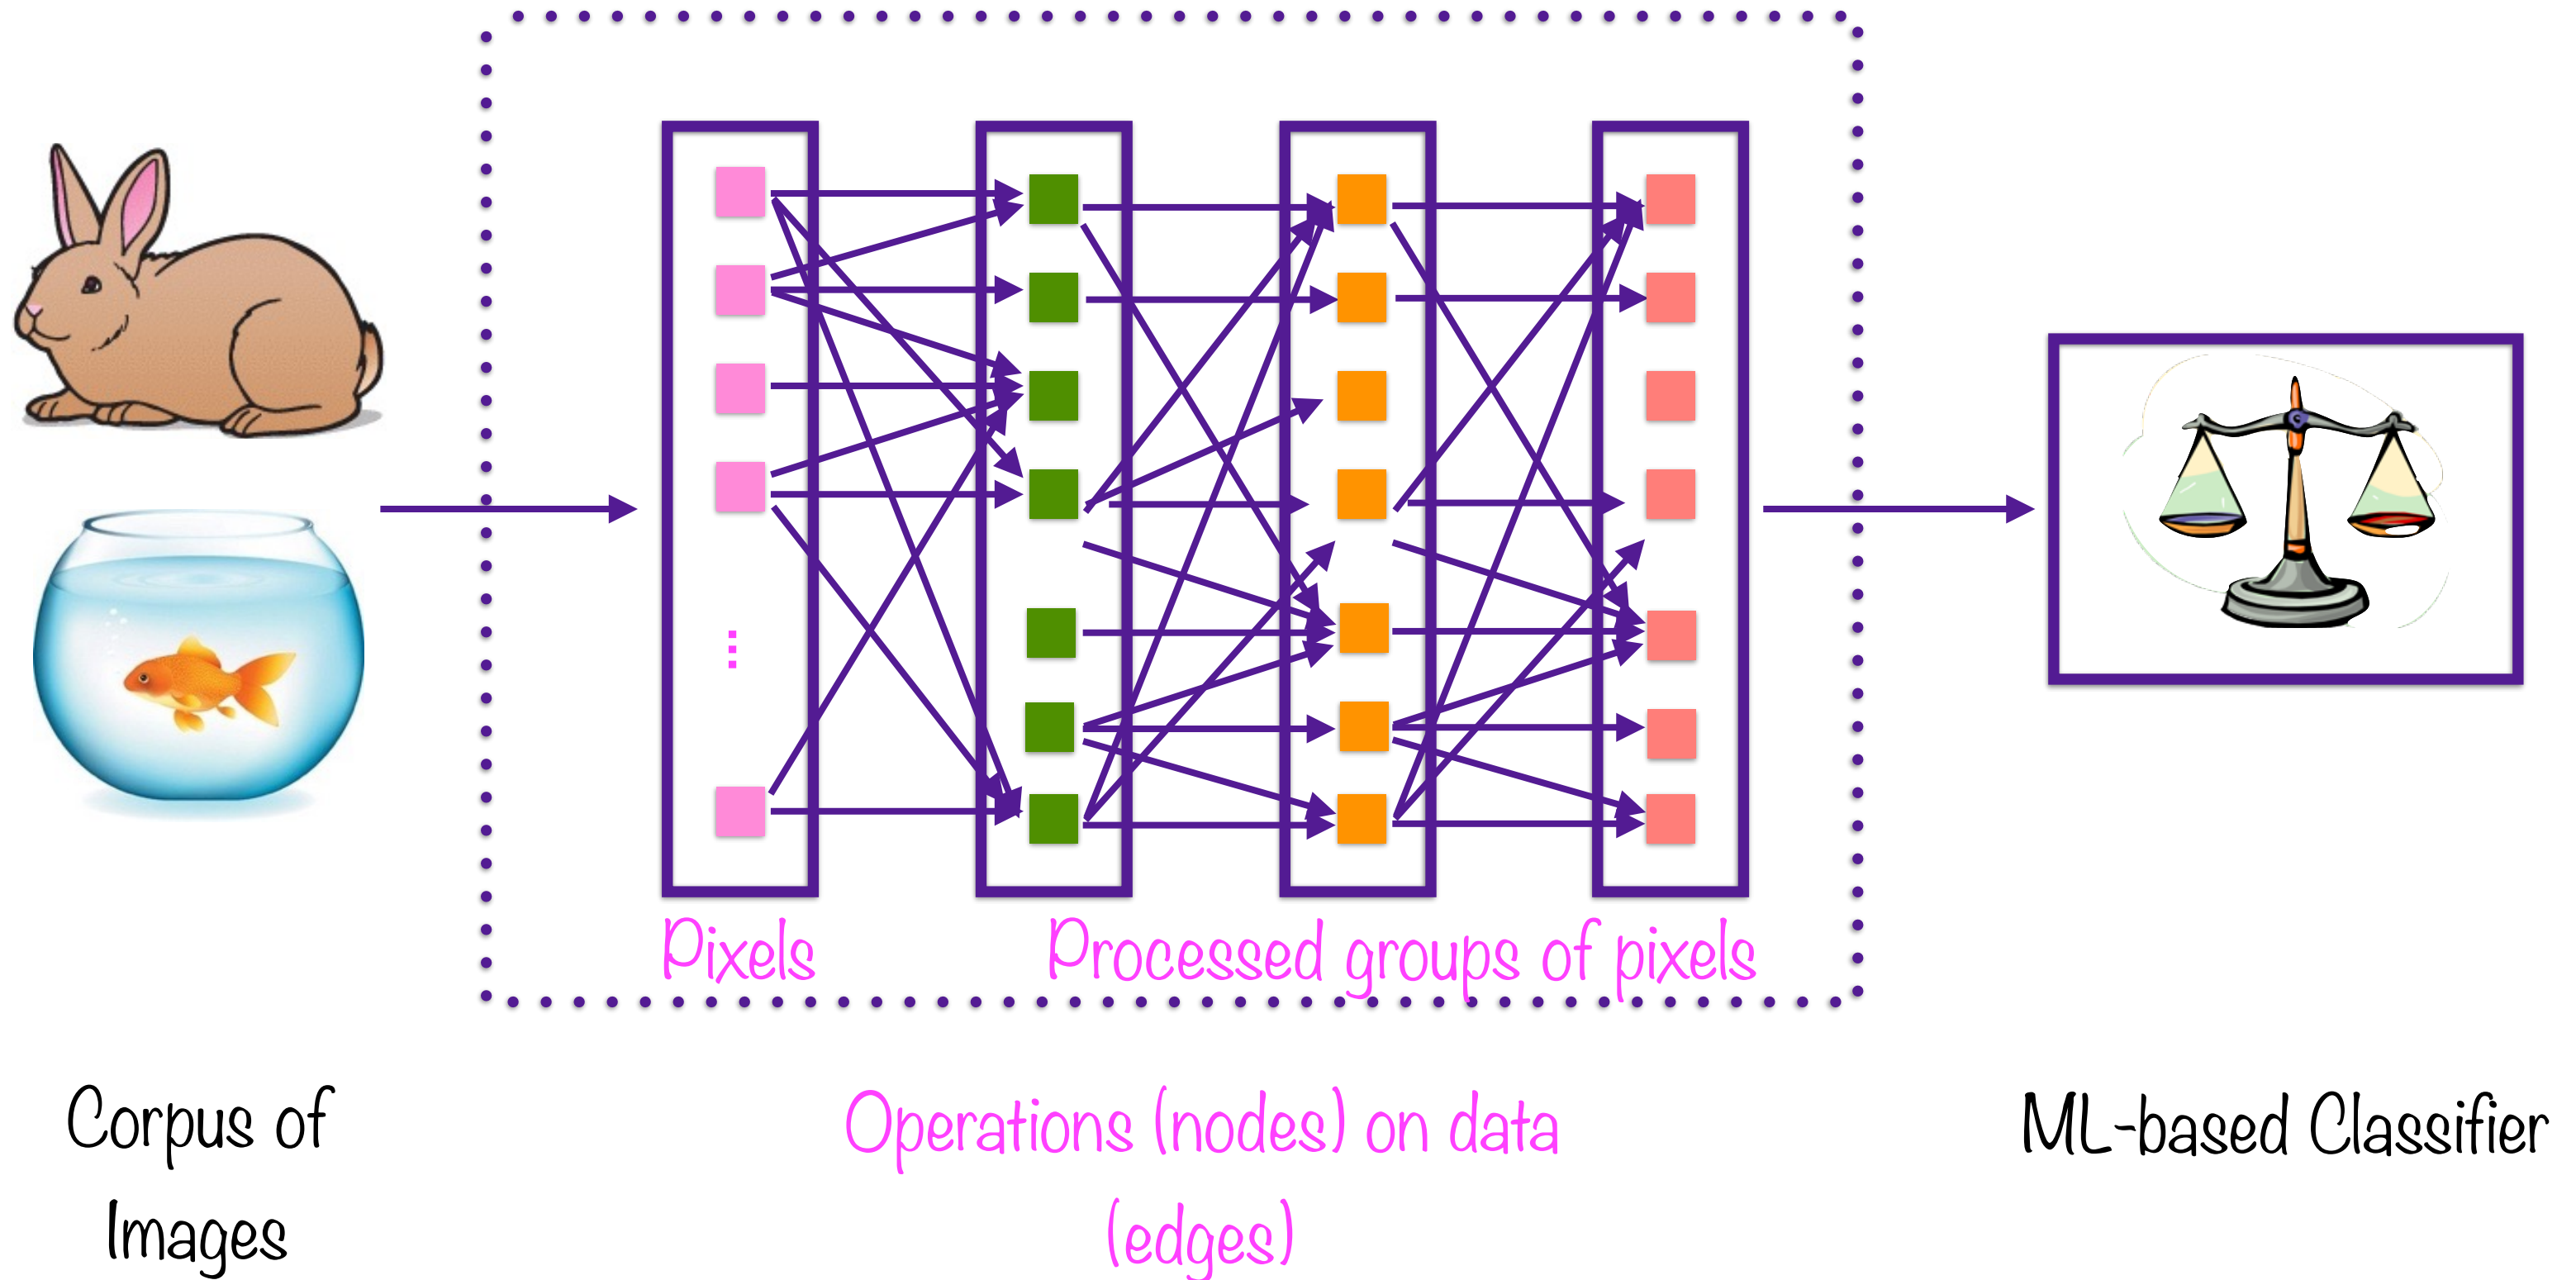
Neural Networks Introduced




Neural Networks Introduced



The Computational Graph



The Computational Graph




Corpus of
Images

The nodes in the computation graph are neurons
(simple building blocks)

ML-based Classifier

The Computational Graph



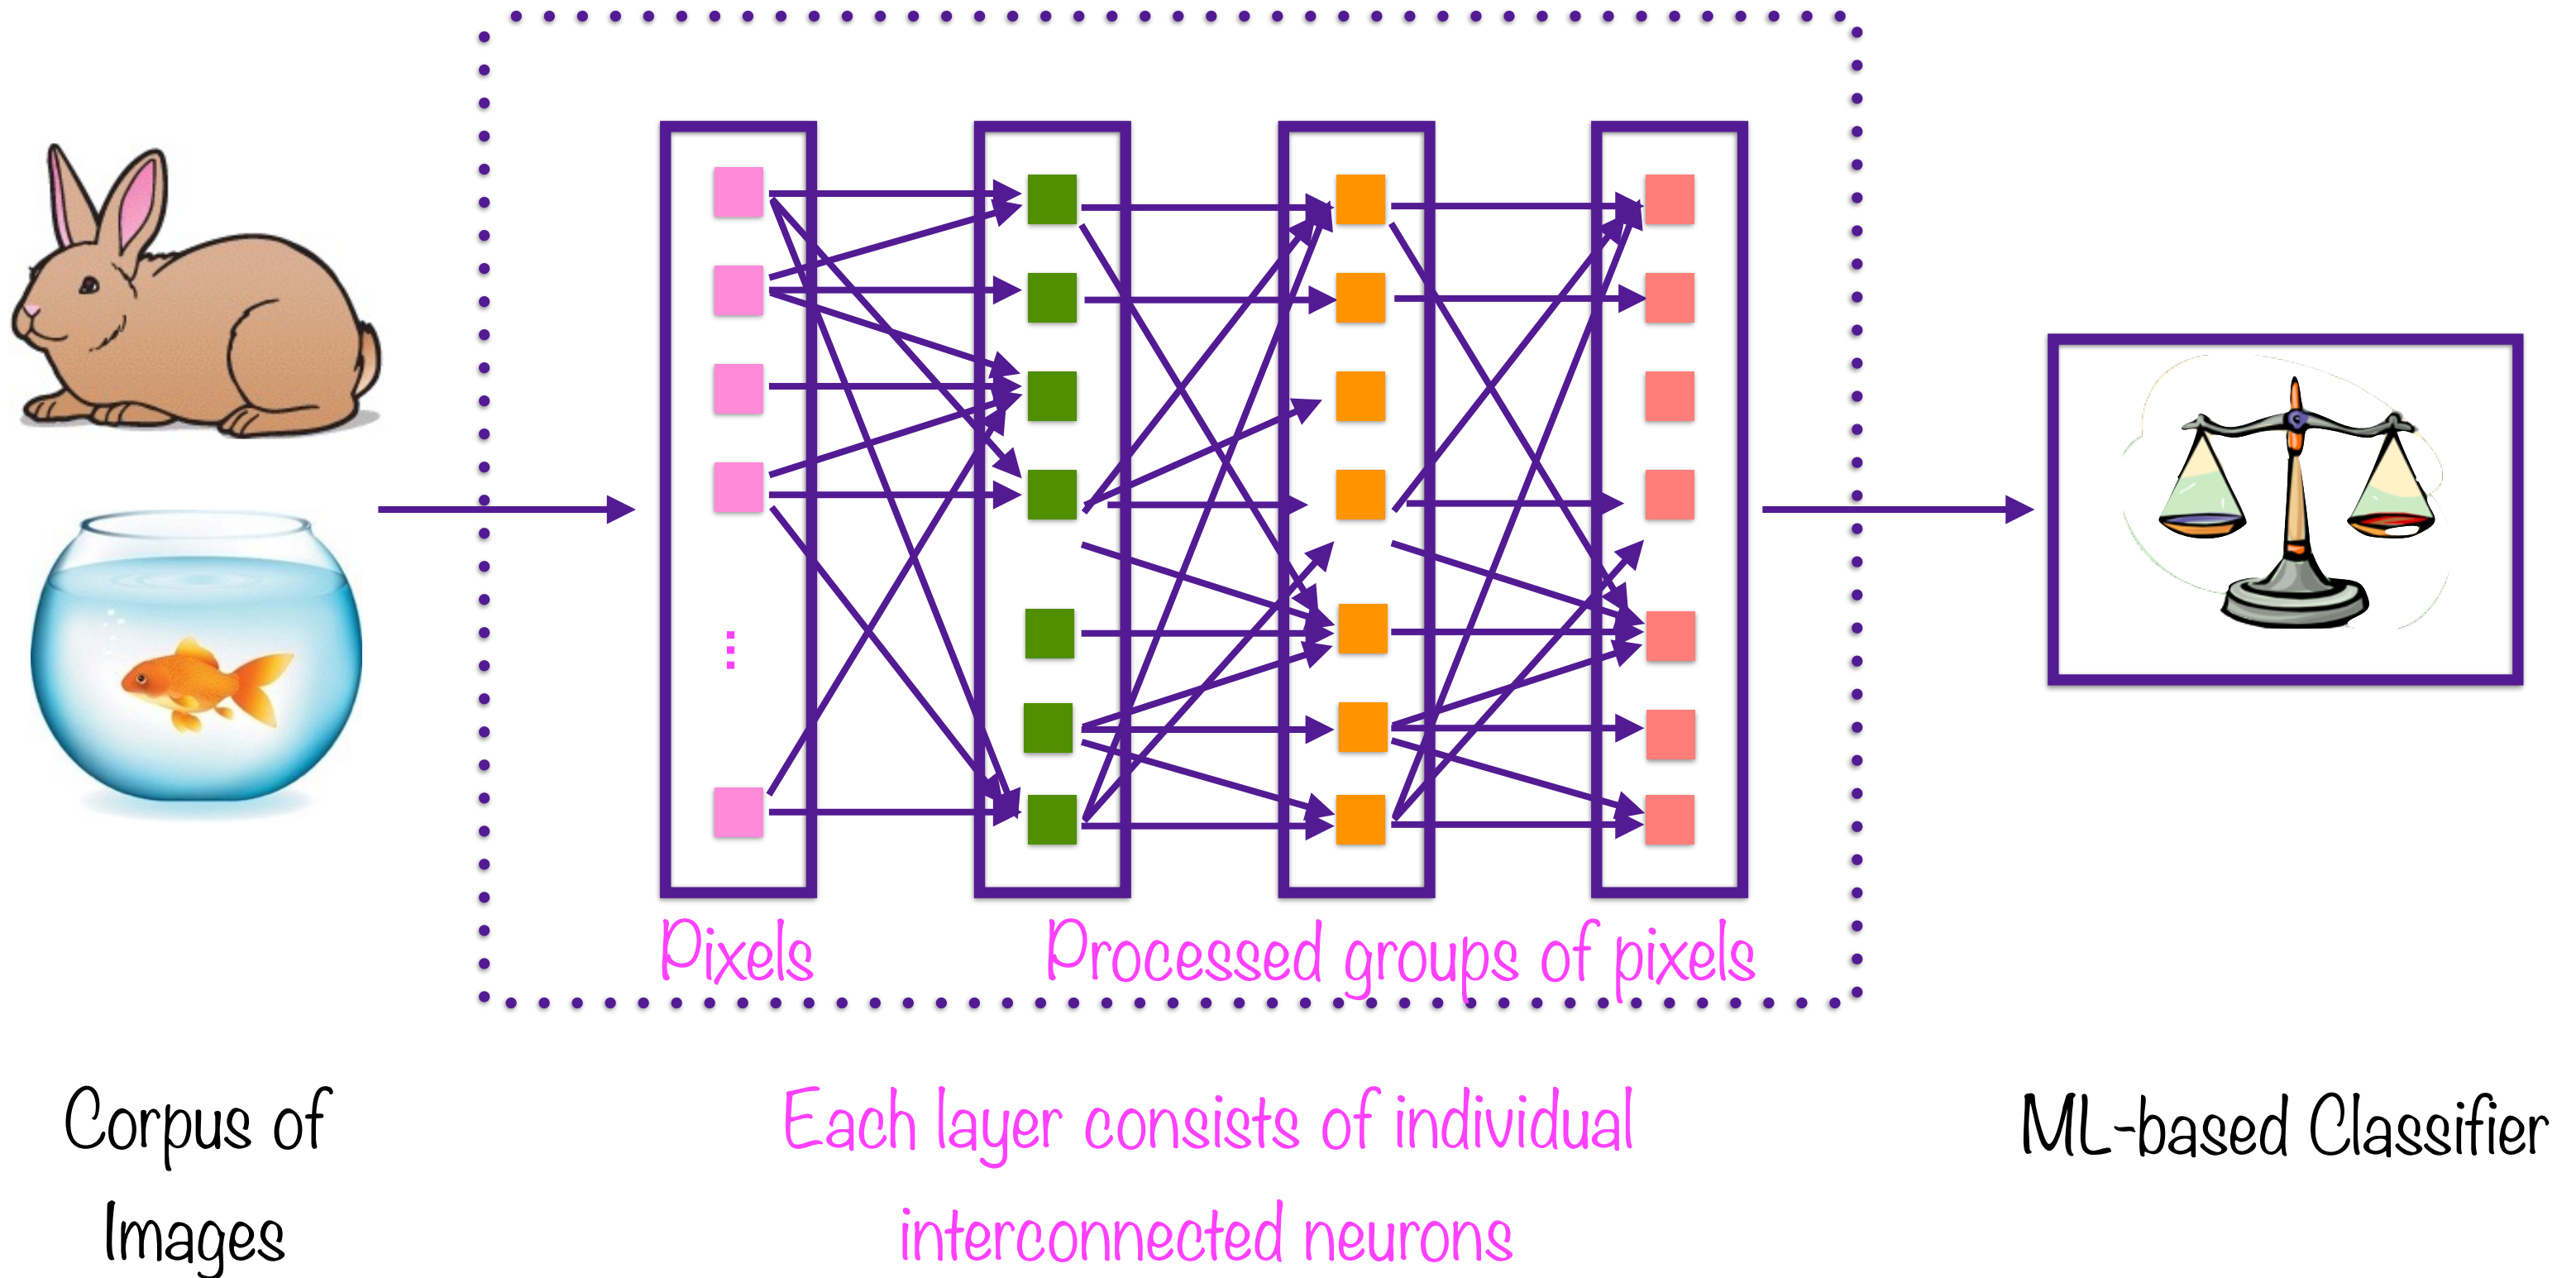
Corpus of
Images

The edges in the computation graph are data items
called tensors


ML-based Classifier

Neuron as a Learning Unit

A Neural Network




Operation of a Single Neuron




For an active neuron a change in inputs should trigger a corresponding change in the outputs

Operation of a Single Neuron




The outputs of neurons feed into the neurons from the next layer

Operation of a Single Neuron



Each connection is associated with a weight


Operation of a Single Neuron



If the second neuron is sensitive to the output of the first neuron the connection between them gets stronger


W increases

Operation of a Single Neuron



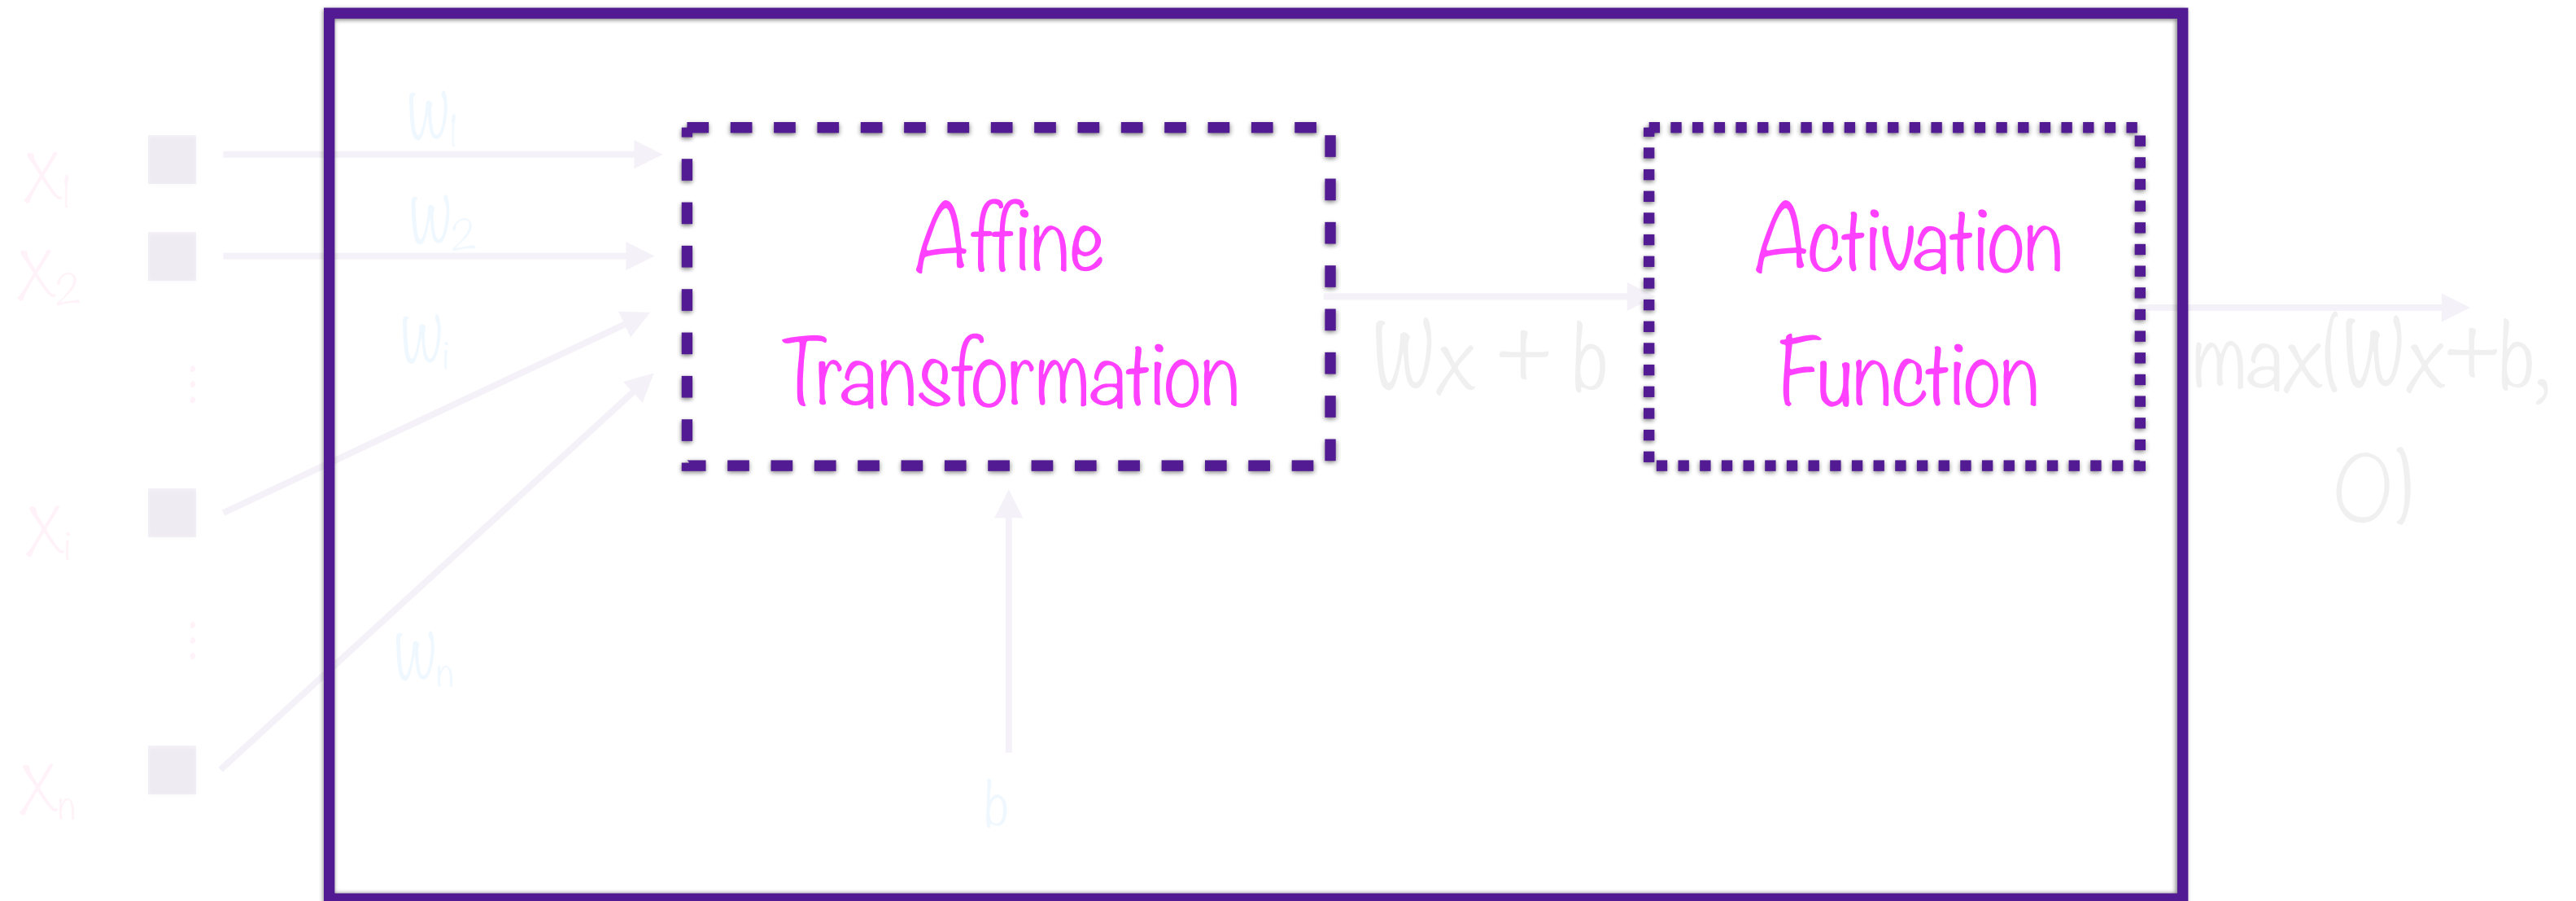
A single neuron is generally connected to multiple neurons in the next layer.
Each connection will have its **own weight**

A Neural Network



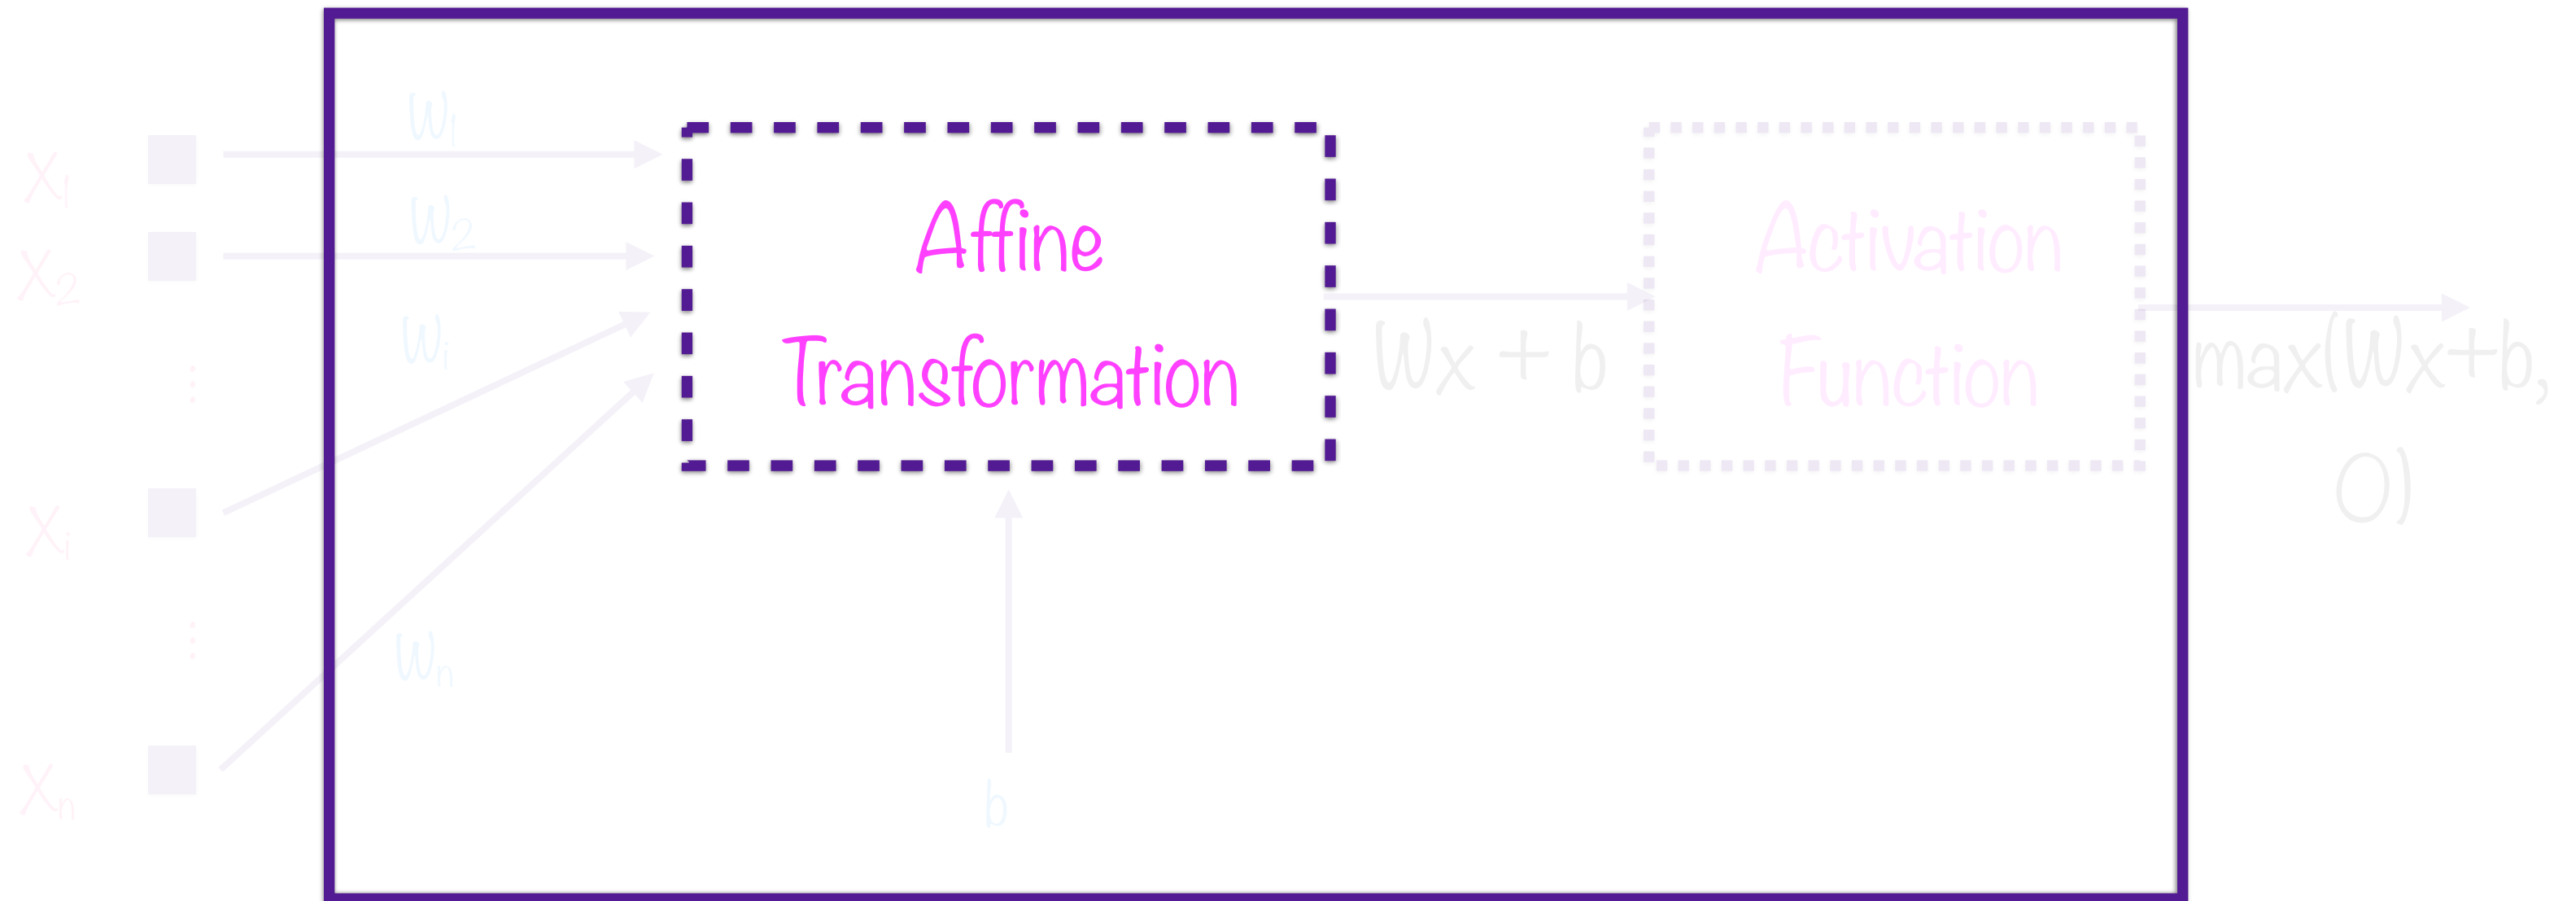
Once a neural network is **trained** all edges have weights which help it make predictions

Operation of a Single Neuron



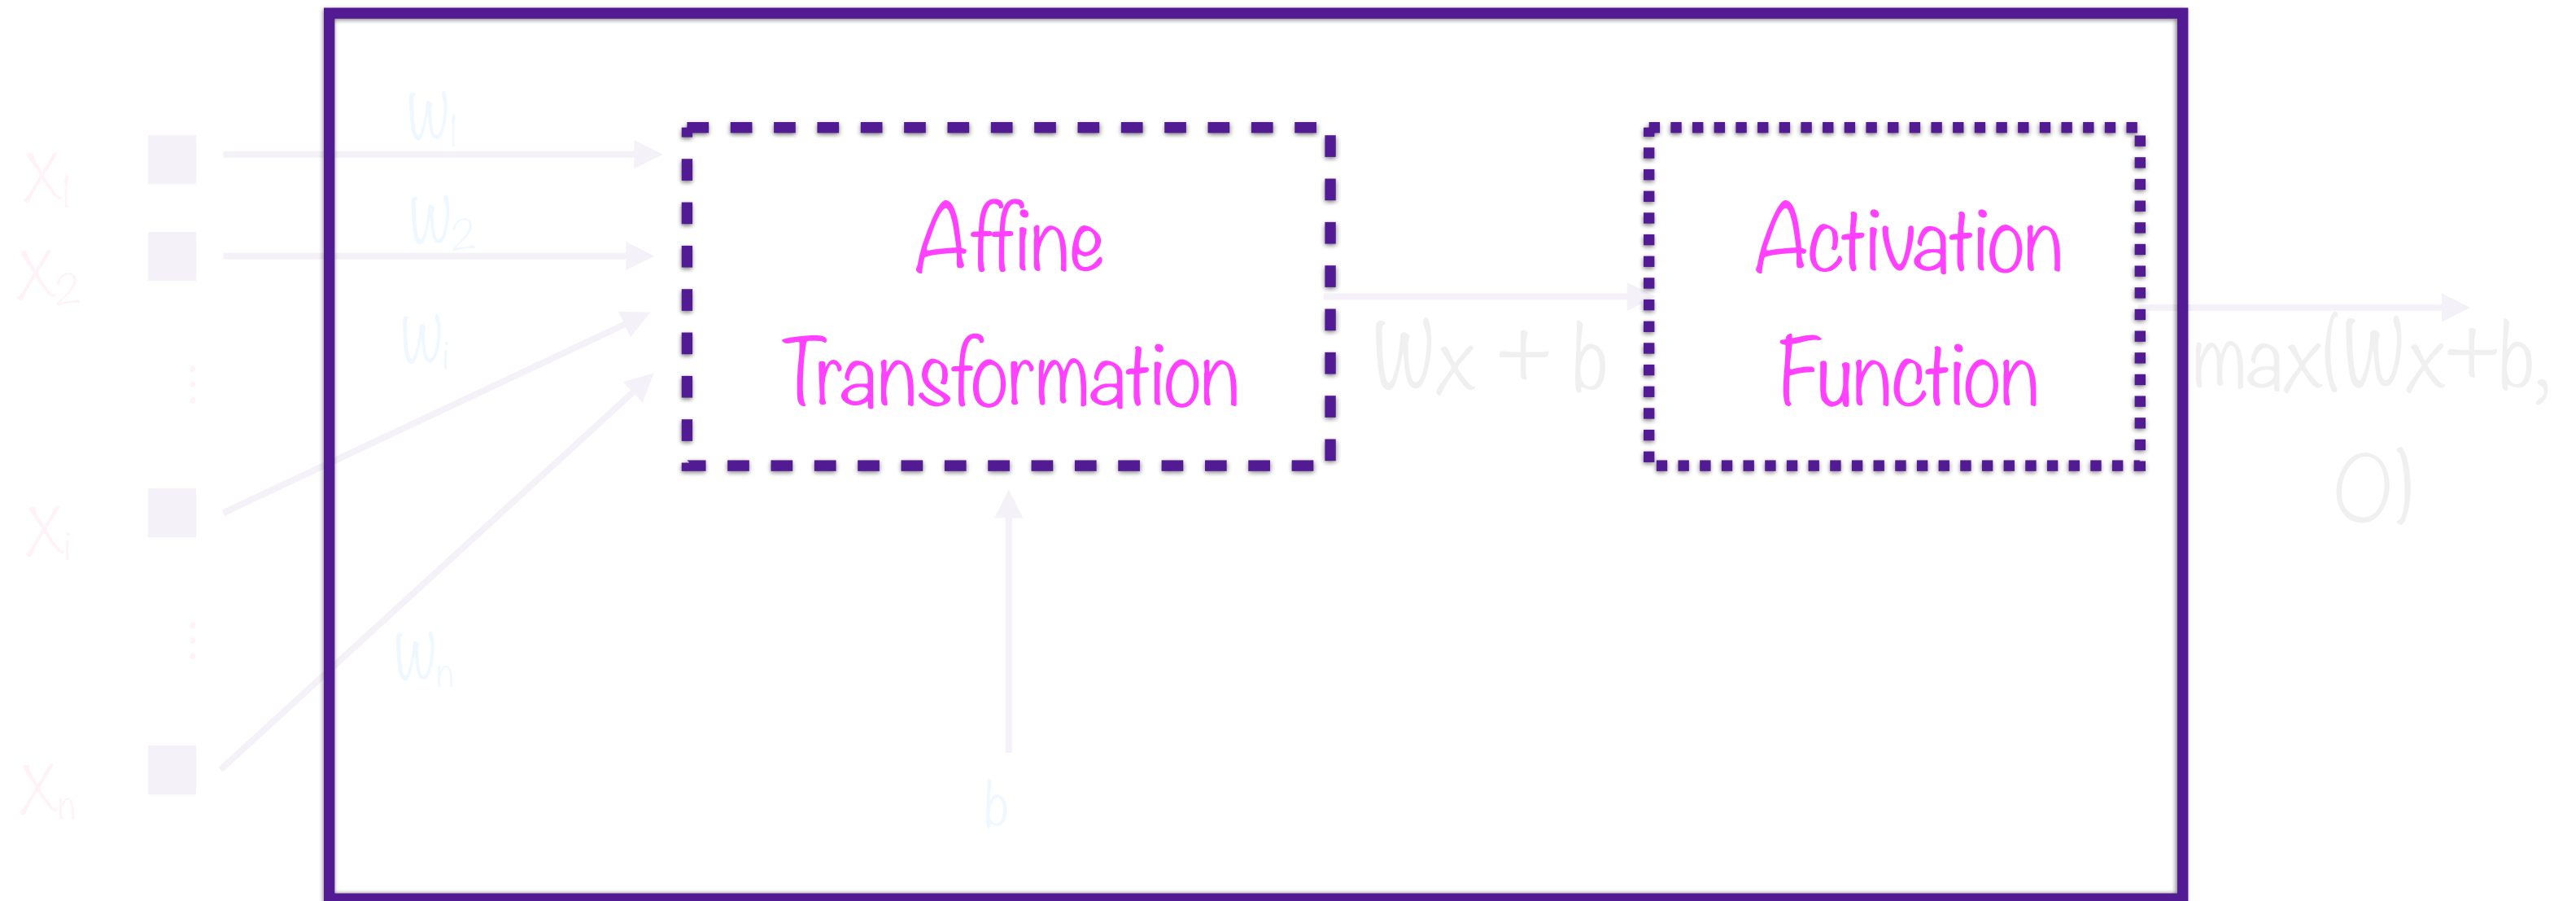
Each neuron only applies two simple functions to its inputs

Operation of a Single Neuron



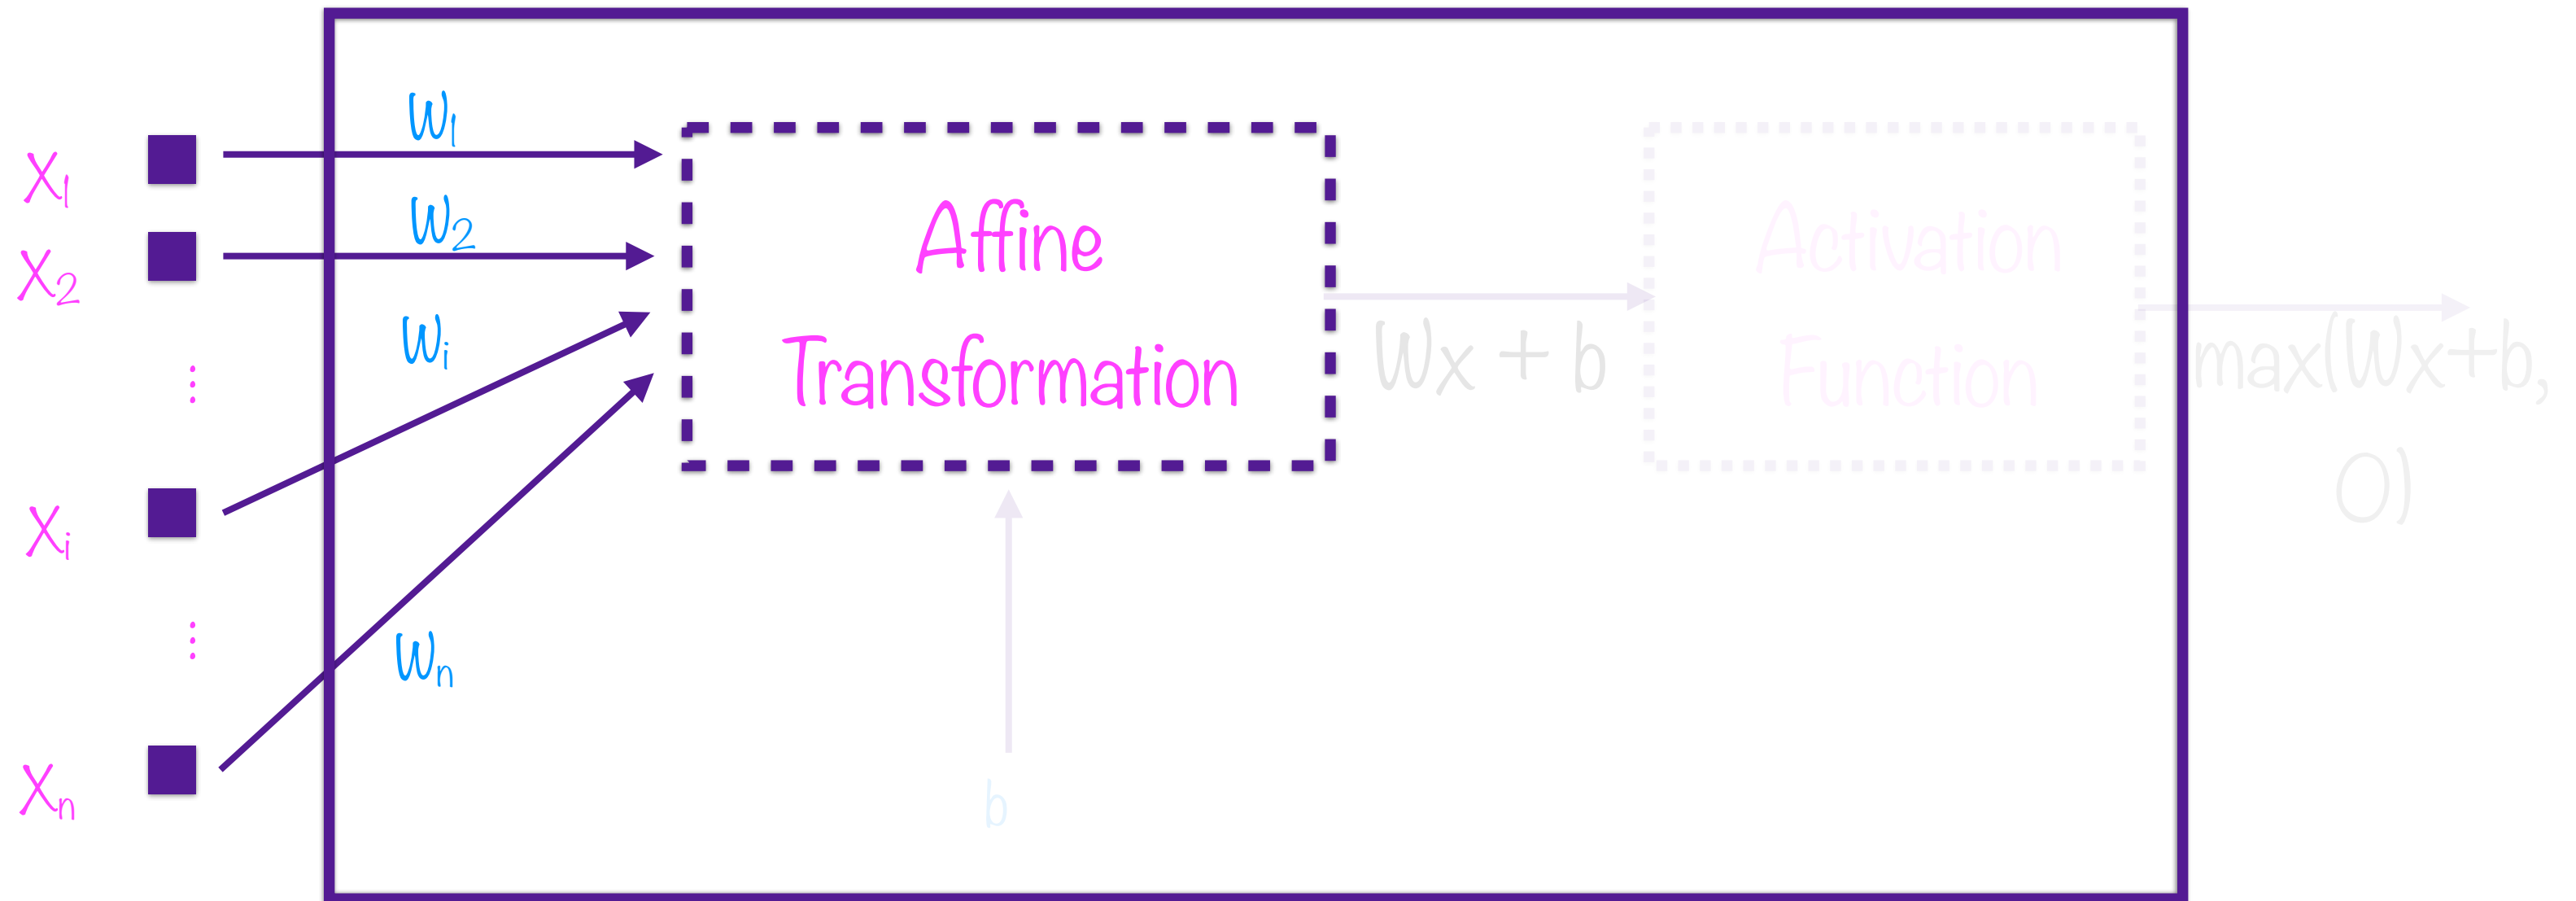
The affine transformation alone can **only** learn **linear** relationships between the inputs and the output

Operation of a Single Neuron



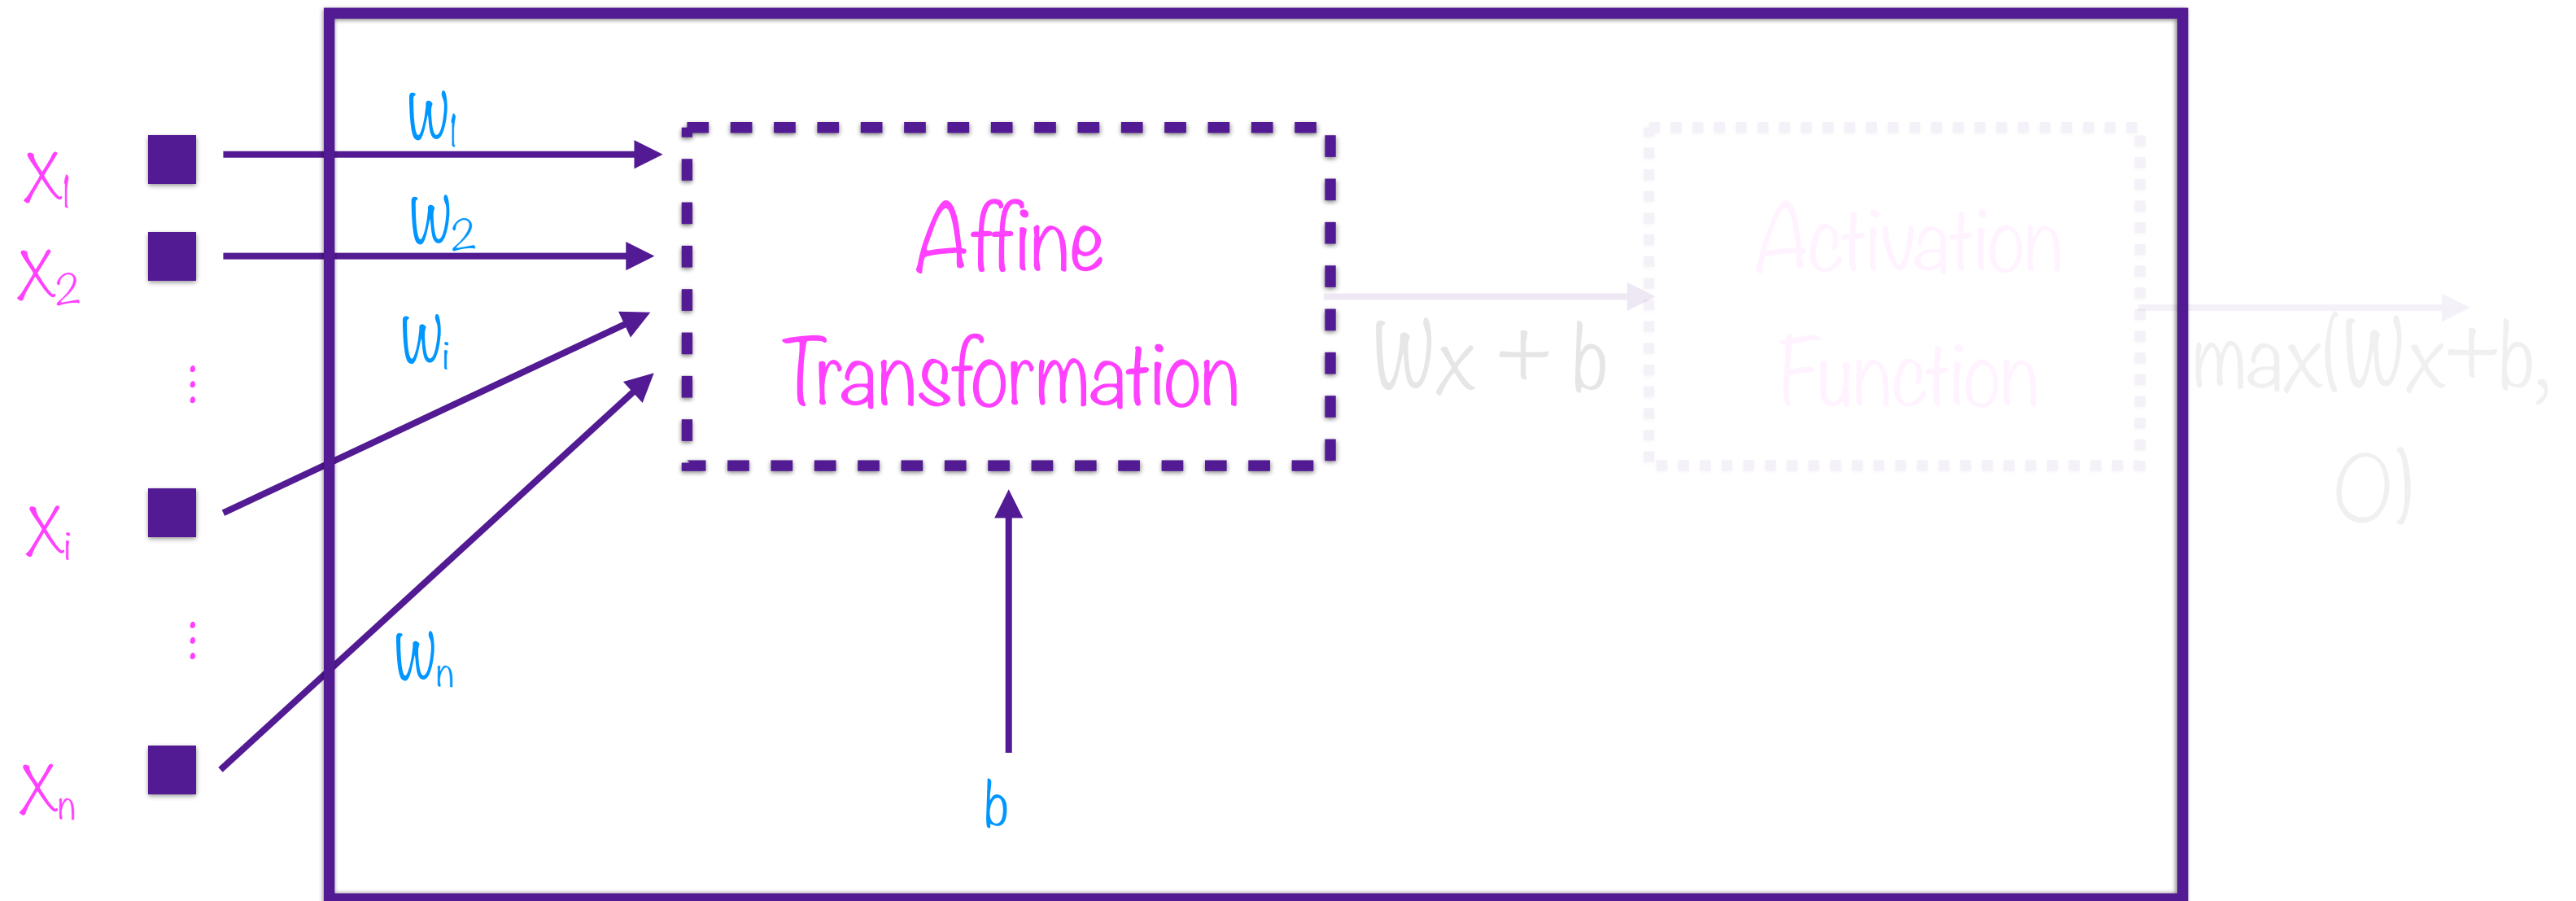
The **combination** of the affine transformation and the activation function can **learn any arbitrary relationship**

Operation of a Single Neuron




The values $W_1, W_2 \dots W_n$ are called the weights

Operation of a Single Neuron



The value b is called the bias

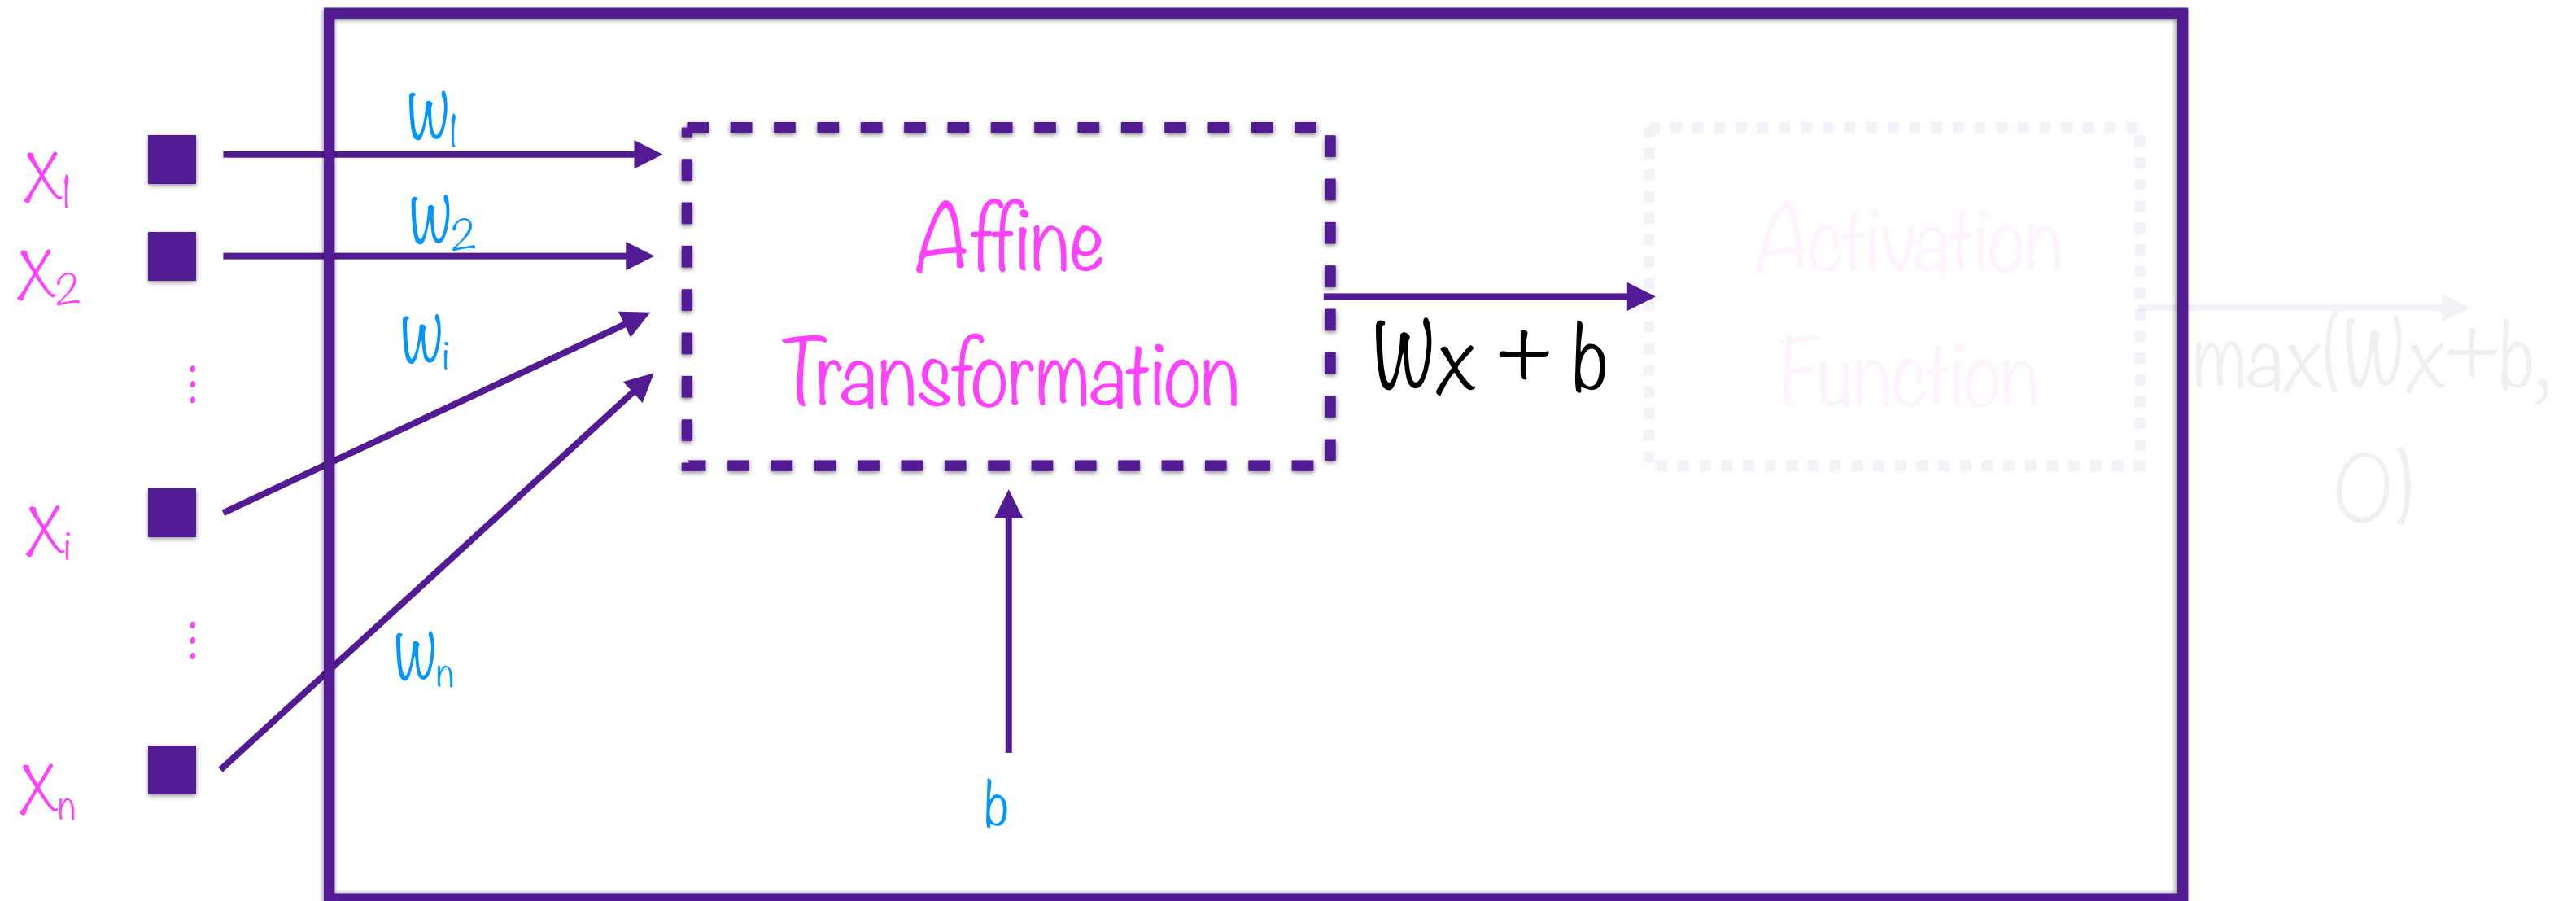
Operation of a Single Neuron



The affine transformation is just a weighted sum with a bias added:

$$w_1x_1 + w_2x_2 + \dots + w_nx_n + b$$

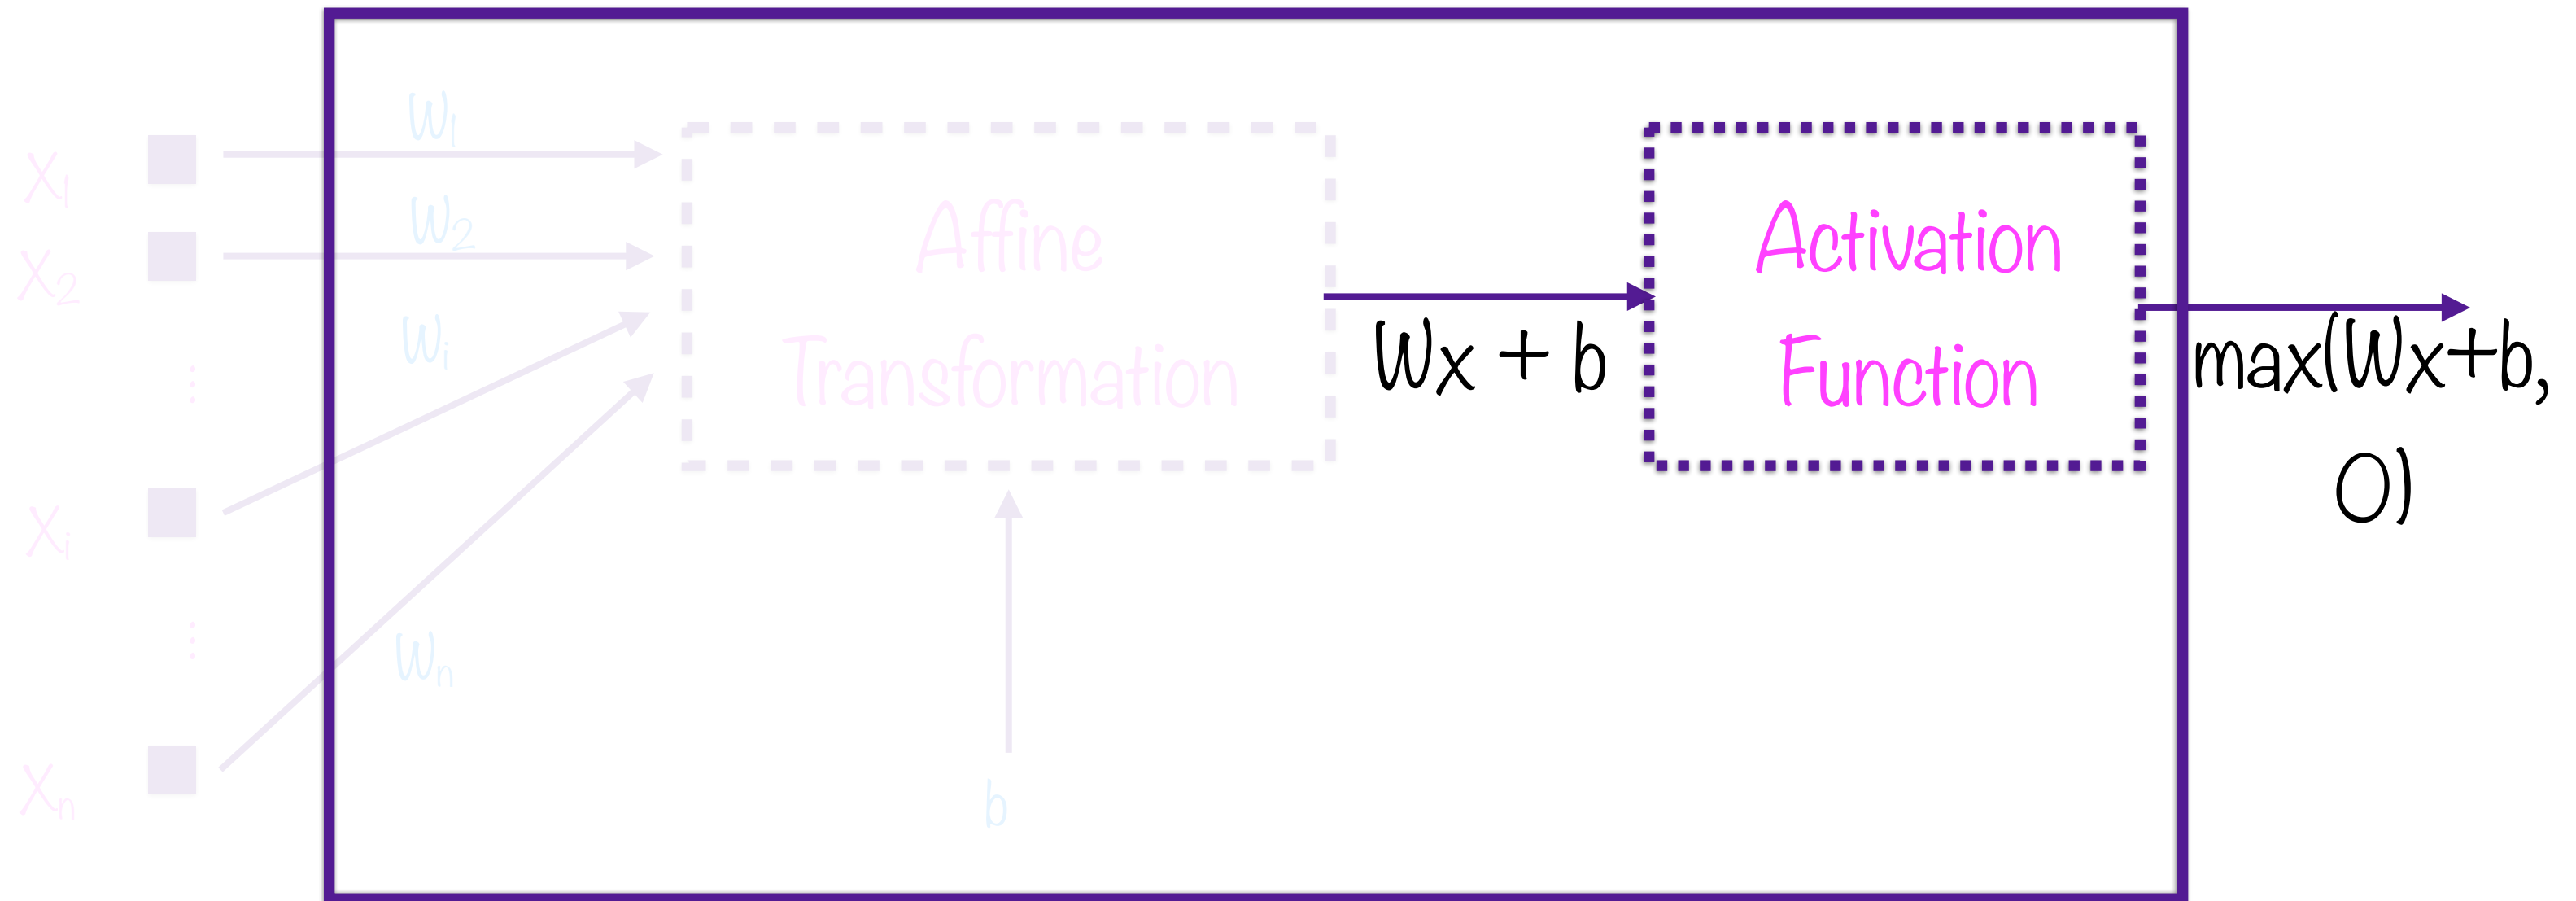
Operation of a Single Neuron



Where do the values of W and b come from?

**The weights and biases of a neuron are determined by
the training process**


Operation of a Single Neuron



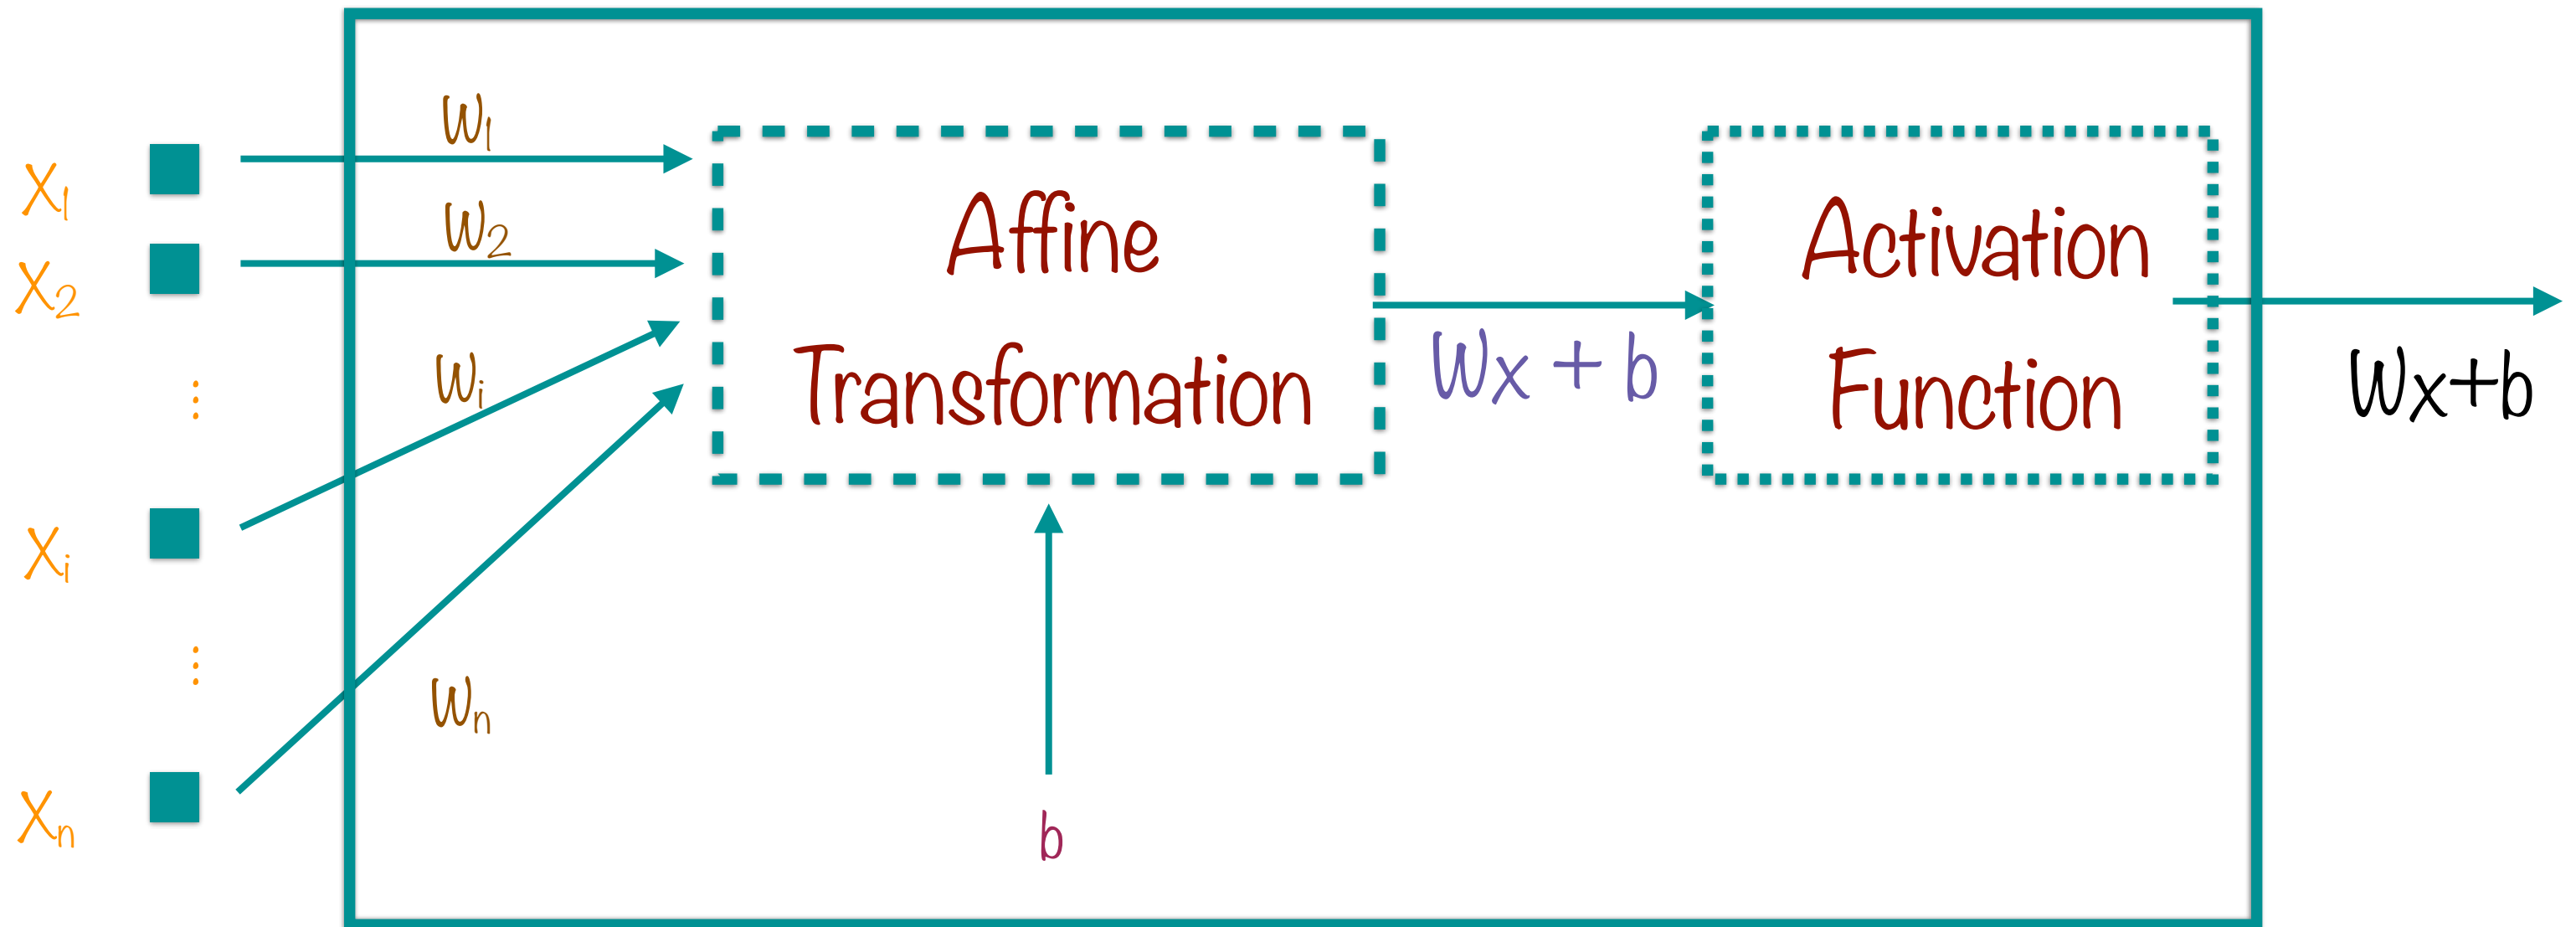
The combination of the affine transformation and the activation function can learn any arbitrary relationship

Activation Functions for Non-Linear Relationships

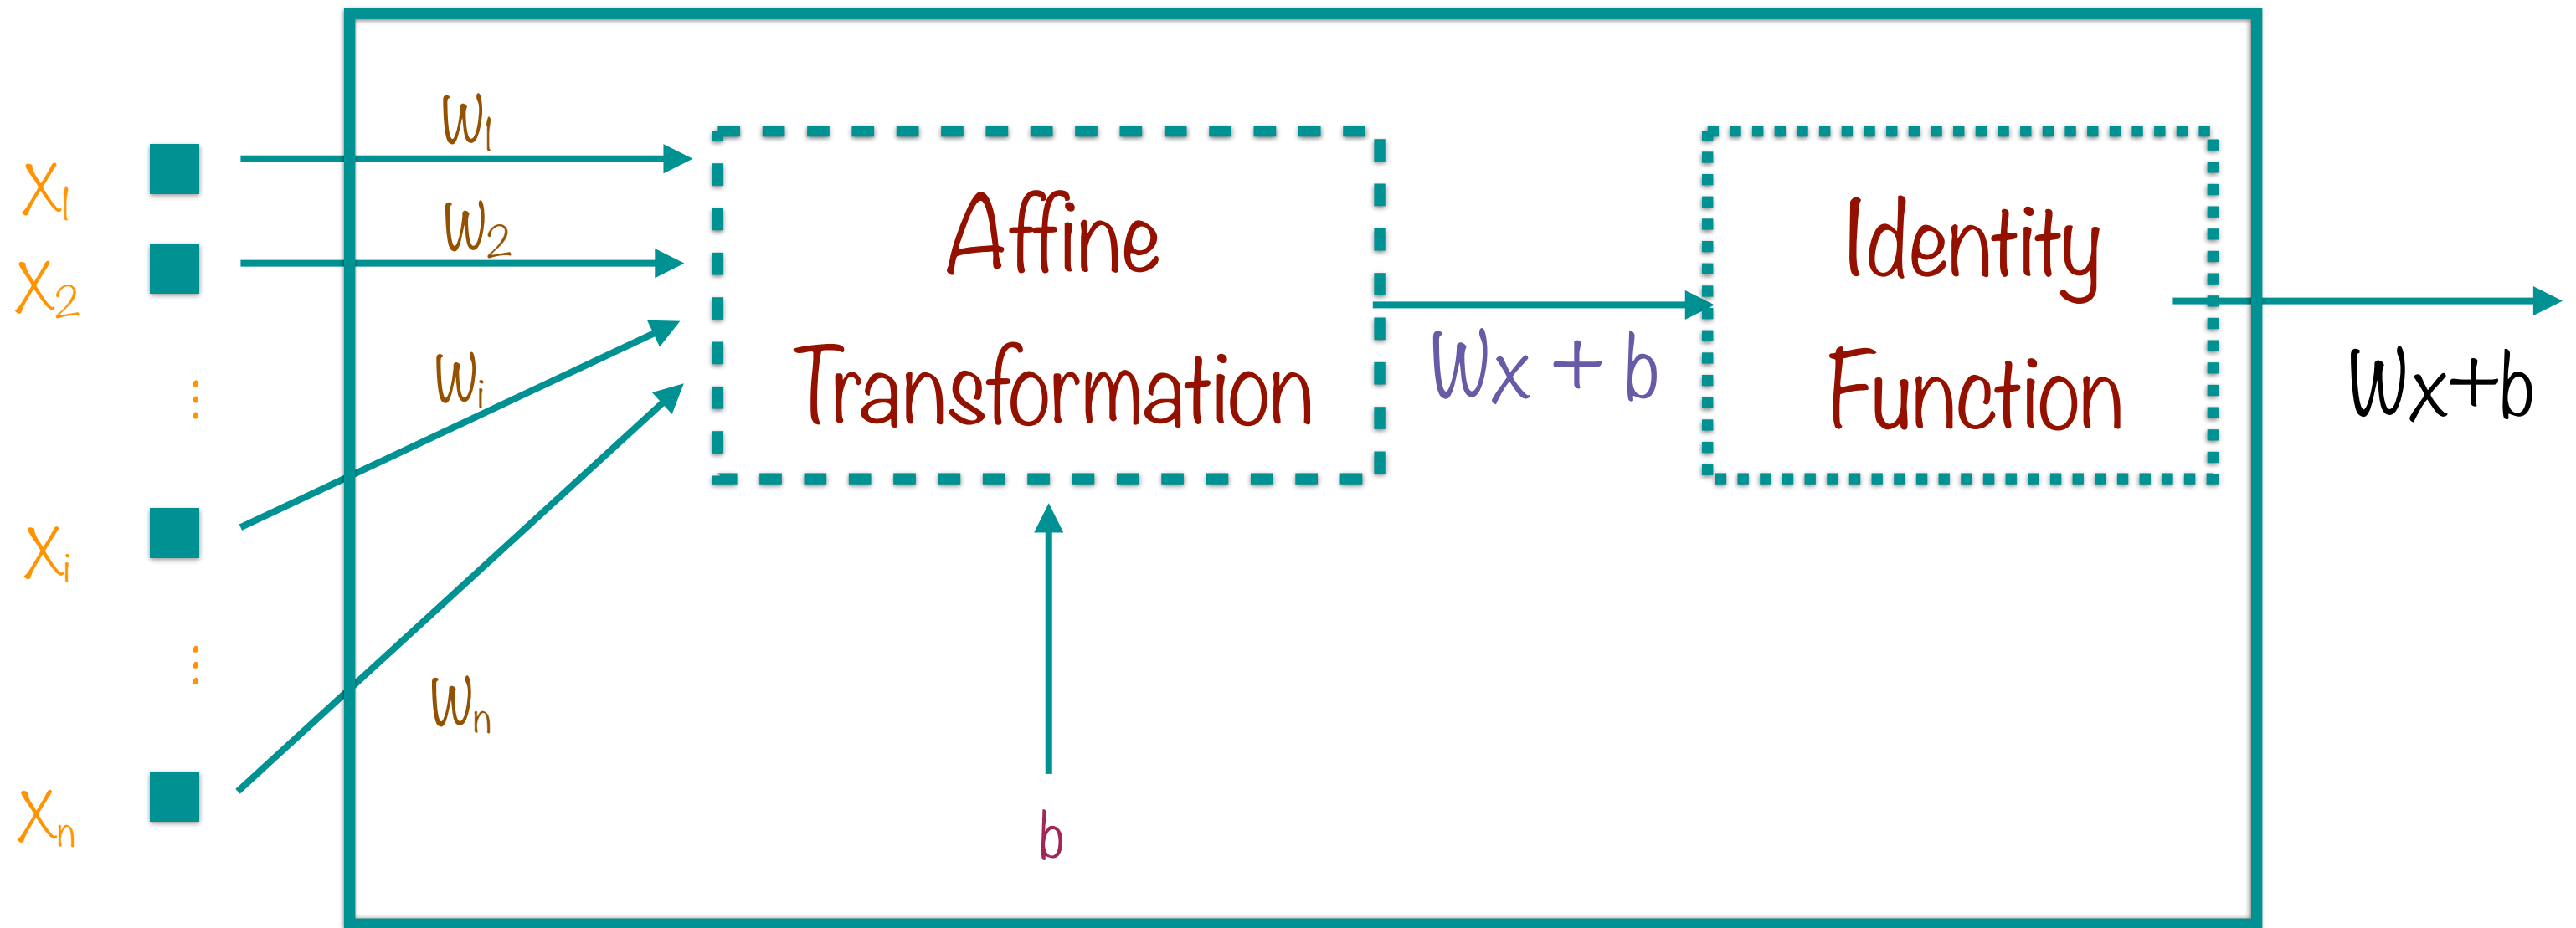
Regression: The Simplest Neural Network



Regression: The Simplest Neural Network




Regression: The Simplest Neural Network



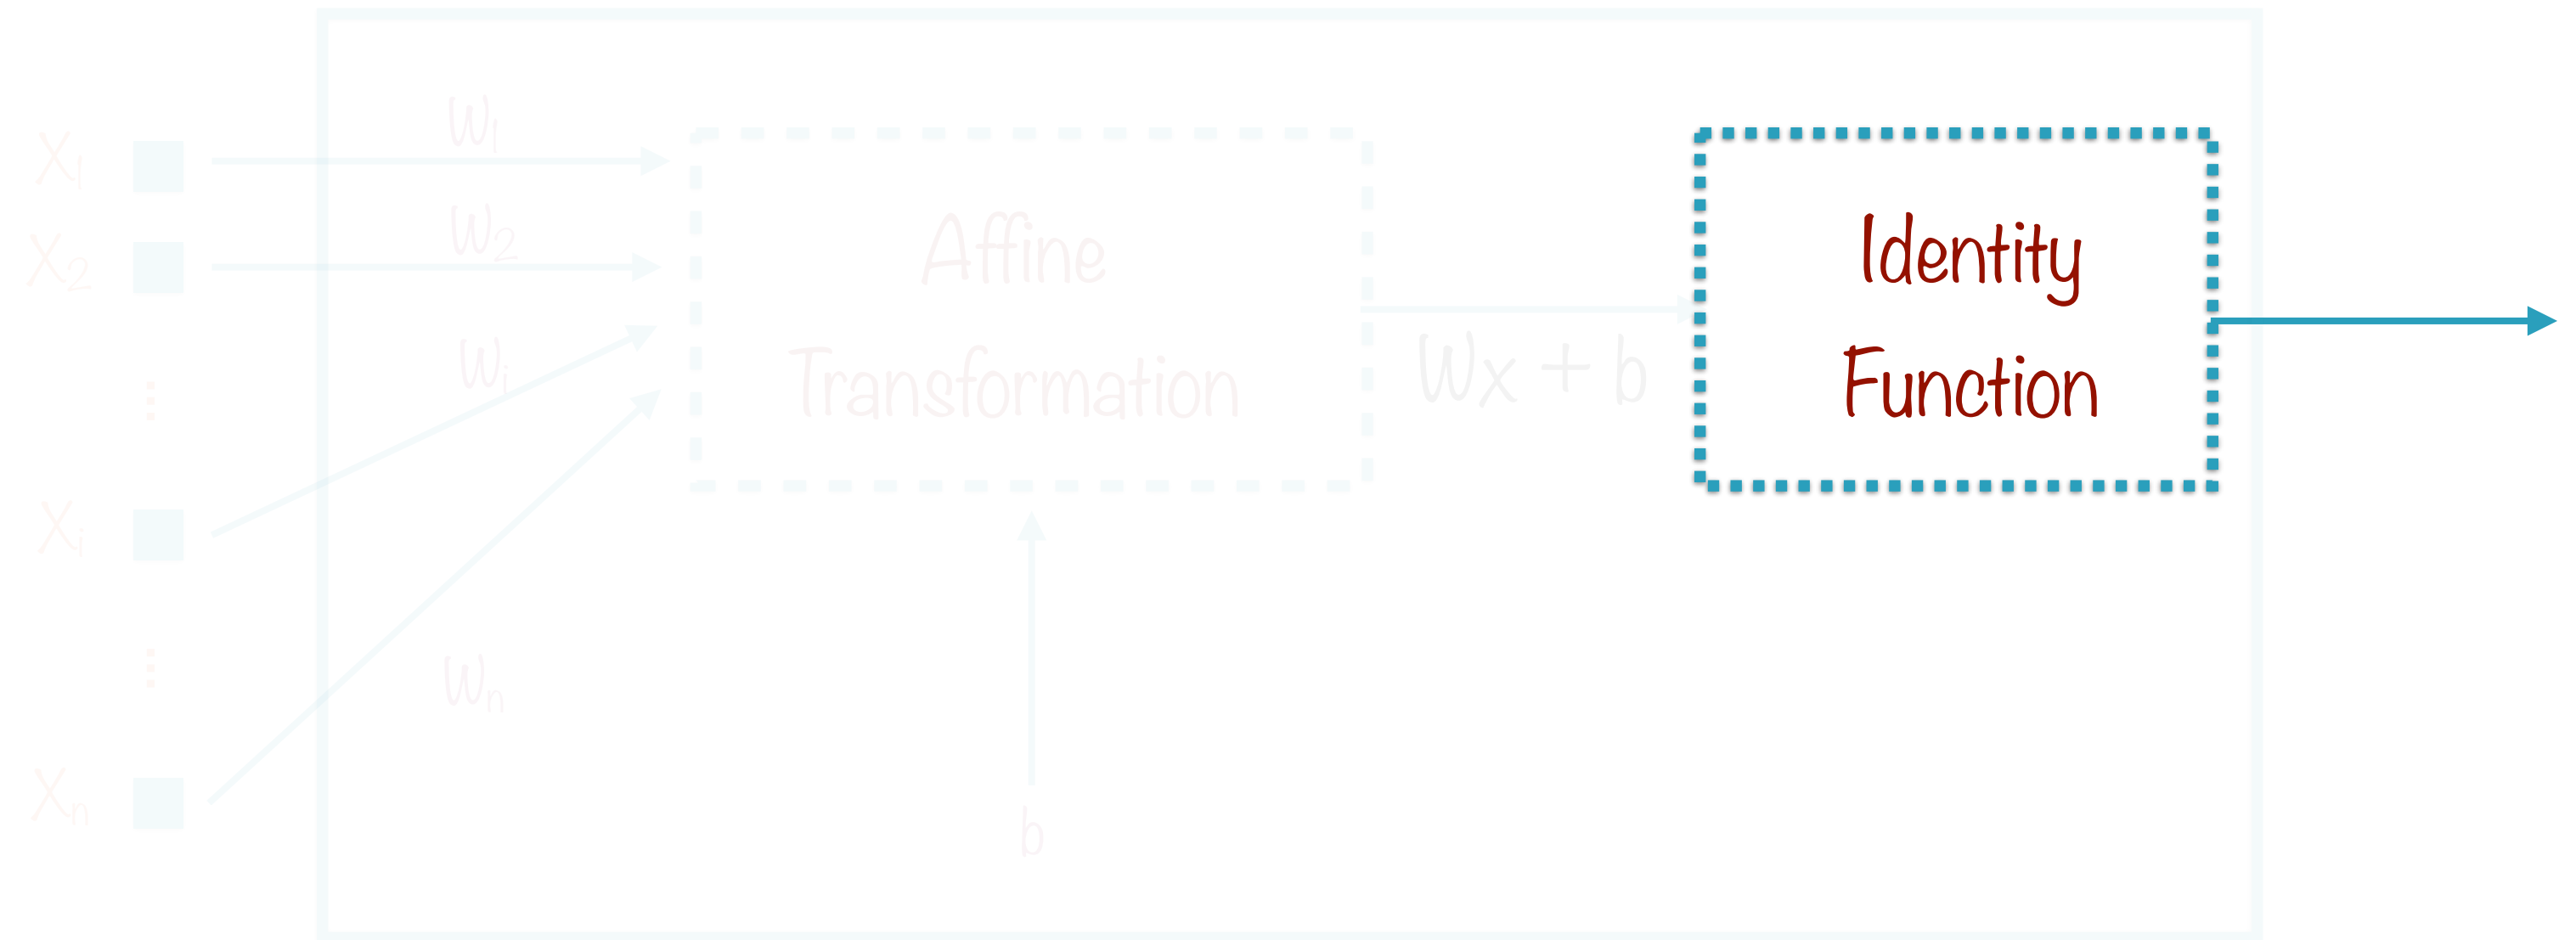
activation=None, simply passes the output of the linear affine transformation to the output of the neuron

Regression: The Simplest Neural Network

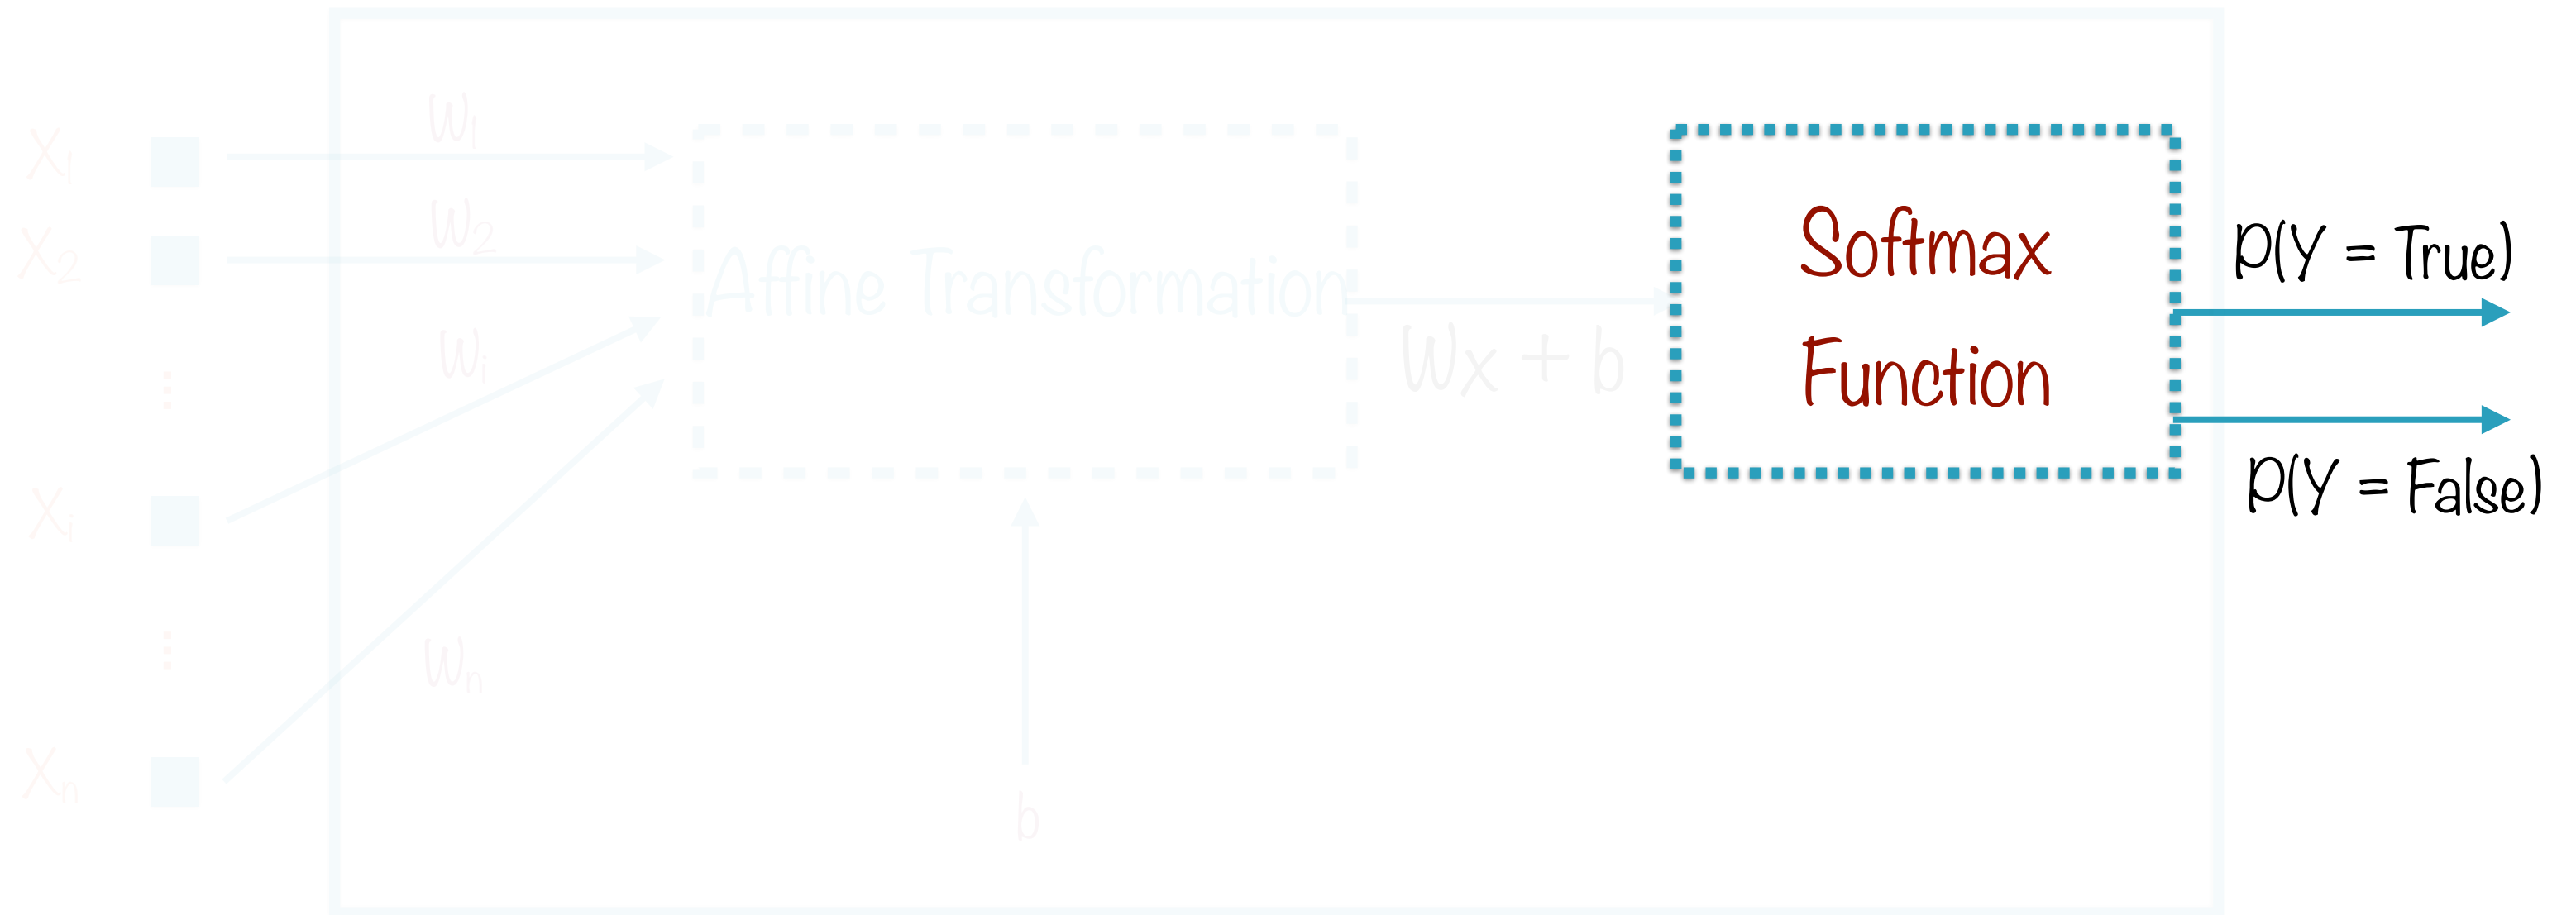


Also called a **linear neuron**

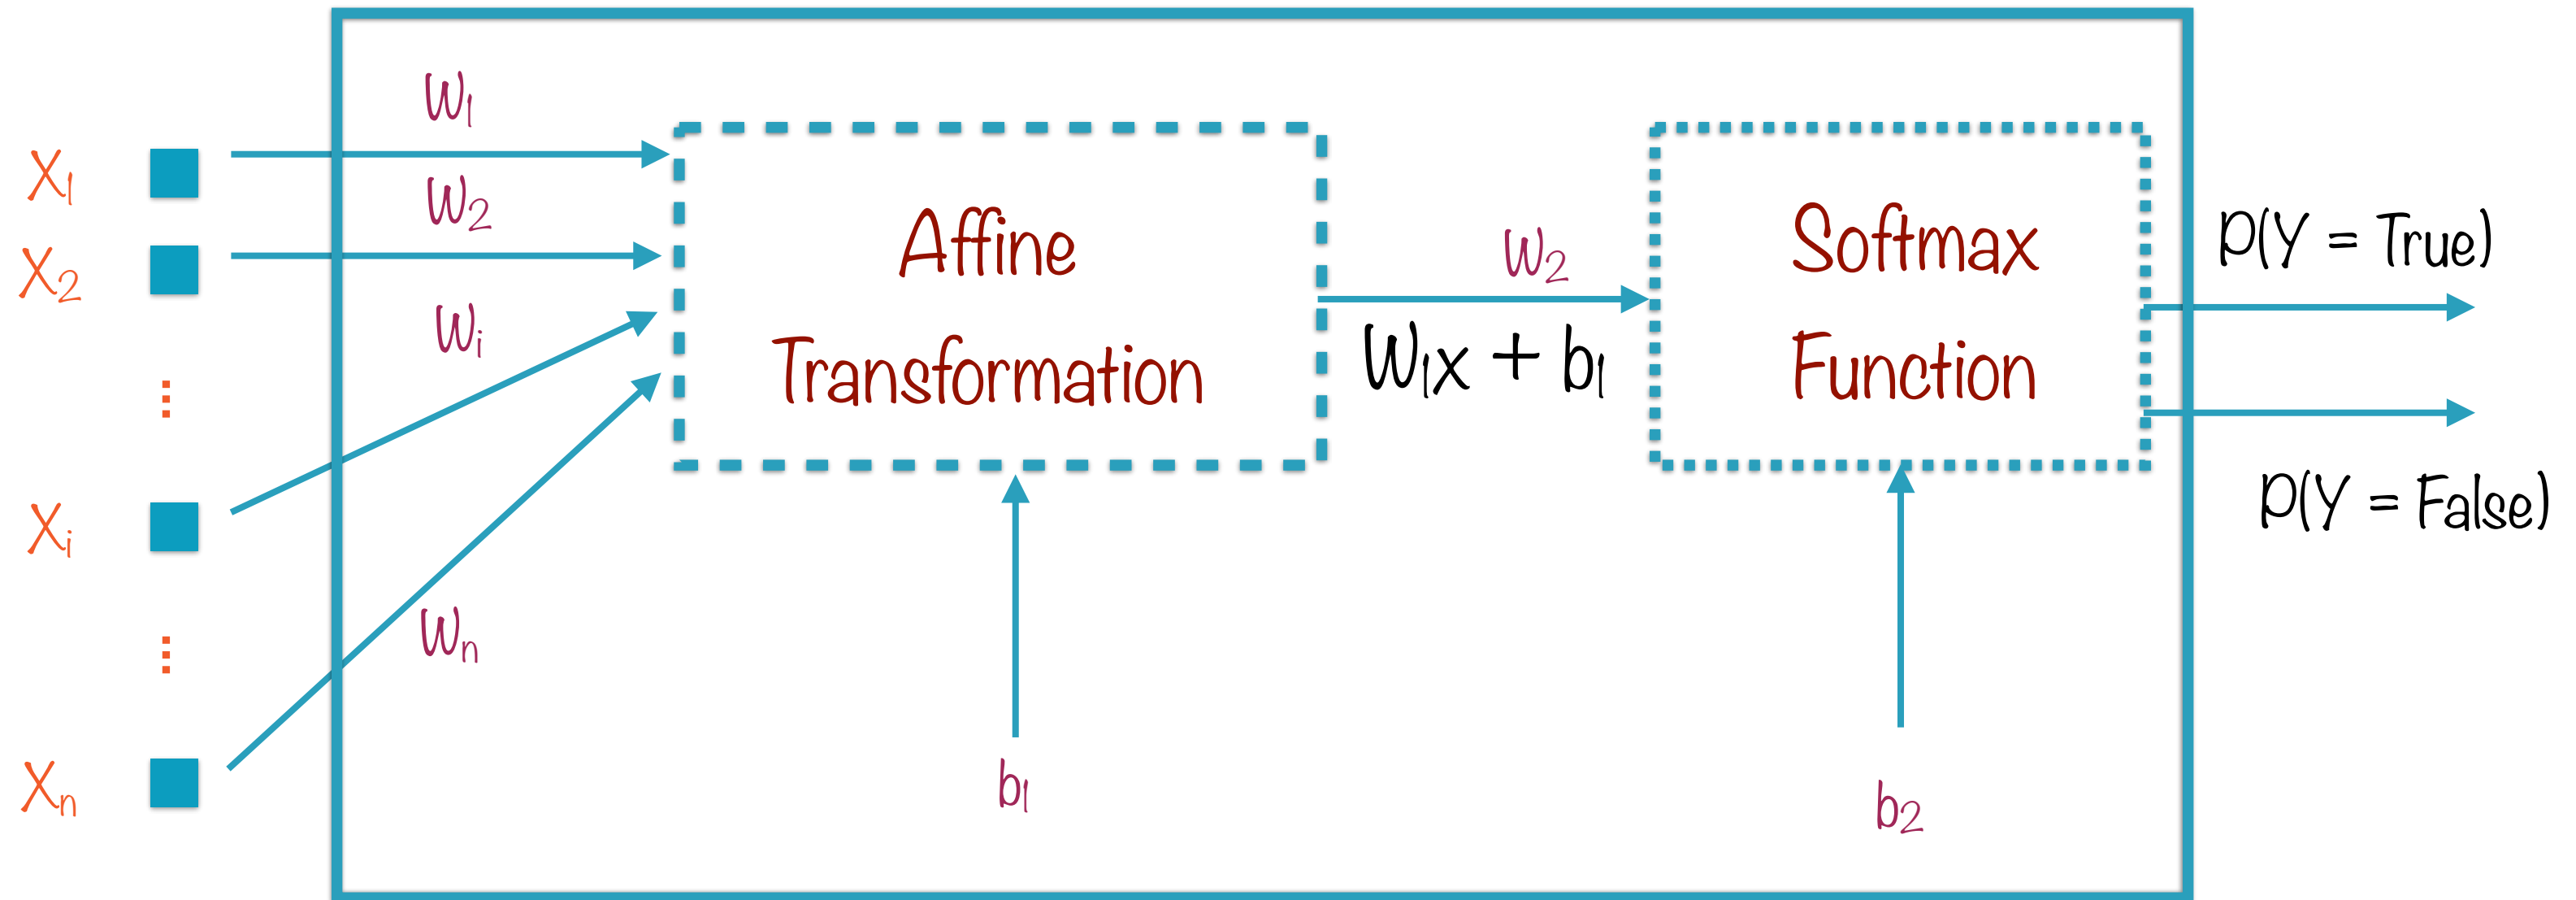
Linear Regression with One Neuron



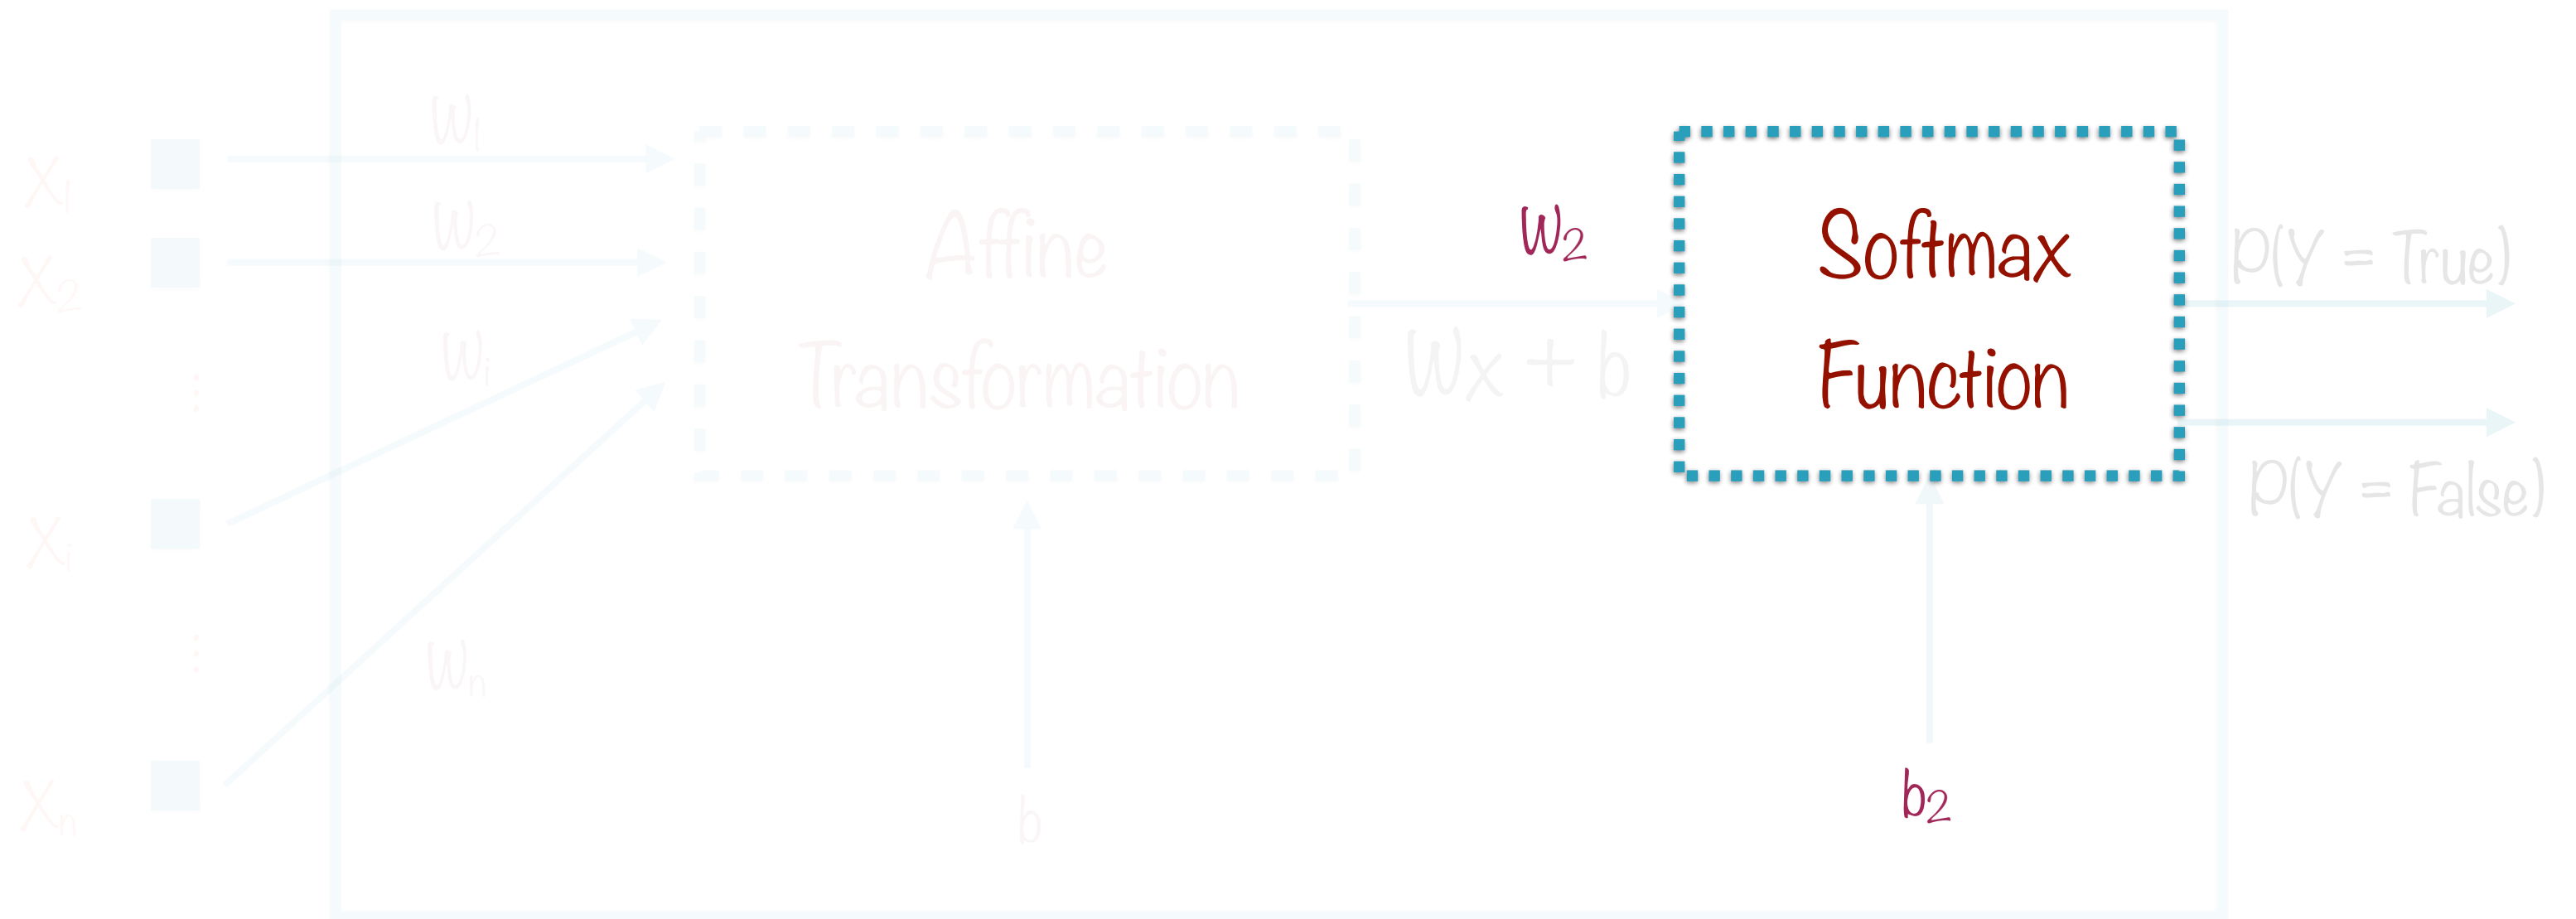
Logistic Regression with One Neuron



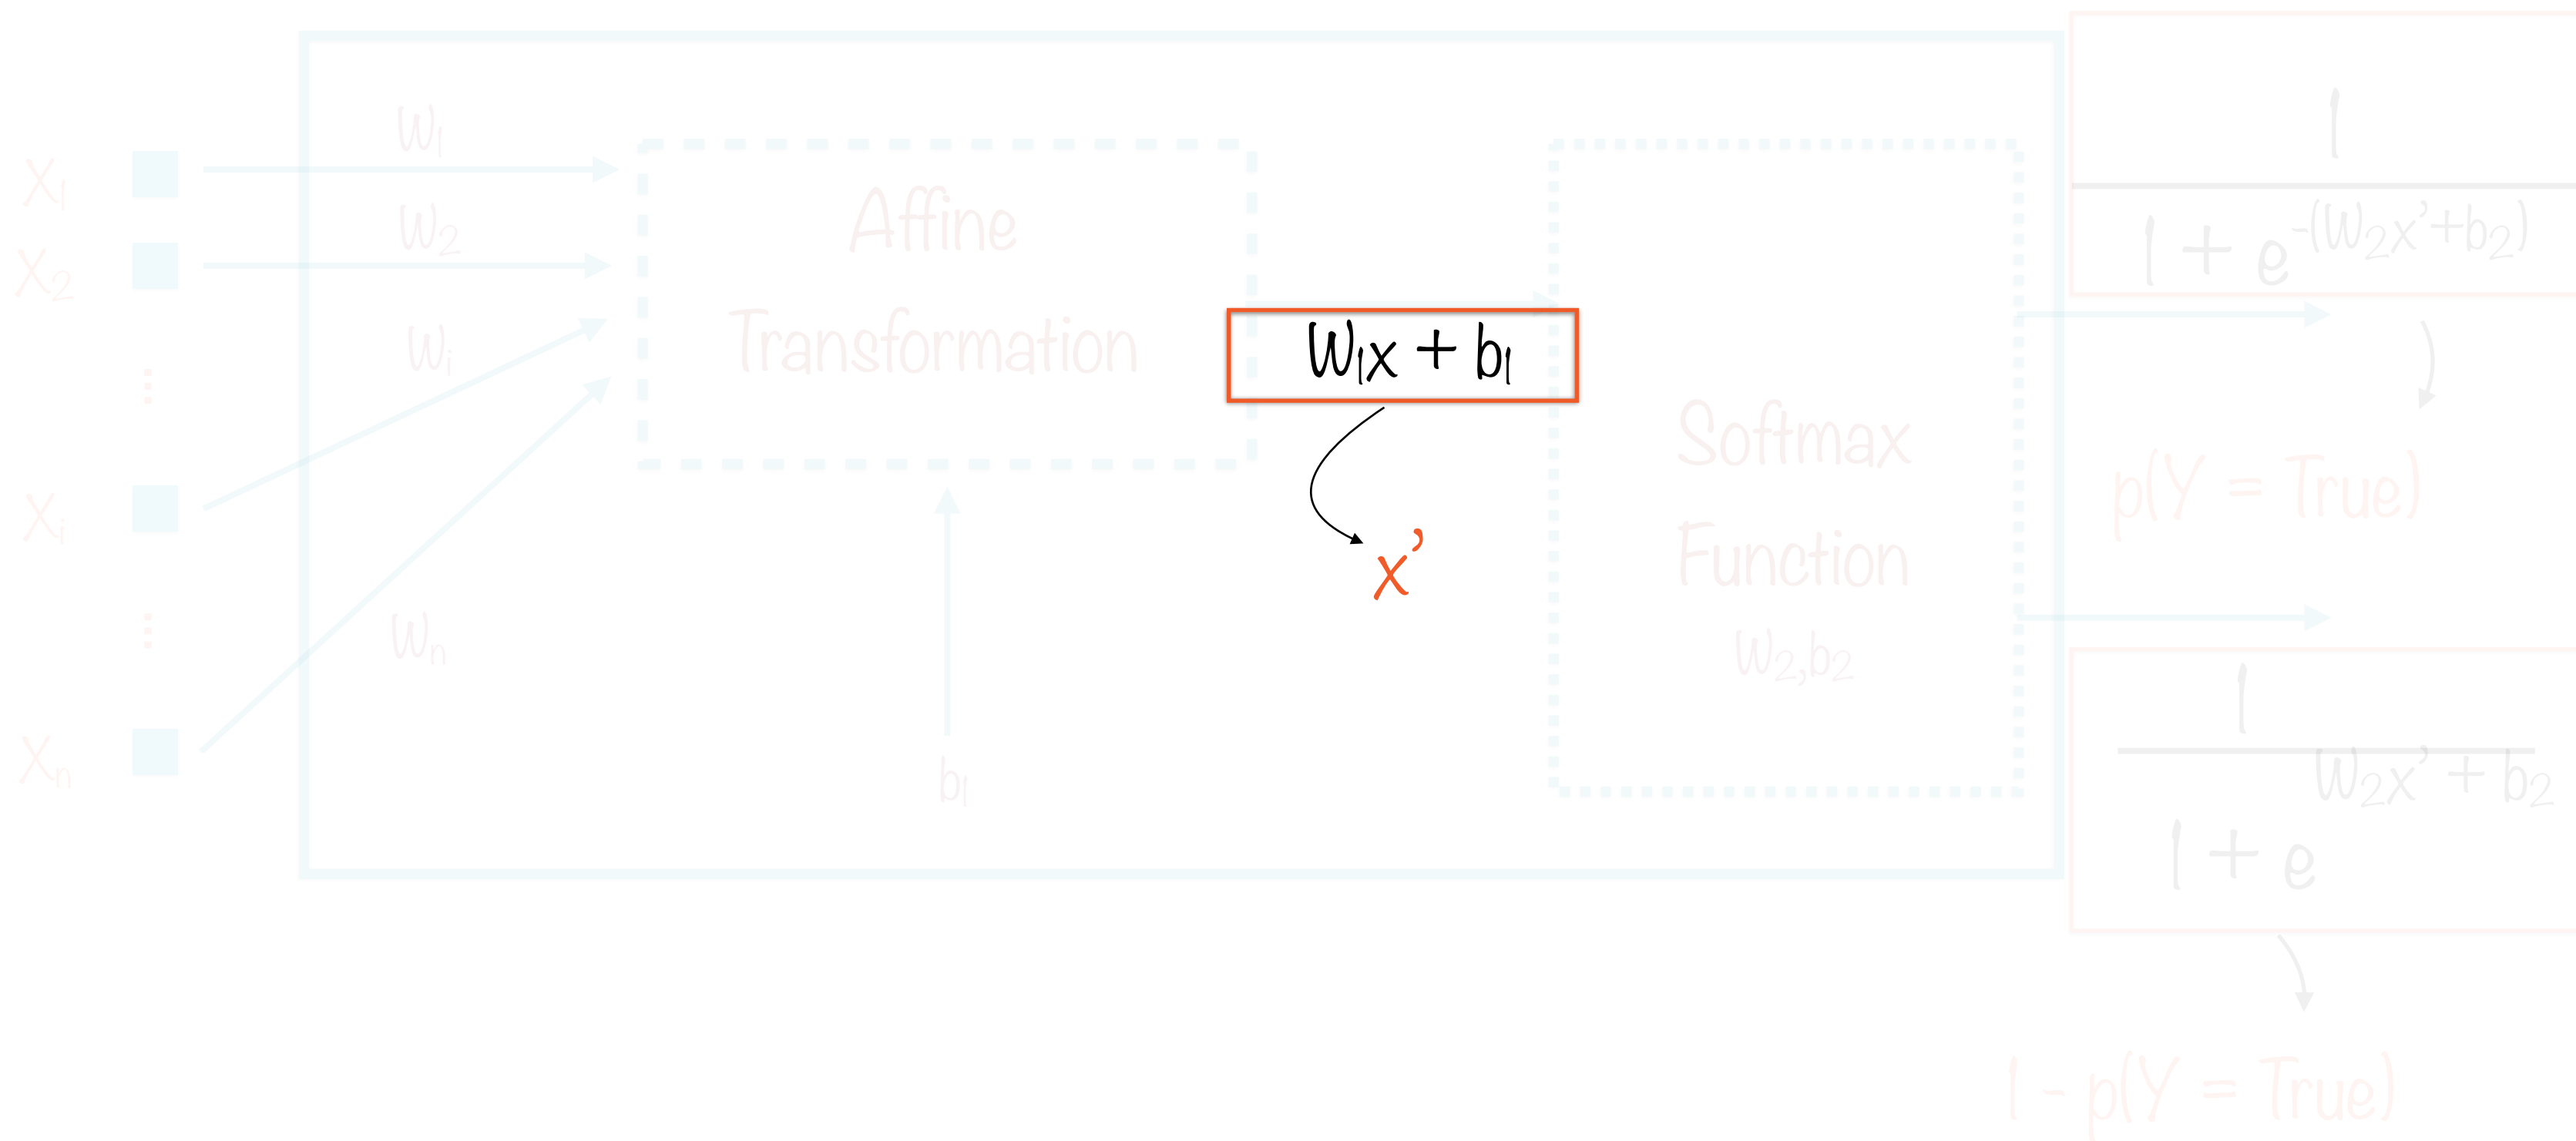
Logistic Regression with One Neuron



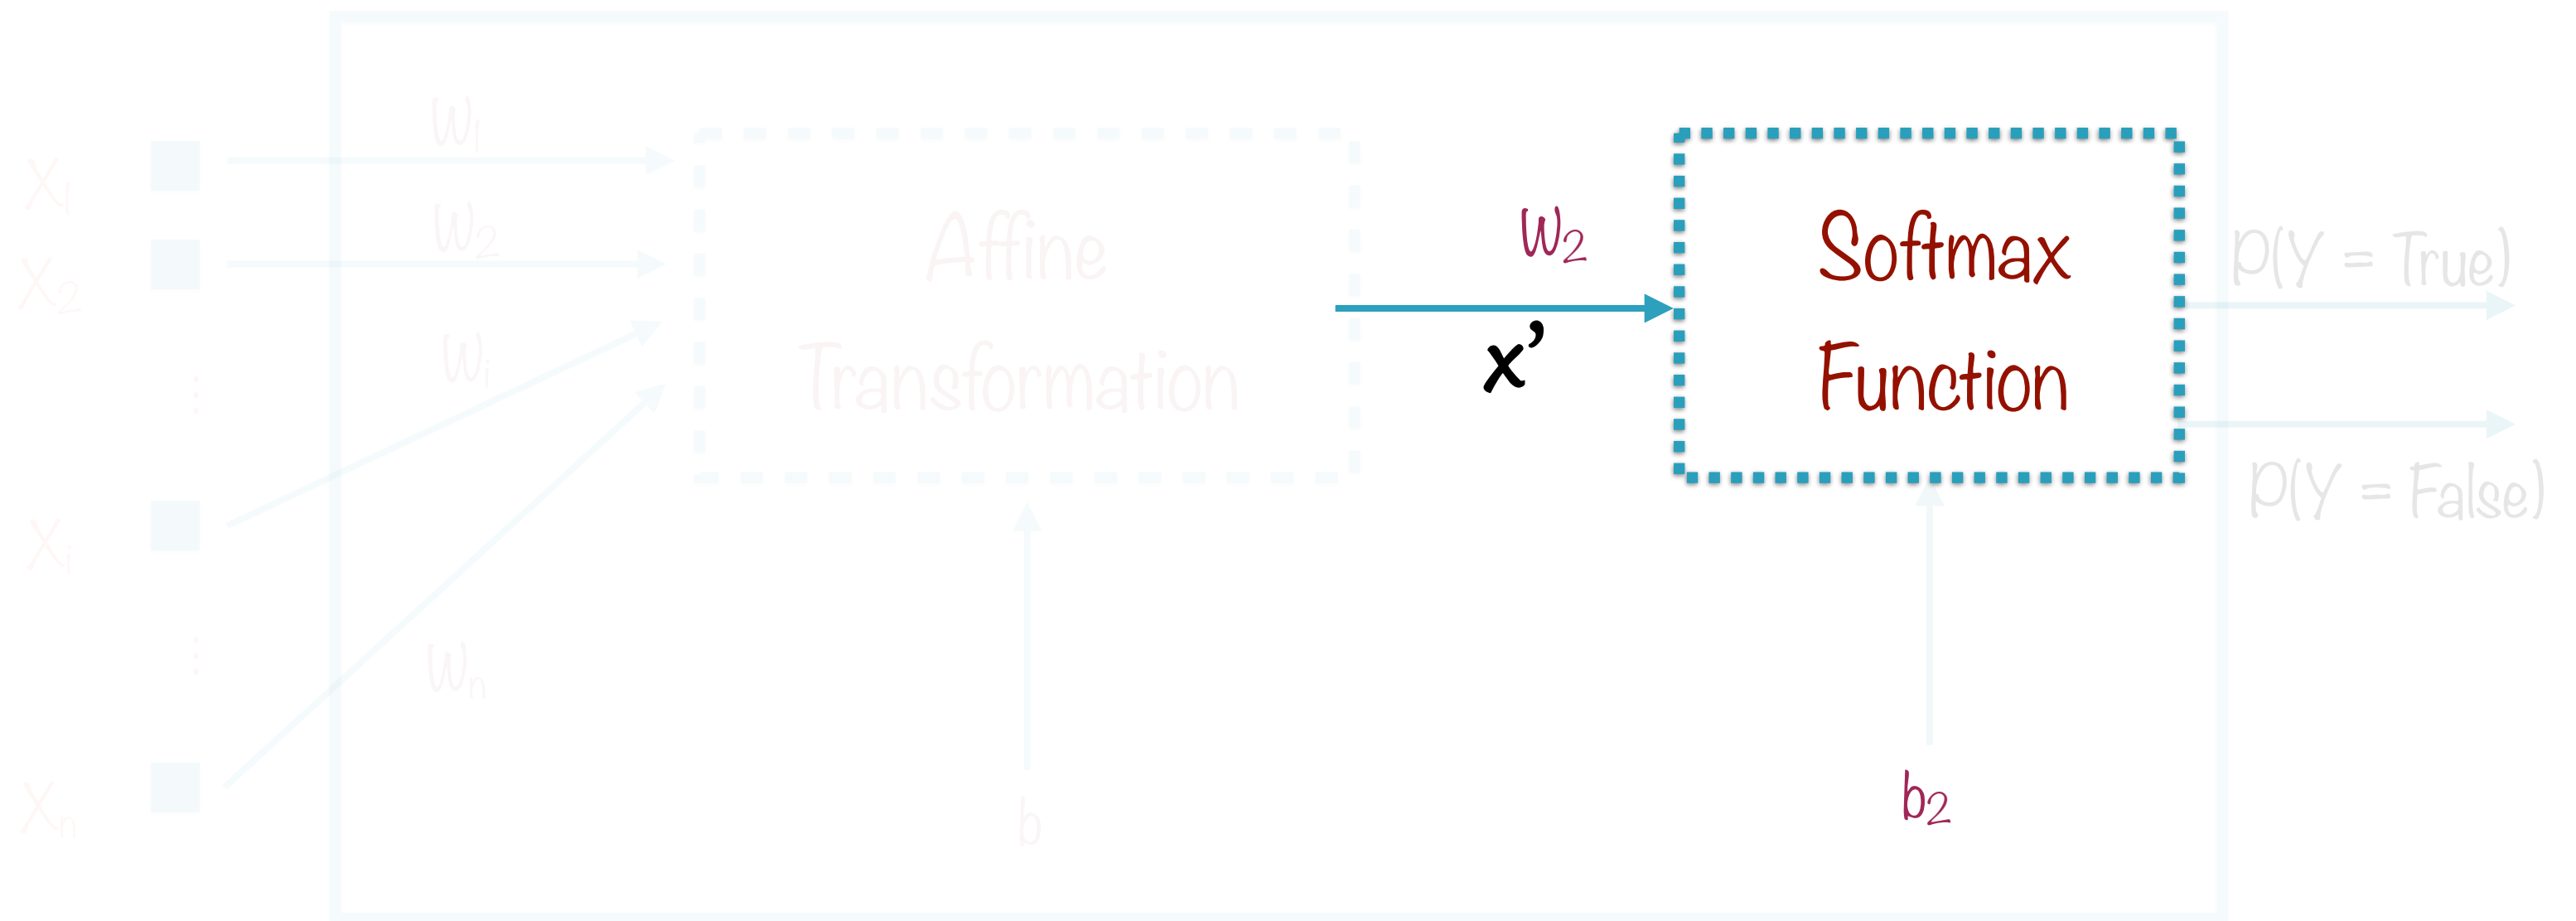
Logistic Regression with One Neuron



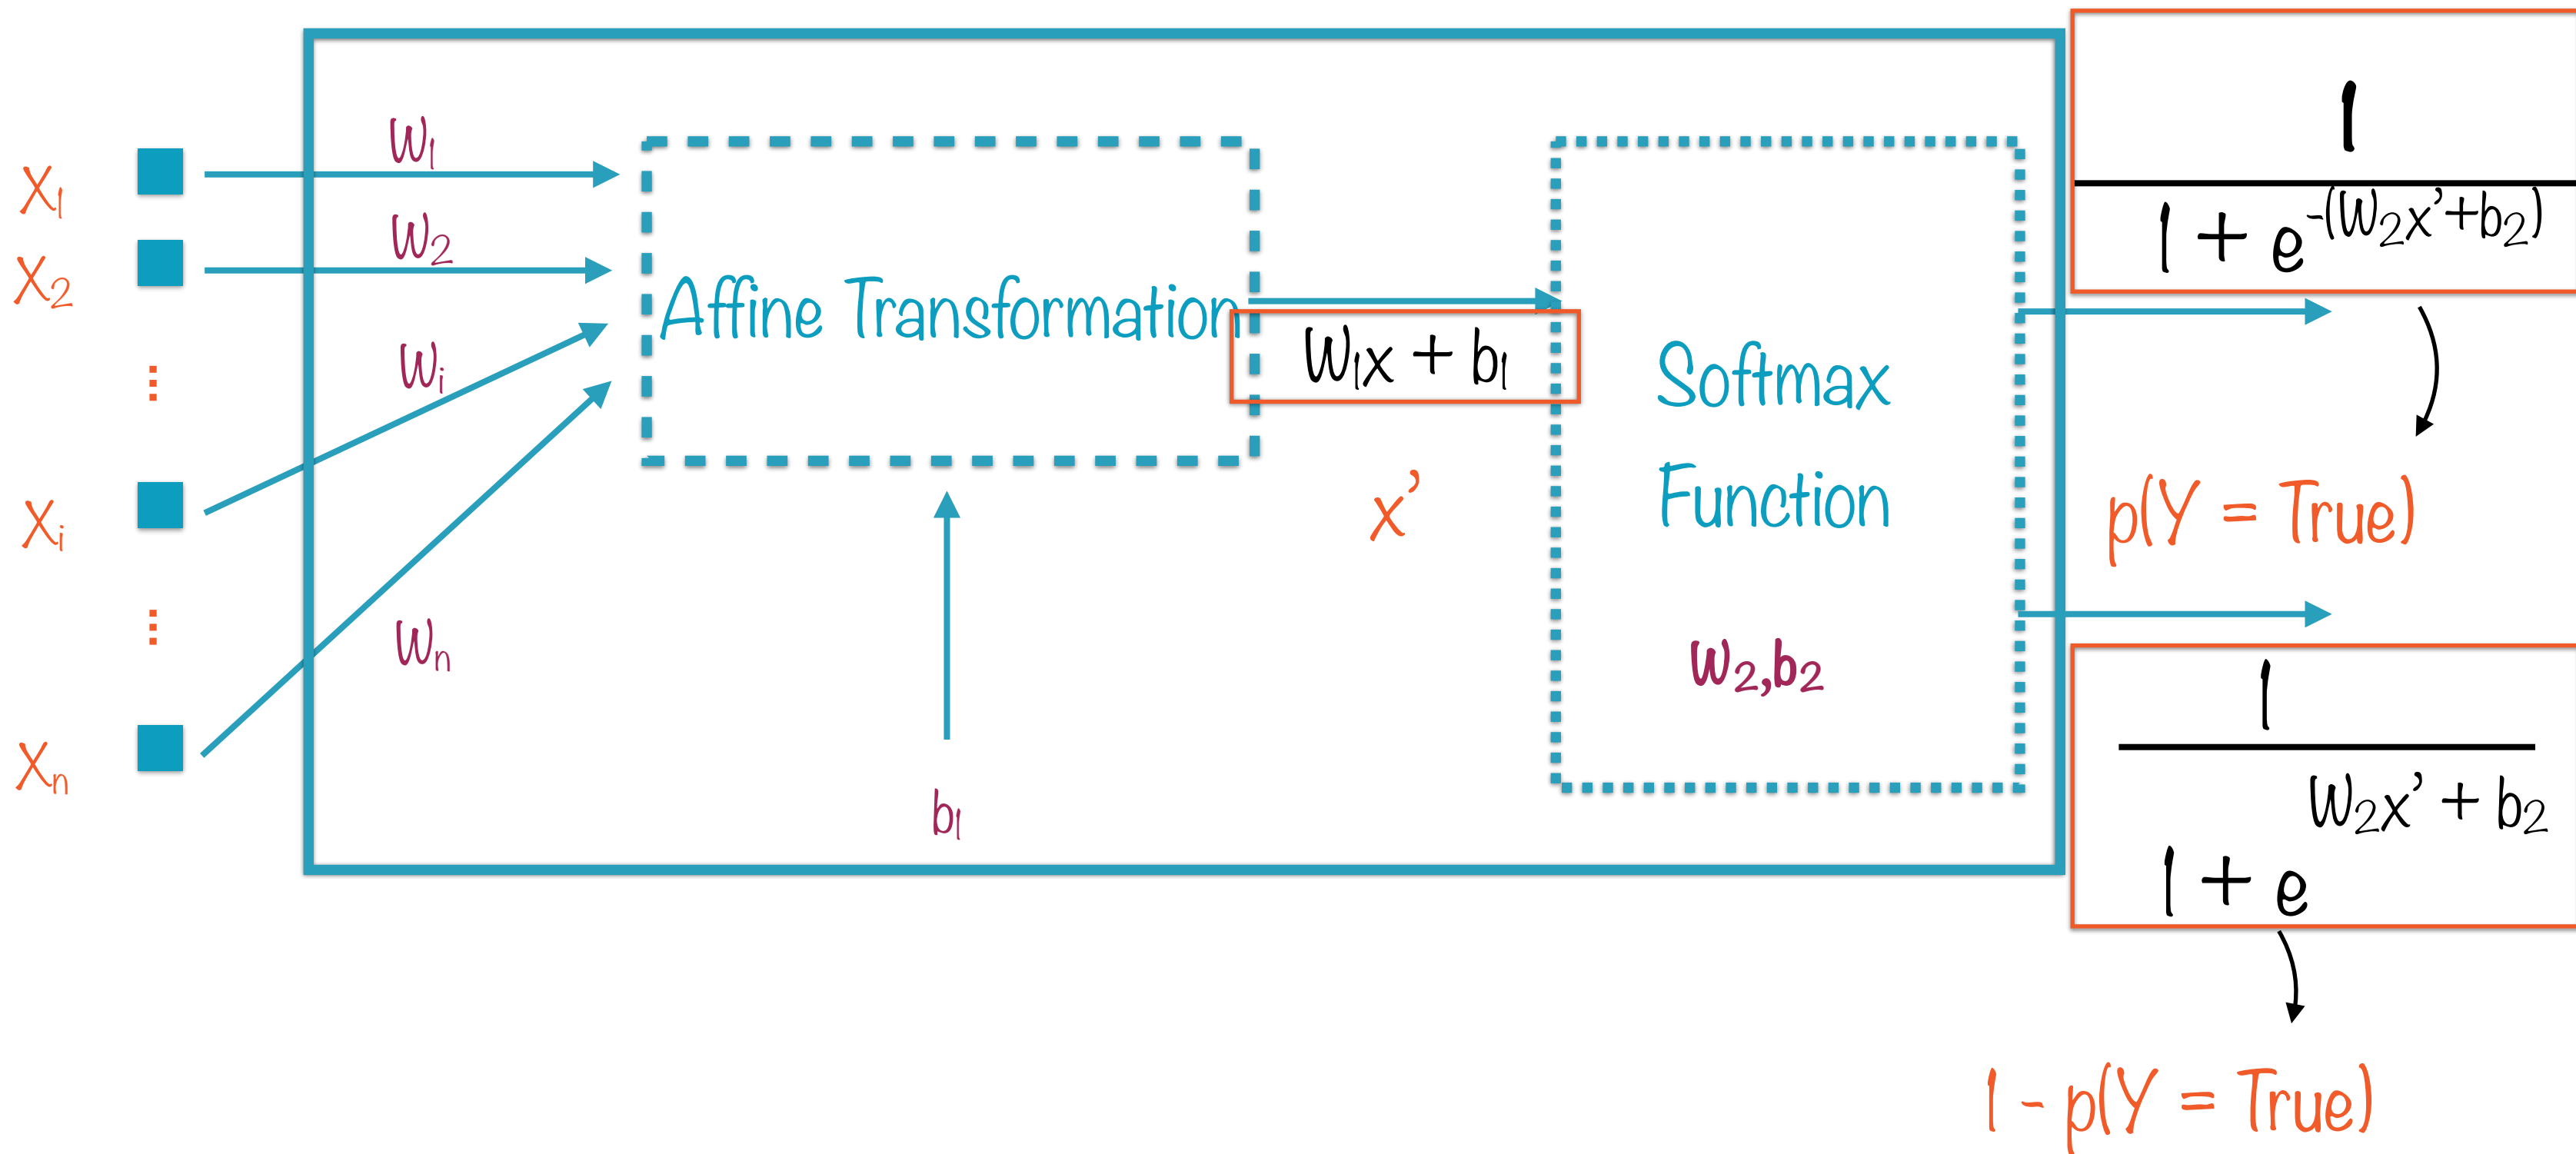
Logistic Regression with One Neuron



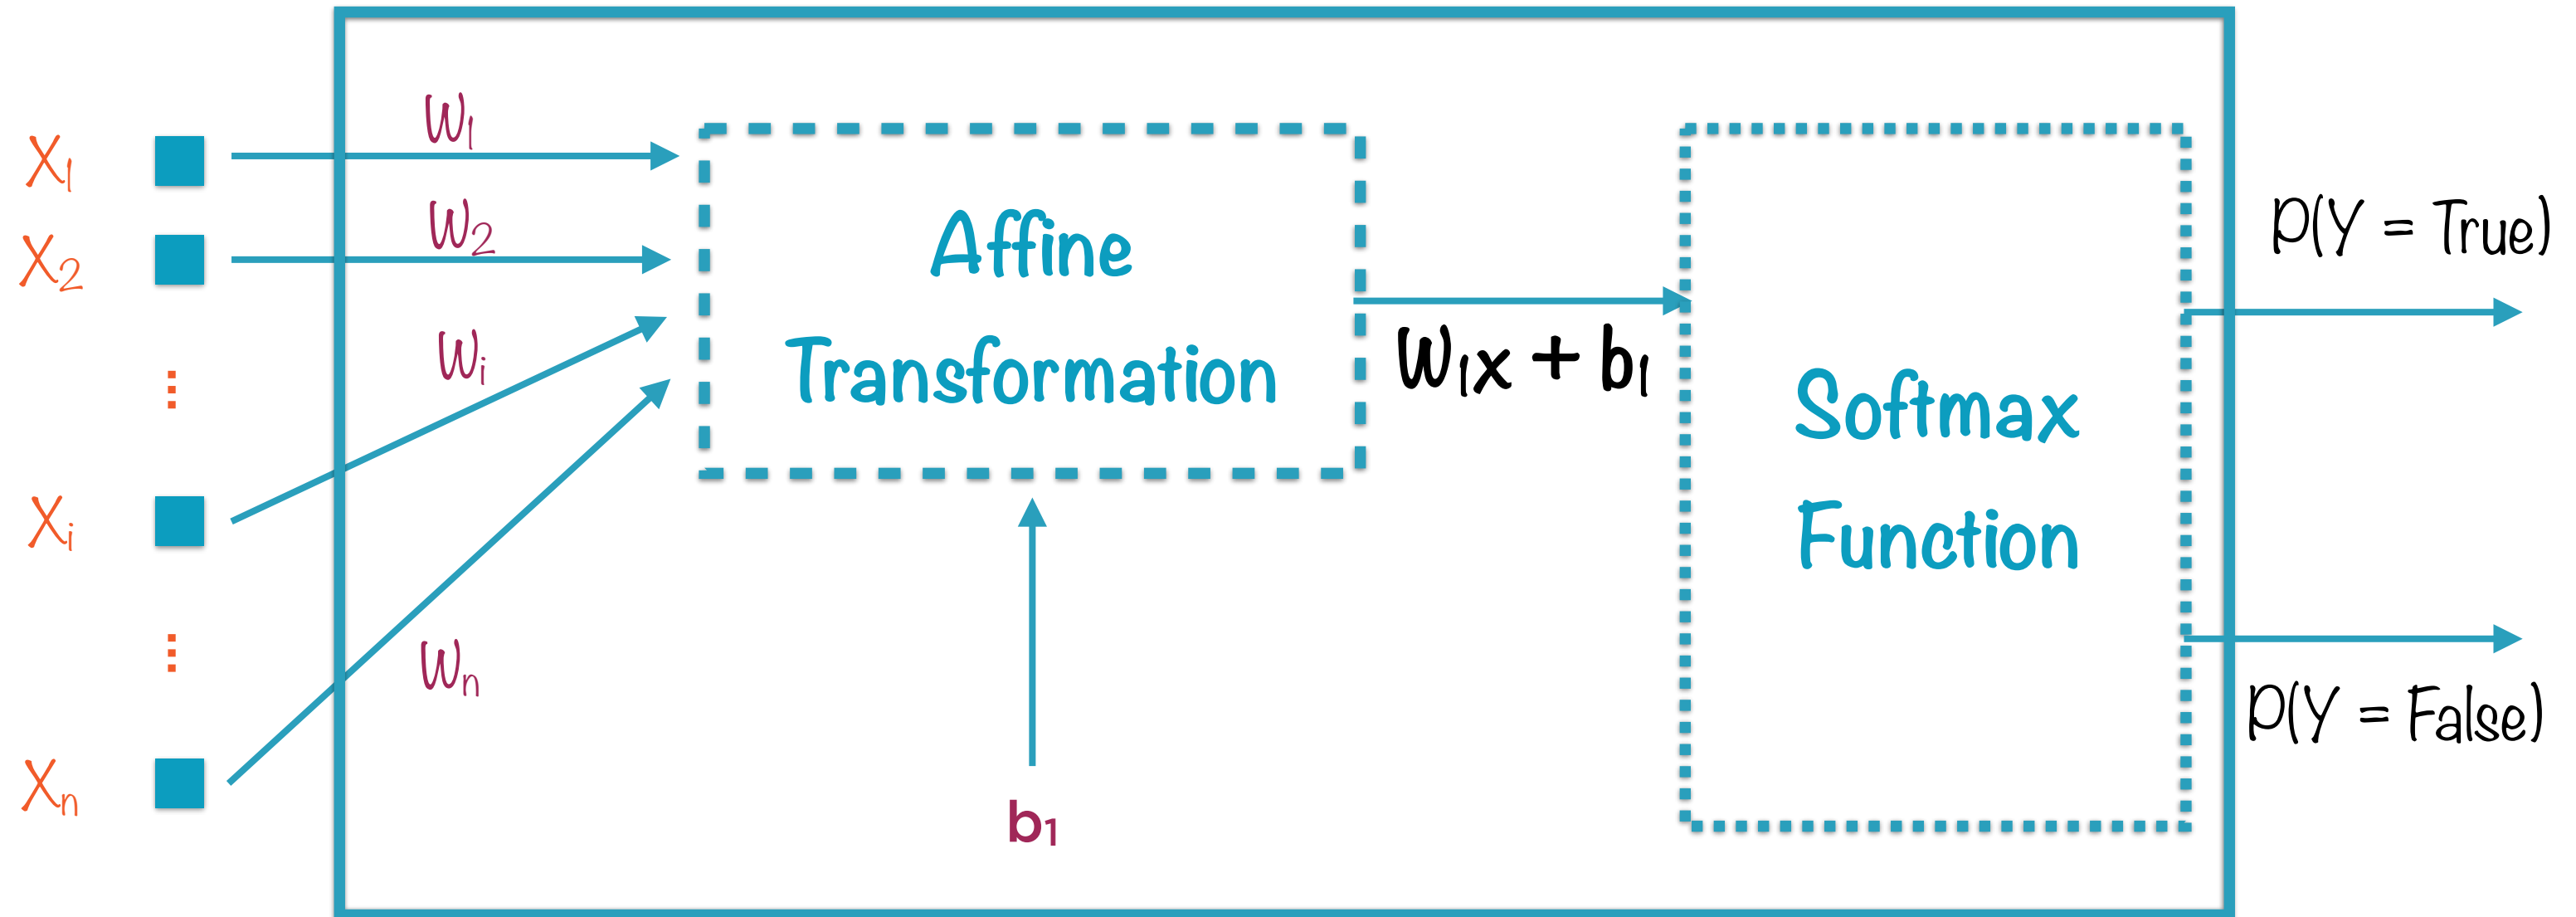
Logistic Regression with One Neuron



Logistic Regression with One Neuron



Logistic Regression with One Neuron



Logistic Regression



$p(y)$




(x_1, y_1)



(x_2, y_2)

(x_n, y_n)



Regression Curve


$p(y) =$

$$\frac{1}{1 + e^{-(A+Bx)}}$$


x




SoftMax for True/False Classification



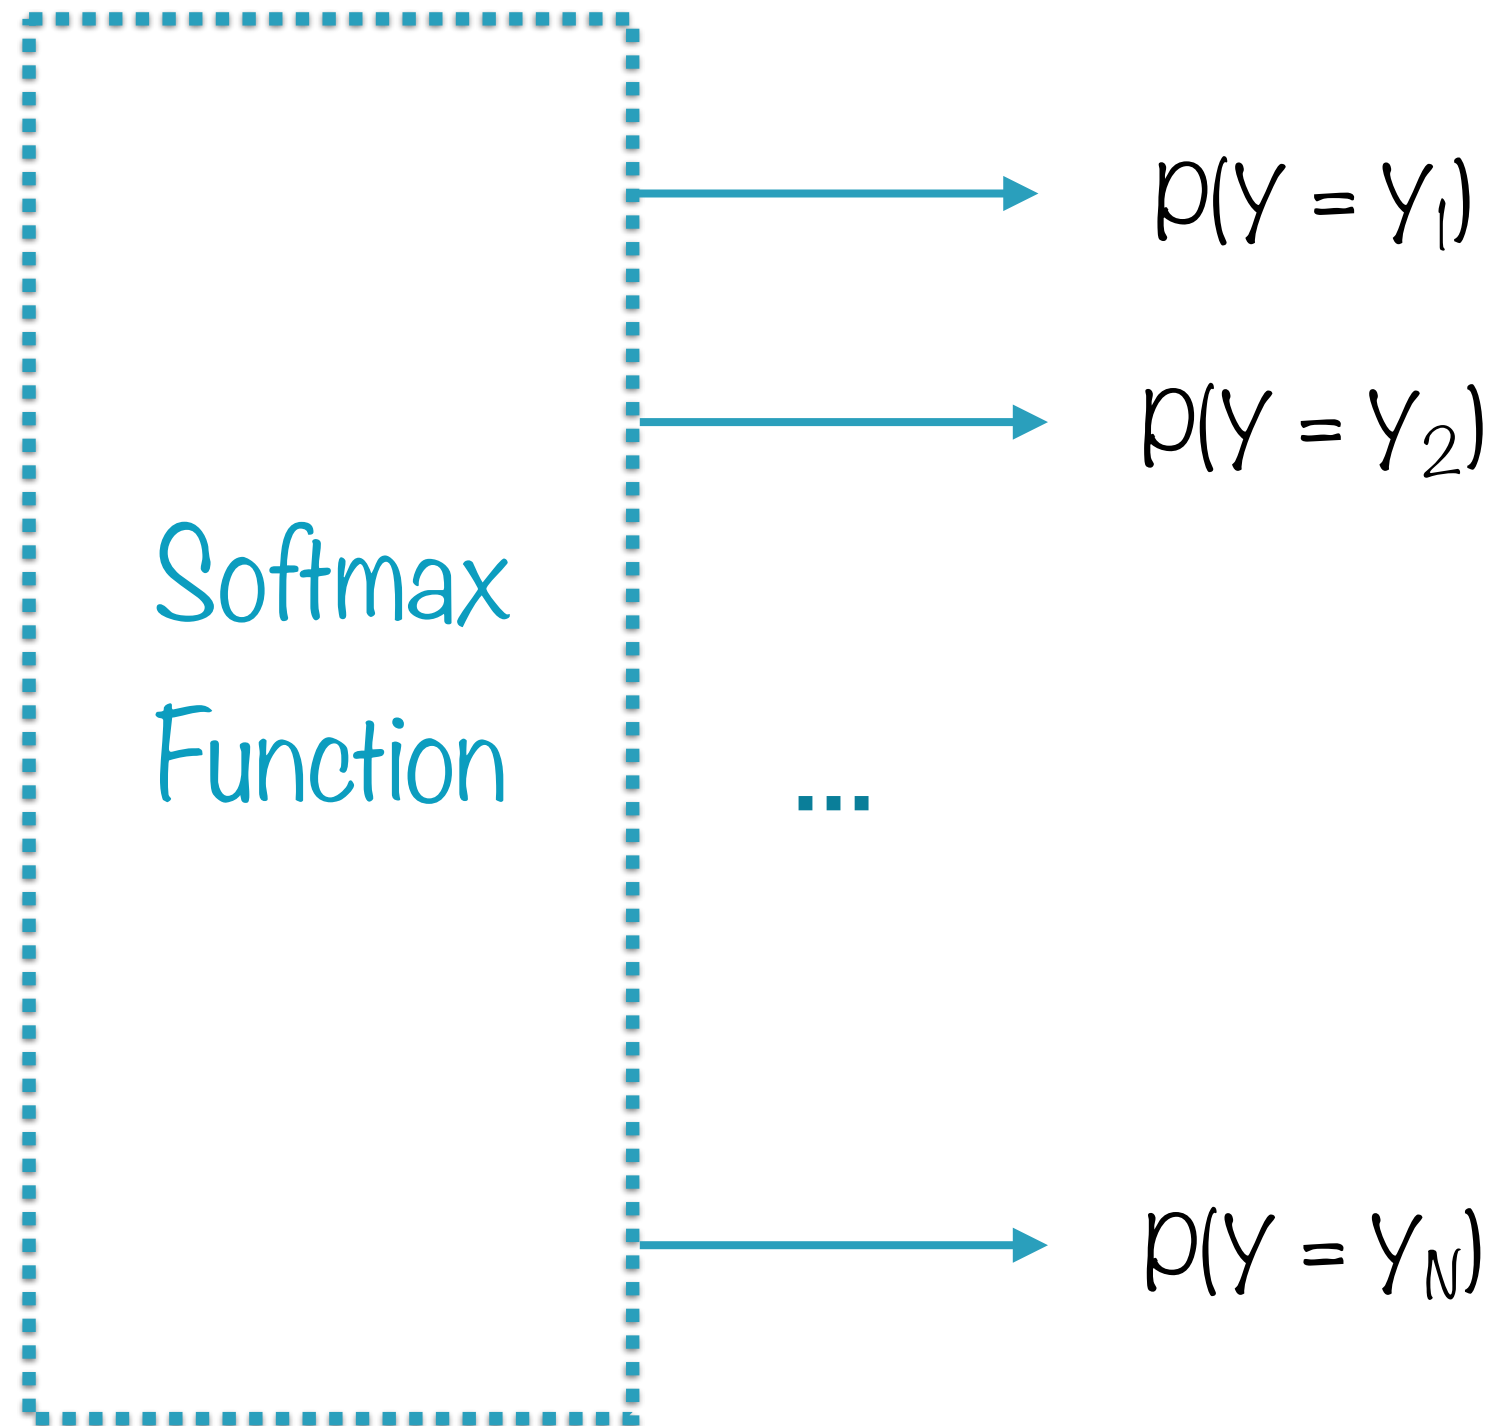
Linear Regression with One Neuron



Logistic Regression with One Neuron



SoftMax N-category Classification



Multilabel Digit Classification



One-versus-all: Train 10 binary classifiers

- 0-detector, 1-detector...
- Predicted label = output of detector with highest score


One-versus-one: Train 45 binary classifiers

- One detector for each pair of digits
- For N labels, need $N(N-1)/2$ classifiers
- Predicted label = output of digit that wins most duels


The logistic or softmax function is just one of many
that can be used for activation

Activation Function


ReLU




logit



tanh




step



Various choices of activation functions exist and drive the design of your neural network

ReLU Activation




The most common form of the activation function is the ReLU

ReLU : Rectified Linear Unit

$$\text{ReLU}(x) = \max(0, x)$$

Tanh Activation




S-shaped, continuous and differentiable

Output ranges from -1 to 1

Makes each layer's output normalized (centered around 0)

SoftMax Activation




Another very common form of the activation function is the SoftMax

$\text{SoftMax}(x)$ outputs a number between 0 and 1

This output can be interpreted as a **probability**

This curve is also called a logit curve


Importance of Activation



The choice of activation function is crucial in determining performance

To see why, we must understand the training process of a Neural Network


Neuron as a Learning Unit




Many of these simple neurons arranged in layers can do magical stuff

Training a Neural Network: Optimization and Back Propagation

A Neural Network



Example: Training for Linear Regression

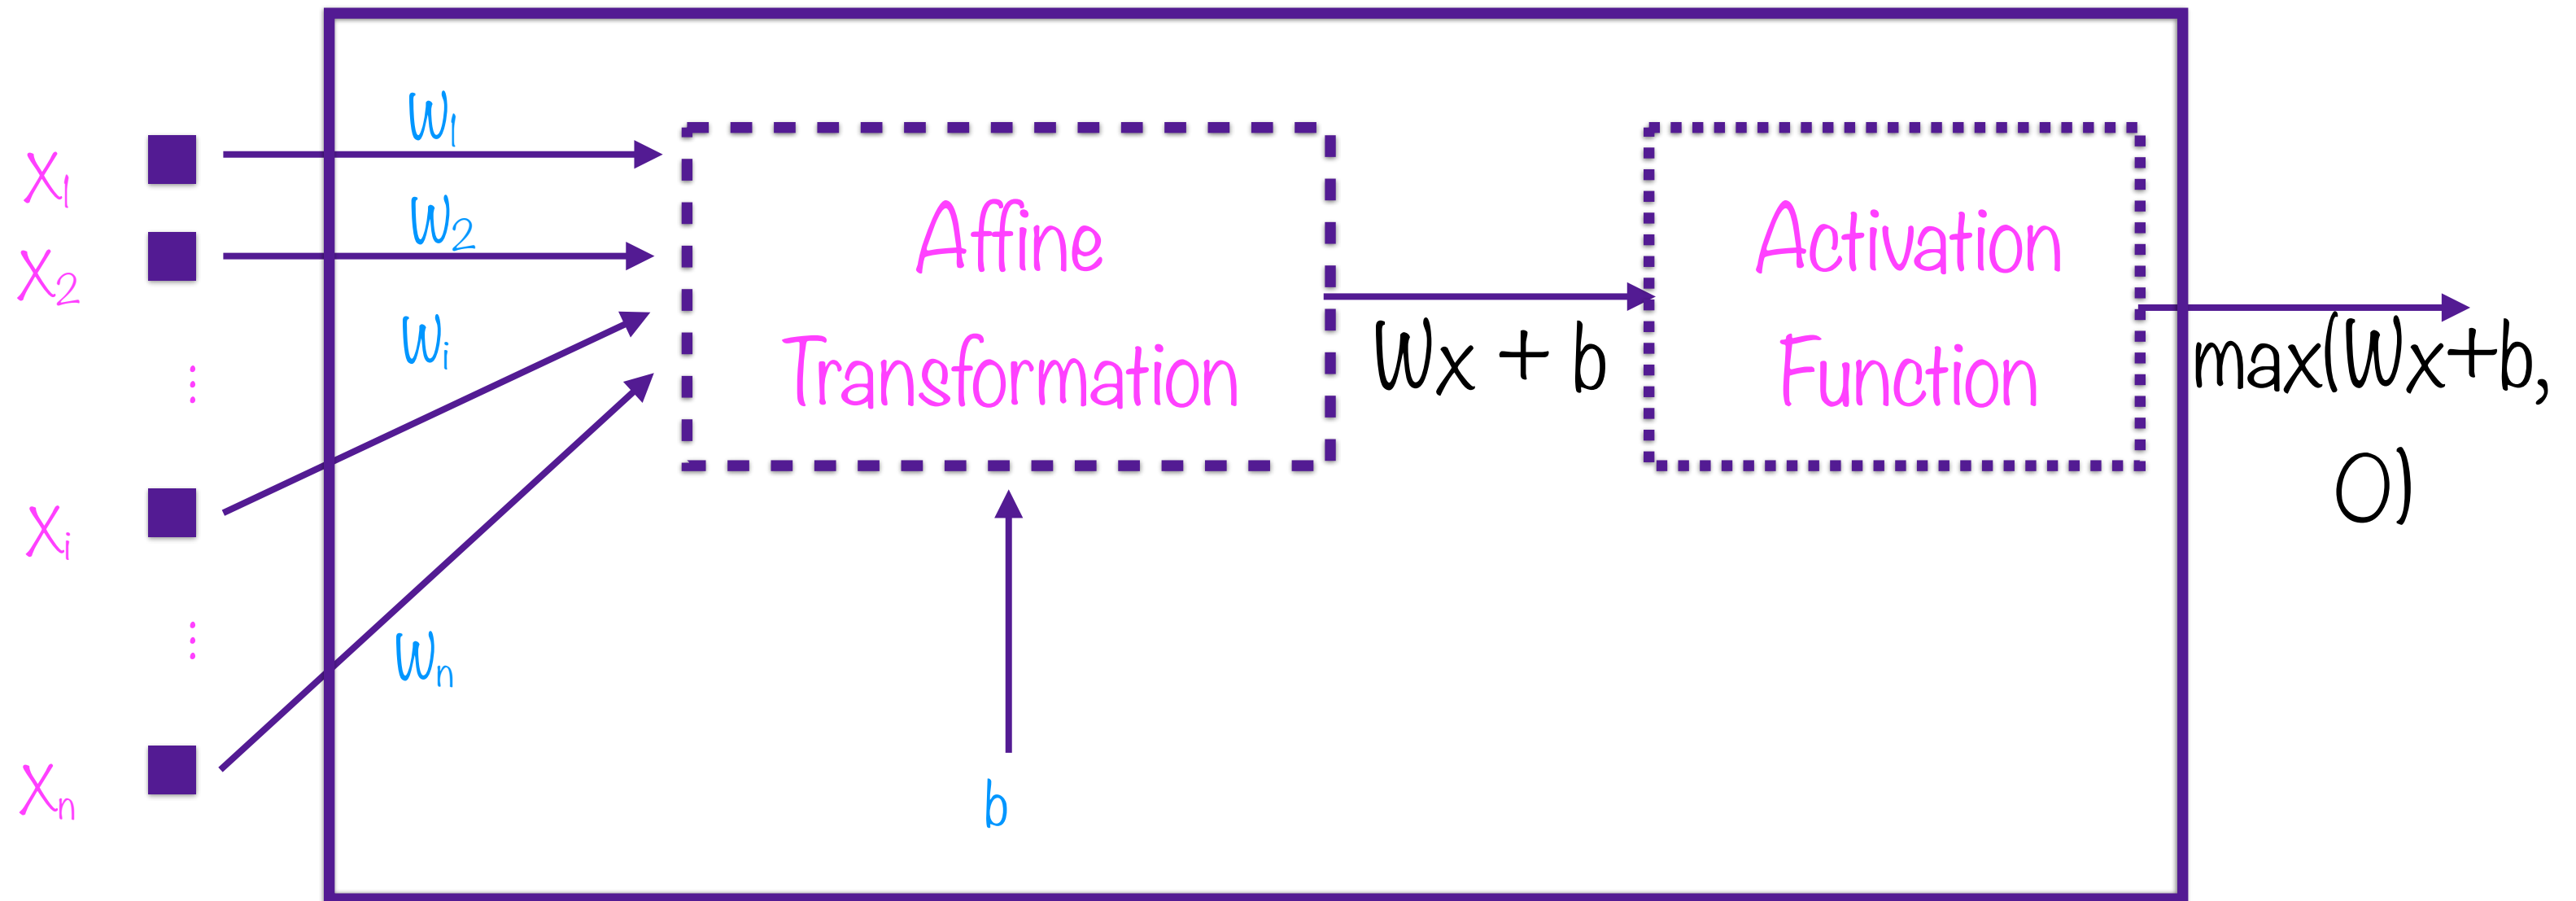


Set of
Points

Single neuron with no
activation function

Regression Line

Example: Training for Linear Regression



The activation function to learn linear regression is simply the identity function

Training as an Optimization Problem



Objective Function

Minimize variance of the residuals (MSE)



Constraints

Express relationship as a straight line

$$y = Wx + b$$



Decision Variables

Values of W and b


The “Best” Regression Line




The “best fit” line is called the regression line

The actual training of a neural network happens via
Gradient Descent Optimization


Minimizing MSE




Minimizing MSE



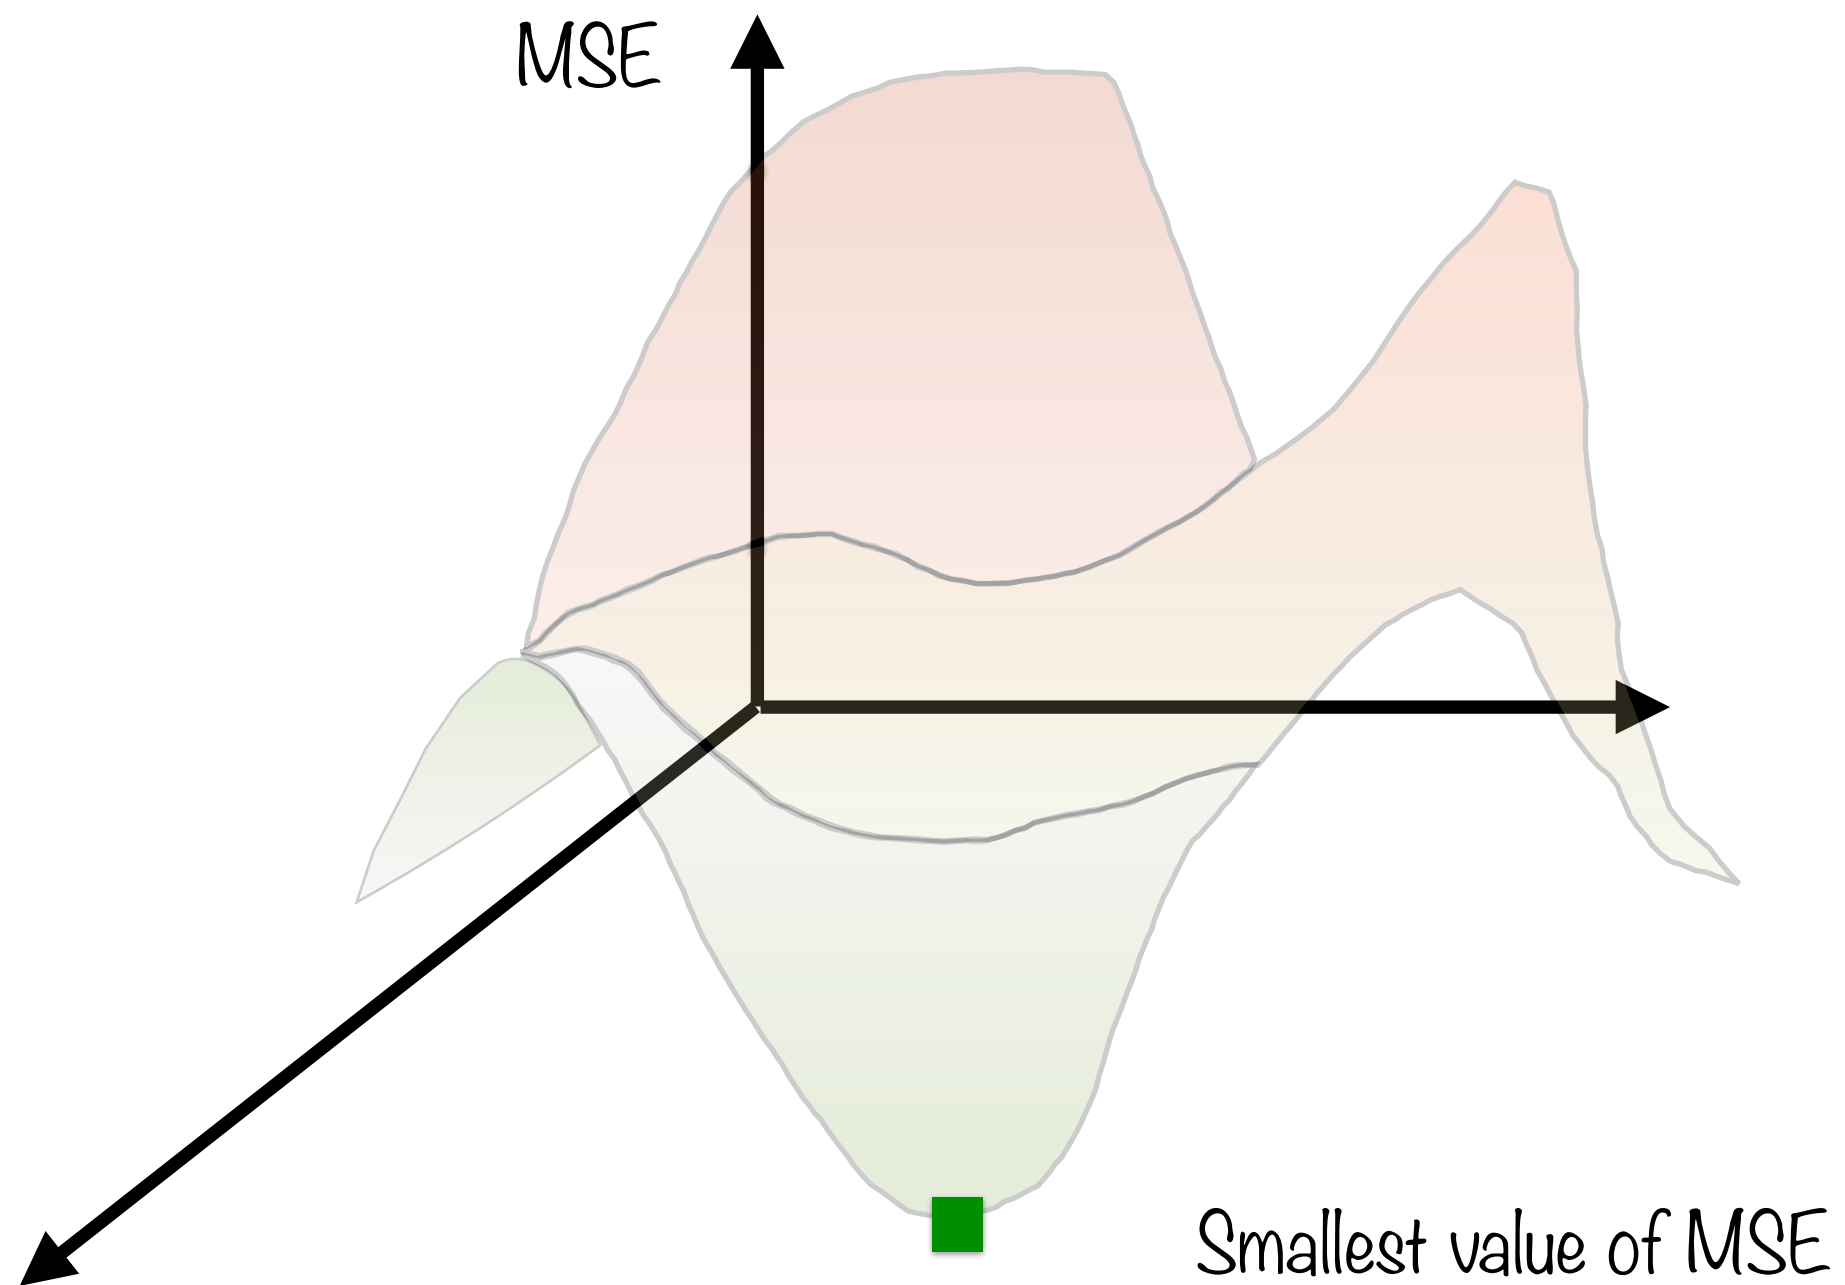
Minimizing MSE




Minimizing MSE



Minimizing MSE




Minimizing MSE




"Gradient Descent"


Converging on the "best" value
using an optimization algorithm




Minimizing MSE



"Training" the Algorithm




Start Somewhere




"Gradient Descent"

Converging on the "best" value
using an optimization algorithm




"Gradient Descent"


Converging on the "best" value
using an optimization algorithm




Training via Back Propagation




Training via Back Propagation




Training via Back Propagation




Training via Back Propagation



Training via Back Propagation



Training via Back Propagation



Back propagation allows the weights and biases of the neurons to **converge** to their final values

Hyperparameters

Decisions in Traditional ML Models

Initial values

Type of optimizer

Number of epochs

Batch size

More Decisions in Neural Networks

Network Topology i.e. neuron connections

Number of layers

Number of neurons in each layer

Activation function

How well the model performs is sensitive to
these decisions

These are **hyperparameters** of our model

Model Parameters vs. Hyperparameters

Model parameters

The weights and biases determined **during the training** process

Result of the training process

Used to **make predictions**

Measure using **validation datasets** to find the best possible model

Hyperparameters

The **design of the actual model** determined before the training process begins


Input to the training process

Used to **generate** the **best** possible model

Hyperparameter tuning to generate the model which is then evaluated using validation datasets

Vanishing, Exploding Gradients, Dying Neurons

Training via Back Propagation




Back Propagation


This is an iterative process

Fails either if

- gradients don't change at all
- gradients change too fast




“Vanishing Gradient Problem”



The weights of the earlier layers remain unchanged

The algorithm **never converges** to a
good solution


“Exploding Gradient Problem”



The weights of the layers become larger and
meaningless

The algorithm **diverges** and never
reaches a good solution

Vanishing and Exploding Gradients



Back propagation fails if

- gradients are vanishing
- gradients are exploding

**This was an important reason why DNNs
were mostly abandoned for a long time**

Coping with Vanishing/Exploding Gradients

Proper initialisation

Gradient clipping

Batch normalisation

Non-saturating activation function

Xavier and He Initialization

Proper
initialisation

The variance of the outputs in each direction is equal to variance of inputs

Connections weights must be initialized randomly

Xavier and He Initialization

Proper
initialisation

Normal distribution:

- mean 0
- standard deviation based on num_inputs and num_outputs for that layer

Xavier and He Initialization

Proper
initialisation

Uniform distribution:

- between $-r$ and $+r$
- r based on num_inputs and num_outputs for that layer

Coping with Vanishing/Exploding Gradients

Proper initialisation

Gradient clipping

Batch normalisation

Non-saturating activation function

Gradient Clipping

Gradient clipping

Limit the gradients to under a threshold during back propagation

Most often used with recurrent neural networks

Coping with Vanishing/Exploding Gradients

Proper initialisation

Gradient clipping

Batch normalisation

Non-saturating activation function

Batch Normalisation



Batch
normalisation

Zero center the inputs before passing to the activation functions

Subtract the mean and divide by the standard deviation

Allows use of saturating activation functions as well

Coping with Vanishing/Exploding Gradients

Proper initialisation

Gradient clipping

Batch normalisation

Non-saturating activation function

Coping with Vanishing/Exploding Gradients


Proper initialisation

Gradient clipping


Batch normalisation

Non-saturating activation function

A Neural Network




Unresponsive Neurons



What if the weights of the connections do not change in response to changing input?


Unresponsive Neurons




Neurons may be dead

Activation Functions


ReLU




logit



tanh




step




Various choices of activation functions exist and drive the design of your neural network

Activation Functions




Consider an S-shaped (sigmoid) activation function

Activation Functions




This is the **active** or **responsive** region of the function

Saturating Activation Functions




The activation function saturates at either end

Saturating Activation Functions




If a neuron operates within these **saturation regions** throughout training it might become unresponsive

Dying Neurons




- Neuron might become unresponsive - output won't change as input changes
- If this continues throughout training, neuron is "dead"
- Saturation of neuron occurs at both ends of S-curve, for instance

Saturating Activation Functions



Logit Activation


Saturates for very large and very small values of input



ReLU Activation

Saturates for very small (negative values) of inputs

ELU Activation



Mitigates dying-neuron problem of ReLU


- Linear for positive values
- Exponential for negative values

ELU is the new favorite activation function

Dying and unresponsive neuron issues can be mitigated
by using activations functions such as ELU


Overfitting and the Bias-Variance Trade-off

Connecting the Dots




Challenge: Fit the “best” curve through these points

Good Fit?




A curve has a “good fit” if the distances of points from the curve are small

Connecting the Dots




We could draw a pretty complex curve

Connecting the Dots




We can even make it pass through every single point

Connecting the Dots




But given a new set of points, this curve might perform quite poorly

Connecting the Dots




The original points were “training data”, the new points are “test data”

Overfitting




Great performance in training, poor performance in real
usage

Connecting the Dots




A simple straight line performs worse in training, but better with
test data

Overfitting



Low Training Error


Model does very well in training...



High Test Error

...but poorly with real data

Cause of Overfitting




Sub-optimal choice in the **bias-variance** trade-off

An overfitted model has:

- high variance error
- low bias error




Bias



Low bias

Few assumptions about the underlying data




High bias

More assumptions about the underlying data




Bias



Model too complex

Training data all-important, model parameter counts for little




Model too simple

Model parameter all-important, training data counts for little




Variance



High variance

The model changes significantly when training data changes




Low variance

The model doesn't change much when the training data changes




Variance



Model too complex


Model varies too much with changing training data



Model too simple


Model not very sensitive to training data

Bias-Variance Trade-off



Model too complex


High variance error



Model too simple

High bias error

Bias-Variance Trade-off




- High-bias algorithms: simple parameters
 - Regression
- High-variance algorithms: complex parameters
 - Decision trees
 - Dense neural networks

Mitigating Overfitting

Preventing Overfitting


- Regularisation
- Cross-validation
- Ensemble learning
 - Dropout



Preventing Overfitting



Regularisation - Penalise complex models



Cross-validation - Distinct training and validation phases



Dropout - Intentionally turn off some neurons during training

Regularisation



Penalise complex models


Simple in Gradient Descent

Add penalty to objective function

Penalty as function of neuron weights

Forces optimiser to keep it simple

Cross-Validation




Distinct training and validation phases

Train different models (with training data only)

Select model that does best on validation data

“Hyperparameter tuning”



Dropout


Specify a fraction of neurons that will stay off in each training step

“Dropout” neurons chosen at random

Different neurons off in each training step

In effect, each training step builds different network configuration

Densely Connected Neural Network




Corpus of
Images


High risk of overfitting during training due to
dense, complex network

ML-based Classifier

Dropout = 50%



Corpus of
Images




Randomly switch off say 50% of neurons in
each training step




ML-based Classifier

Dropout = 50%



Corpus of
Images




Training forced to rely on a much simpler neural
network




ML-based Classifier

Dropout = 50%



Corpus of
Images




Each training step will build a different
configuration




ML-based Classifier

Dropout = 50%



Corpus of
Images




Each training step will build a different
configuration



ML-based Classifier

Dropout During Training Only



Corpus of
Images

During actual usage in test mode, full dense
neural network is used

ML-based Classifier