

Convolutional Neural Networks in TensorFlow


Overview

Convolutional NNs are one kind of NN architecture which work well with 2D data

Modeled on the visual cortex, they are amazing at image classification


How Do We See?

Viewing an Image




All neurons in the eye don't see the entire image

Viewing an Image




Each neuron has its own local receptive field

Viewing an Image




It reacts only to visual stimuli located in its receptive field

Viewing an Image



Some neurons react to more complex patterns that are
combinations of lower level patterns

Neural Networks



Sounds like a classic neural network problem

Two Kinds of Layers in CNNs

Convolution

Local receptive field

Pooling

Subsampling of inputs

Convolution

Two Kinds of Layers in CNNs

Convolution

Local receptive field

Pooling

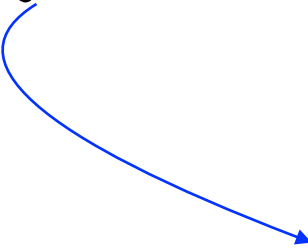
Subsampling of inputs

Convolution

In this context, a sliding window function applied to a matrix

Convolution

In this context, a sliding window function applied to
a matrix



e.g. a matrix of pixels representing
an image

Convolution

In this context, a sliding window **function** applied to
a matrix

Often called a kernel or filter




Convolution

In this context, a sliding window **function** applied to a matrix




Kernel is applied element-wise in sliding-window fashion

Representing Images as Matrices



= 784 pixels

Representing Images as Matrices




= 36 pixels

Representing Images

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix




Kernel

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



Kernel

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



	x1	x0	x1
x0		x1	x0
x1	x0		x1

4

4

1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8	2.0	1.8

Convolution Result

Convolution

0_{xl}	$x0$	0_{xl}	0	0	0
$x0$	0.8_{xl}	$x0$	0.3	0.6	0
0.2_{xl}	$x0$	0_{xl}	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



Convolution Result

Convolution

0_{xl}	$x0$	0_{xl}	0	0	0
$x0$	0.8_{xl}	$x0$	0.3	0.6	0
0.2_{xl}	$x0$	0_{xl}	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1			

Convolution Result

Convolution

0	0_{x1}	$x0$	0_{x1}	0	0
0.2	$x0$	0_{x1}	$x0$	0.6	0
0.2	0_{x1}	$x0$	0.3_{x1}	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1			

Convolution Result

Convolution

0	0_{x1}	$x0$	0_{x1}	0	0
0.2	$x0$	0_{x1}	$x0$	0.6	0
0.2	0_{x1}	$x0$	0.3_{x1}	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2		

Convolution Result

Convolution

0	0	0_{x1}	$x0$	0_{x1}	0
0.2	0.8	$x0$	0.3_{x1}	$x0$	0
0.2	0.9	0_{x1}	$x0$	0.8_{x1}	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2		

Convolution Result

Convolution

0	0	0_{x1}	$x0$	0_{x1}	0
0.2	0.8	$x0$	0.3_{x1}	$x0$	0
0.2	0.9	0_{x1}	$x0$	0.8_{x1}	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	

Convolution Result

Convolution

0	0	0	0_{xl}	$x0$	0_{xl}
0.2	0.8	0	$x0$	0.6_{xl}	$x0$
0.2	0.9	0	0.3_{xl}	$x0$	0_{xl}
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	

Convolution Result

Convolution

0	0	0	0_{xl}	$x0$	0_{xl}
0.2	0.8	0	$x0$	0.6_{xl}	$x0$
0.2	0.9	0	0.3_{xl}	$x0$	0_{xl}
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9

Convolution Result

Convolution

0	0	0	0	0	0
0.2 _{x1}	x0	0 _{x1}	0.3	0.6	0
x0	0.9 _{x1}	x0	0.3	0.8	0
0.3 _{x1}	x0	0.7 _{x1}	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9

Convolution Result

Convolution

0	0	0	0	0	0
0.2 _{x1}	x0	0 _{x1}	0.3	0.6	0
x0	0.9 _{x1}	x0	0.3	0.8	0
0.3 _{x1}	x0	0.7 _{x1}	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9			

Convolution Result

Convolution

0	0	0	0	0	0
0.2	0.8 _{x1}	0 _{x0}	0.3 _{x1}	0.6	0
0.2	0 _{x0}	0 _{x1}	0 _{x0}	0.8	0
0.3	0.8 _{x1}	0 _{x0}	0.8 _{x1}	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9			

Convolution Result

Convolution

0	0	0	0	0	0
0.2	0.8 _{x1}	0 _{x0}	0.3 _{x1}	0.6	0
0.2	0 _{x0}	0 _{x1}	0 _{x0}	0.8	0
0.3	0.8 _{x1}	0 _{x0}	0.8 _{x1}	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7		

Convolution Result

Convolution

0	0	0	0	0	0
0.2	0.8	0_{x1}	$x0$	0.6_{x1}	0
0.2	0.9	$x0$	0.3_{x1}	$x0$	0
0.3	0.8	0.7_{x1}	$x0$	0.9_{x1}	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7		

Convolution Result

Convolution

0	0	0	0	0	0
0.2	0.8	0_{x1}	$x0$	0.6_{x1}	0
0.2	0.9	$x0$	0.3_{x1}	$x0$	0
0.3	0.8	0.7_{x1}	$x0$	0.9_{x1}	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	

Convolution Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3 _{x1}	x0	0 _{x1}
0.2	0.9	0	x0	0.8 _{x1}	x0
0.3	0.8	0.7	0.8 _{x1}	x0	0 _{x1}
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	

Convolution Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3 _{x1}	x0	0 _{x1}
0.2	0.9	0	x0	0.8 _{x1}	x0
0.3	0.8	0.7	0.8 _{x1}	x0	0 _{x1}
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9

Convolution Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2 _{x1}	0 _{x0}	0 _{x1}	0.3	0.8	0
0 _{x0}	0.8 _{x1}	0 _{x0}	0.8	0.9	0
0 _{x1}	0 _{x0}	0 _{x1}	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9

Convolution Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2 _{x1}	x0	0 _{x1}	0.3	0.8	0
x0	0.8 _{x1}	x0	0.8	0.9	0
0 _{x1}	x0	0 _{x1}	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0			

Convolution Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9 _{x1}	x0	0.3 _{x1}	0.8	0
0.3	x0	0.7 _{x1}	x0	0.9	0
0	0 _{x1}	x0	0.2 _{x1}	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0			

Convolution Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9 _{x1}	x0	0.3 _{x1}	0.8	0
0.3	x0	0.7 _{x1}	x0	0.9	0
0	0 _{x1}	x0	0.2 _{x1}	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1		

Convolution Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0 _{x1}	0 _{x0}	0.8 _{x1}	0
0.3	0.8	0 _{x0}	0.8 _{x1}	0 _{x0}	0
0	0	0 _{x1}	0 _{x0}	0.8 _{x1}	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1		

Convolution Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0 _{x1}	x0	0.8 _{x1}	0
0.3	0.8	x0	0.8 _{x1}	x0	0
0	0	0 _{x1}	x0	0.8 _{x1}	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	

Convolution Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3 _{xl}	x0	0 _{xl}
0.3	0.8	0.7	x0	0.9 _{xl}	x0
0	0	0	0.2 _{xl}	x0	0 _{xl}
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	

Convolution Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3 _{x1}	x0	0 _{x1}
0.3	0.8	0.7	x0	0.9 _{x1}	x0
0	0	0	0.2 _{x1}	x0	0 _{x1}
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4

Convolution Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3 _{x1}	x0	0.7 _{x1}	0.8	0.9	0
x0	0 _{x1}	x0	0.2	0.8	0
0 _{x1}	x0	0 _{x1}	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4

Convolution Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3 _{xl}	x0	0.7 _{xl}	0.8	0.9	0
x0	0 _{xl}	x0	0.2	0.8	0
0 _{xl}	x0	0 _{xl}	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0			

Convolution Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8 _{x1}		0.8 _{x1}	0.9	0
0		0 _{x1}		0.8	0
0	0 _{x1}		0.2 _{x1}	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0			

Convolution Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8 _{x1}		0.8 _{x1}	0.9	0
0		0 _{x1}		0.8	0
0	0 _{x1}		0.2 _{x1}	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8		

Convolution Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7 _{x1}	x0	0.9 _{x1}	0
0	0	x0	0.2 _{x1}	x0	0
0	0	0 _{x1}	x0	0.2 _{x1}	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8		

Convolution Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7 _{x1}	x0	0.9 _{x1}	0
0	0	x0	0.2 _{x1}	x0	0
0	0	0 _{x1}	x0	0.2 _{x1}	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8	2.0	

Convolution Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8 _{xl}	x0	0 _{xl}
0	0	0	x0	0.8 _{xl}	x0
0	0	0	0.2 _{xl}	x0	0 _{xl}

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8	2.0	

Convolution Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8 _{xl}	x0	0 _{xl}
0	0	0	x0	0.8 _{xl}	x0
0	0	0	0.2 _{xl}	x0	0 _{xl}

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8	2.0	1.8

Convolution Result

Choice of Kernel Function

x0		x0	x0
x0	x1	x1	x0
x0	x1	x1	x0
x0	x1	x1	x0
x0			x0

Averaging neighbouring pixels ~ **Blurring**

Subtracting neighbouring pixels ~ **Edge detection**

Positive middle, negative neighbours ~ **Sharpen**

Negative corners, zero elsewhere ~ **Edge enhance**

More complex patterns ~ **Emboss**

...

Choice of Kernel Function

[**http://aishack.in/tutorials/image-convolution-examples/**](http://aishack.in/tutorials/image-convolution-examples/)

Blur

$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$



A simple blur done with convolutions

Line Detection

-1	-1	-1
2	2	2
-1	-1	-1

Horizontal lines

-1	2	-1
-1	2	-1
-1	2	-1

Vertical lines

-1	-1	2
-1	2	-1
2	-1	-1

45 degree lines

2	-1	-1
-1	2	-1
-1	-1	2

135 degree lines

Horizontal Lines

-1	-1	-1
2	2	2
-1	-1	-1

Horizontal lines





Edge Detection

-1	-1	-1
-1	8	-1
-1	-1	-1





Zero-padding, Stride Size

Narrow vs. Wide Convolution




Input matrix i.e. image

Narrow vs. Wide Convolution




Convolution result

Narrow vs. Wide Convolution



Narrow Convolution

Little zero padding; output narrower than input



Wide Convolution

Lots of zero padding; output wider than input

Without Zero Padding

6

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

6

Matrix




x'	x0	x'
x0	x'	x0
x'	x0	x'

4


4

Convolution Result


Zero Padding



Matrix




8





Convolution Result

Zero Padding



Matrix



Convolution Result

Zero Padding

x0		x0	x0
x0	x1	x1	x0
x0	x1	x1	x0
x0	x1	x1	x0
x0			x0

With zero-padding, every element of matrix will be passed into filter

Can decide number of zero columns to pad with

Use to get output larger than input

Stride Size

0_{xl}	$x0$	0_{xl}	0	0	0
$x0$	0.8_{xl}	$x0$	0.3	0.6	0
0.2_{xl}	$x0$	0_{xl}	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Stride Size



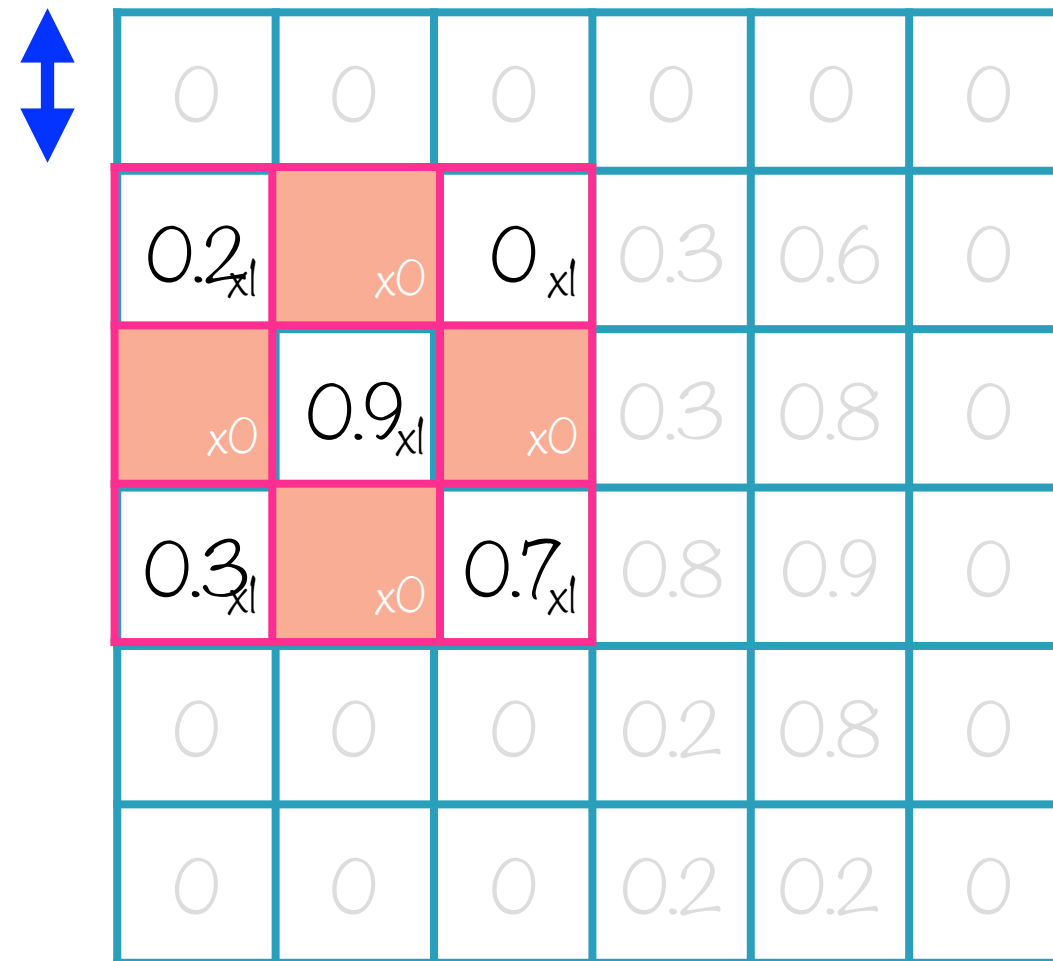
0	0_{xl}	$x0$	0_{xl}	0	0
0.2	$x0$	0_{xl}	$x0$	0.6	0
0.2	0.9_{xl}	$x0$	0.3_{xl}	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Horizontal stride of 1

Stride Size

0_{xl}	$x0$	0_{xl}	0	0	0
$x0$	0.8_{xl}	$x0$	0.3	0.6	0
0.2_{xl}	$x0$	0_{xl}	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

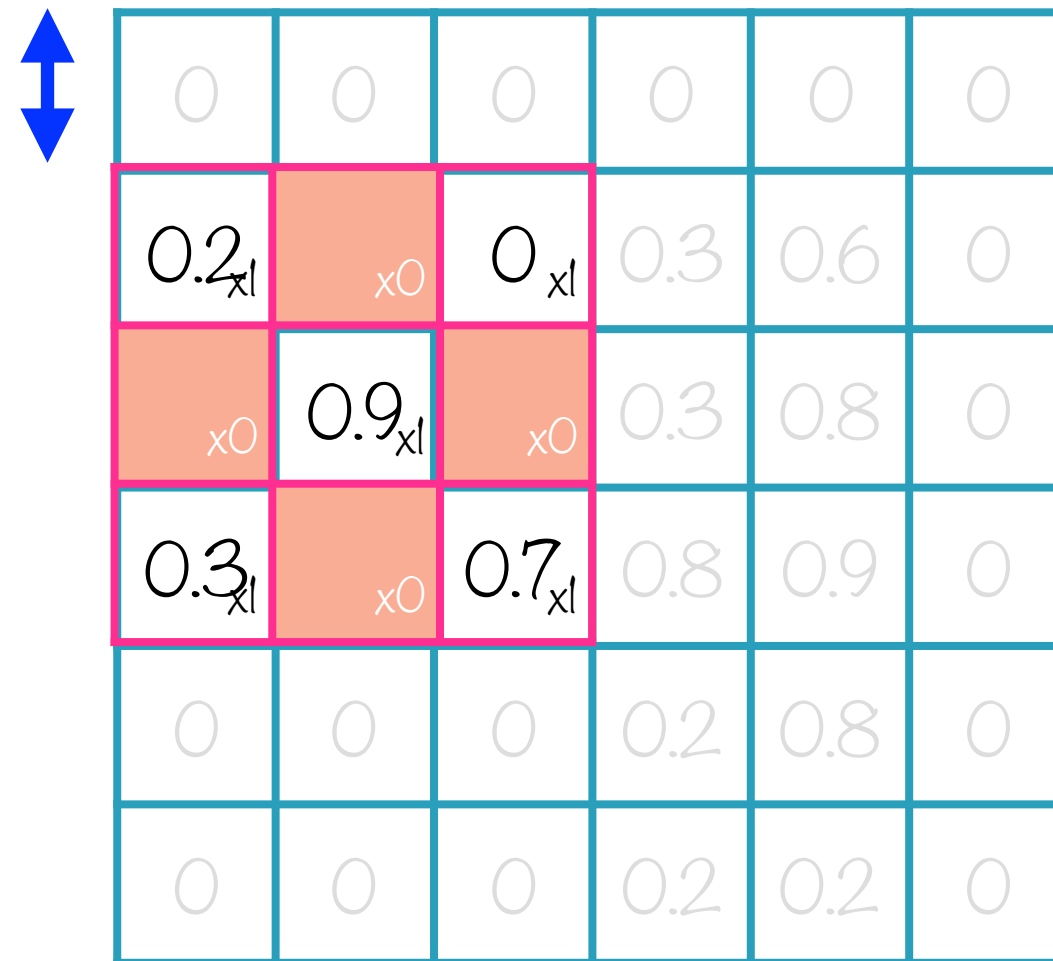
Stride Size



0	0	0	0	0	0
0.2 _{xl}	x0	0 _{xl}	0.3	0.6	0
x0	0.9 _{xl}	x0	0.3	0.8	0
0.3 _{xl}	x0	0.7 _{xl}	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Vertical stride of 1

Stride Size

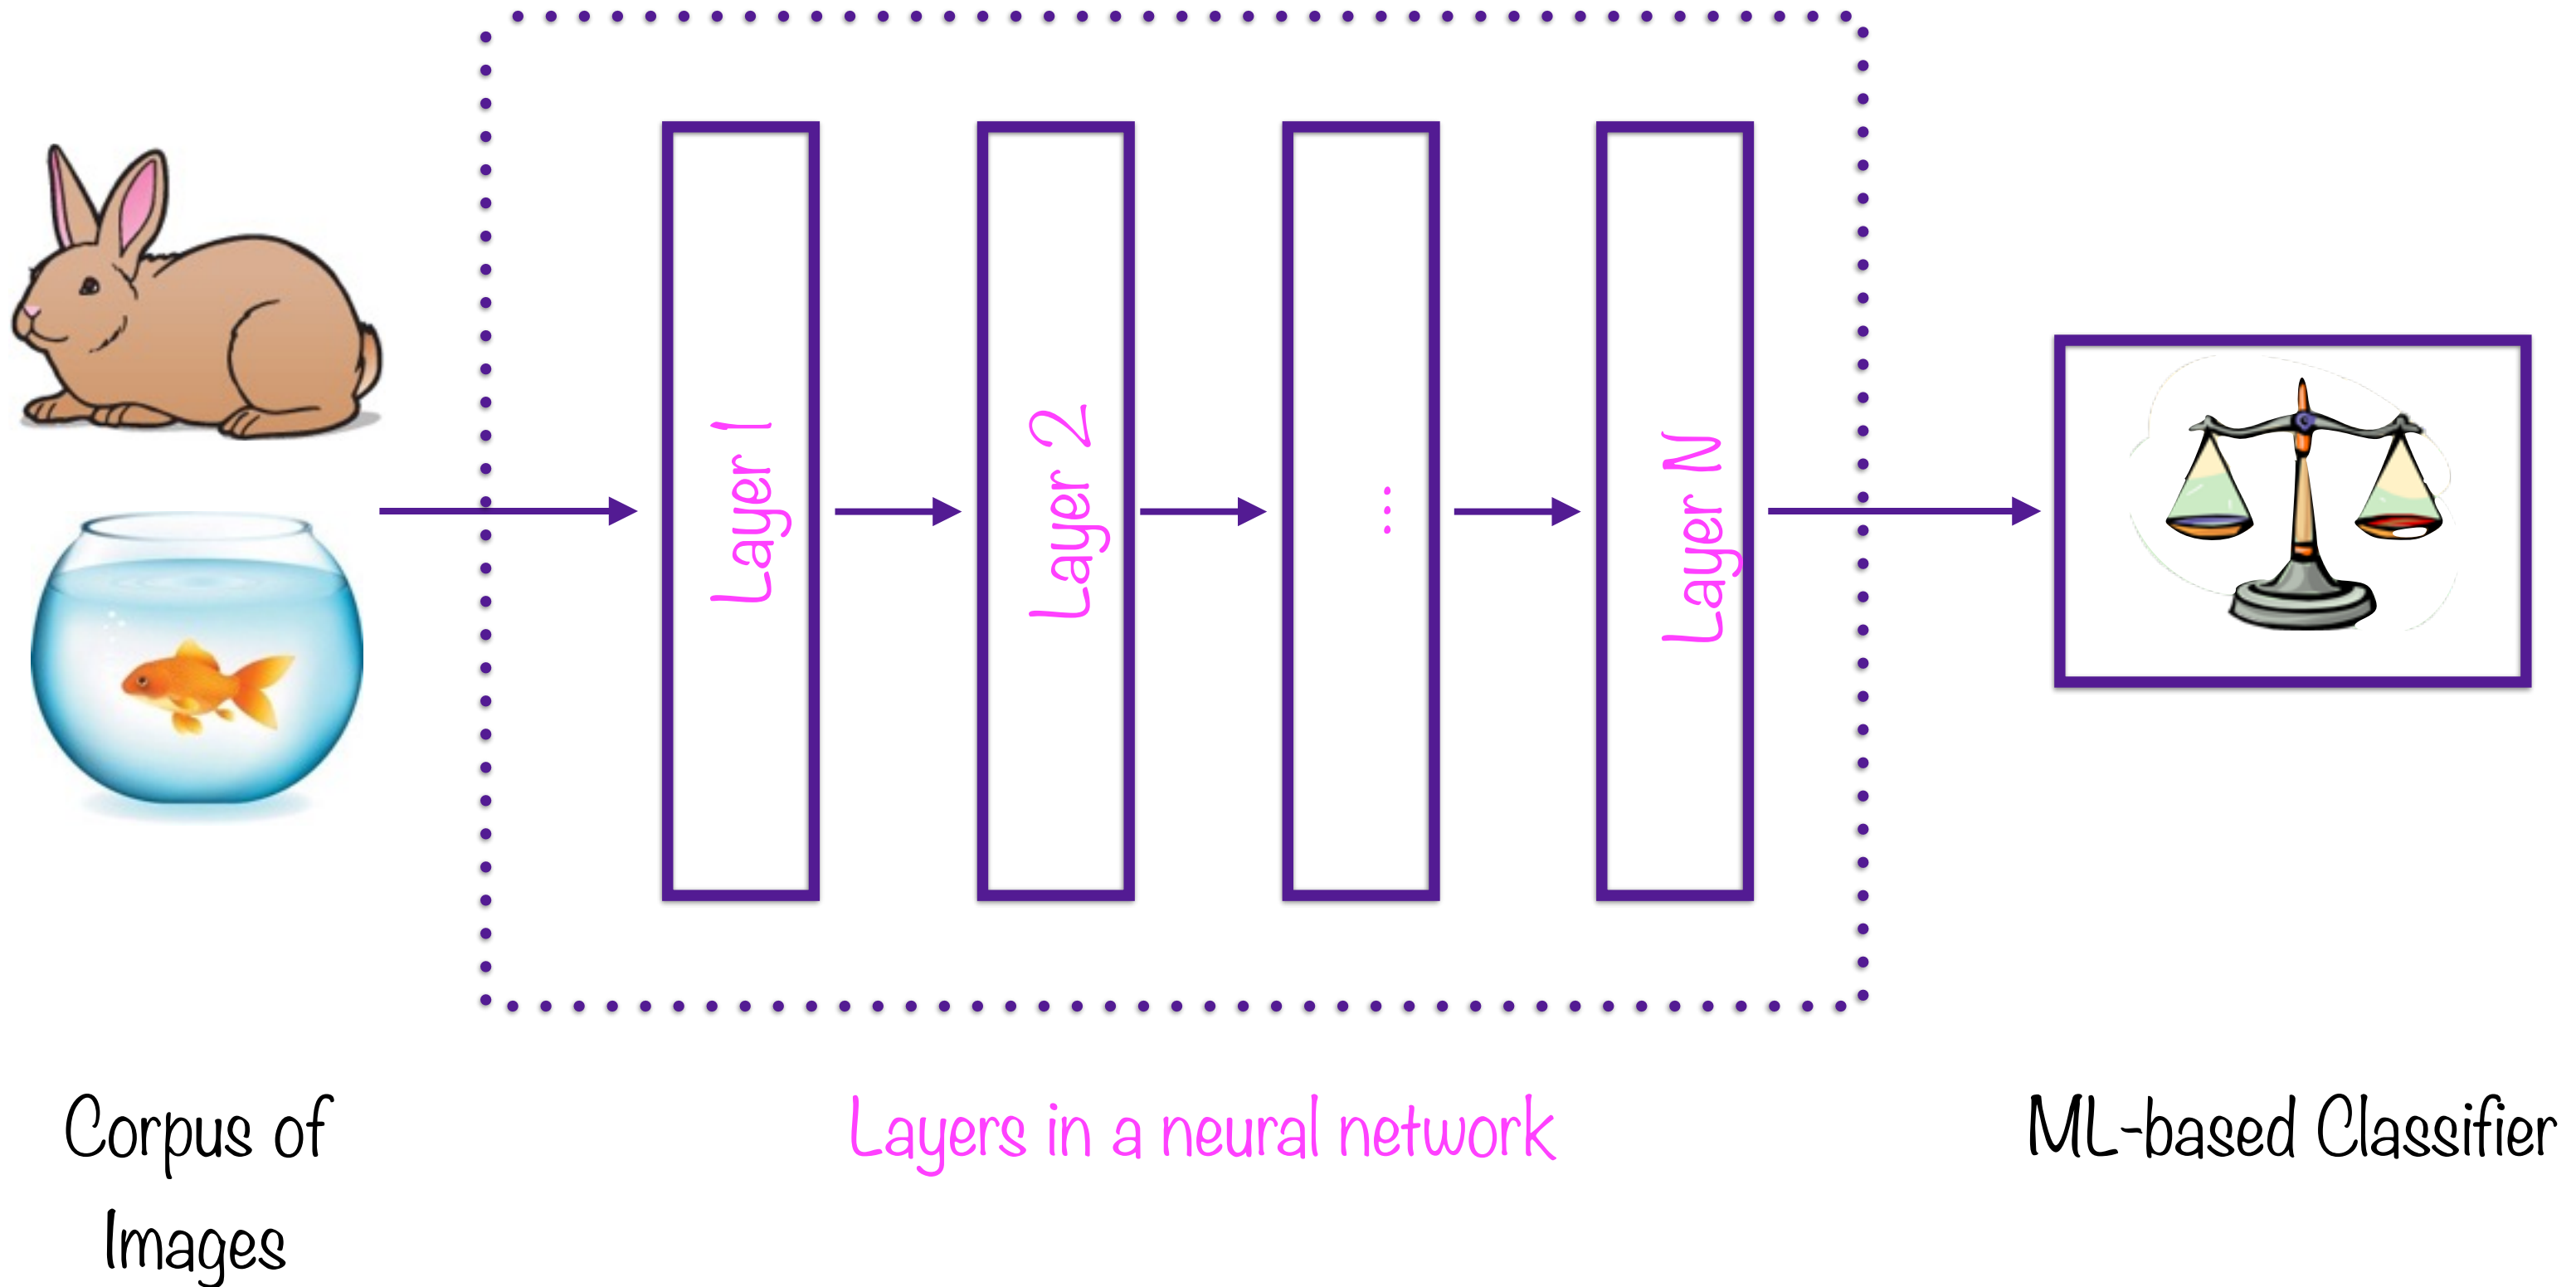


0	0	0	0	0	0
0.2 _{x1}	x0	0 _{x1}	0.3	0.6	0
x0	0.9 _{x1}	x0	0.3	0.8	0
0.3 _{x1}	x0	0.7 _{x1}	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

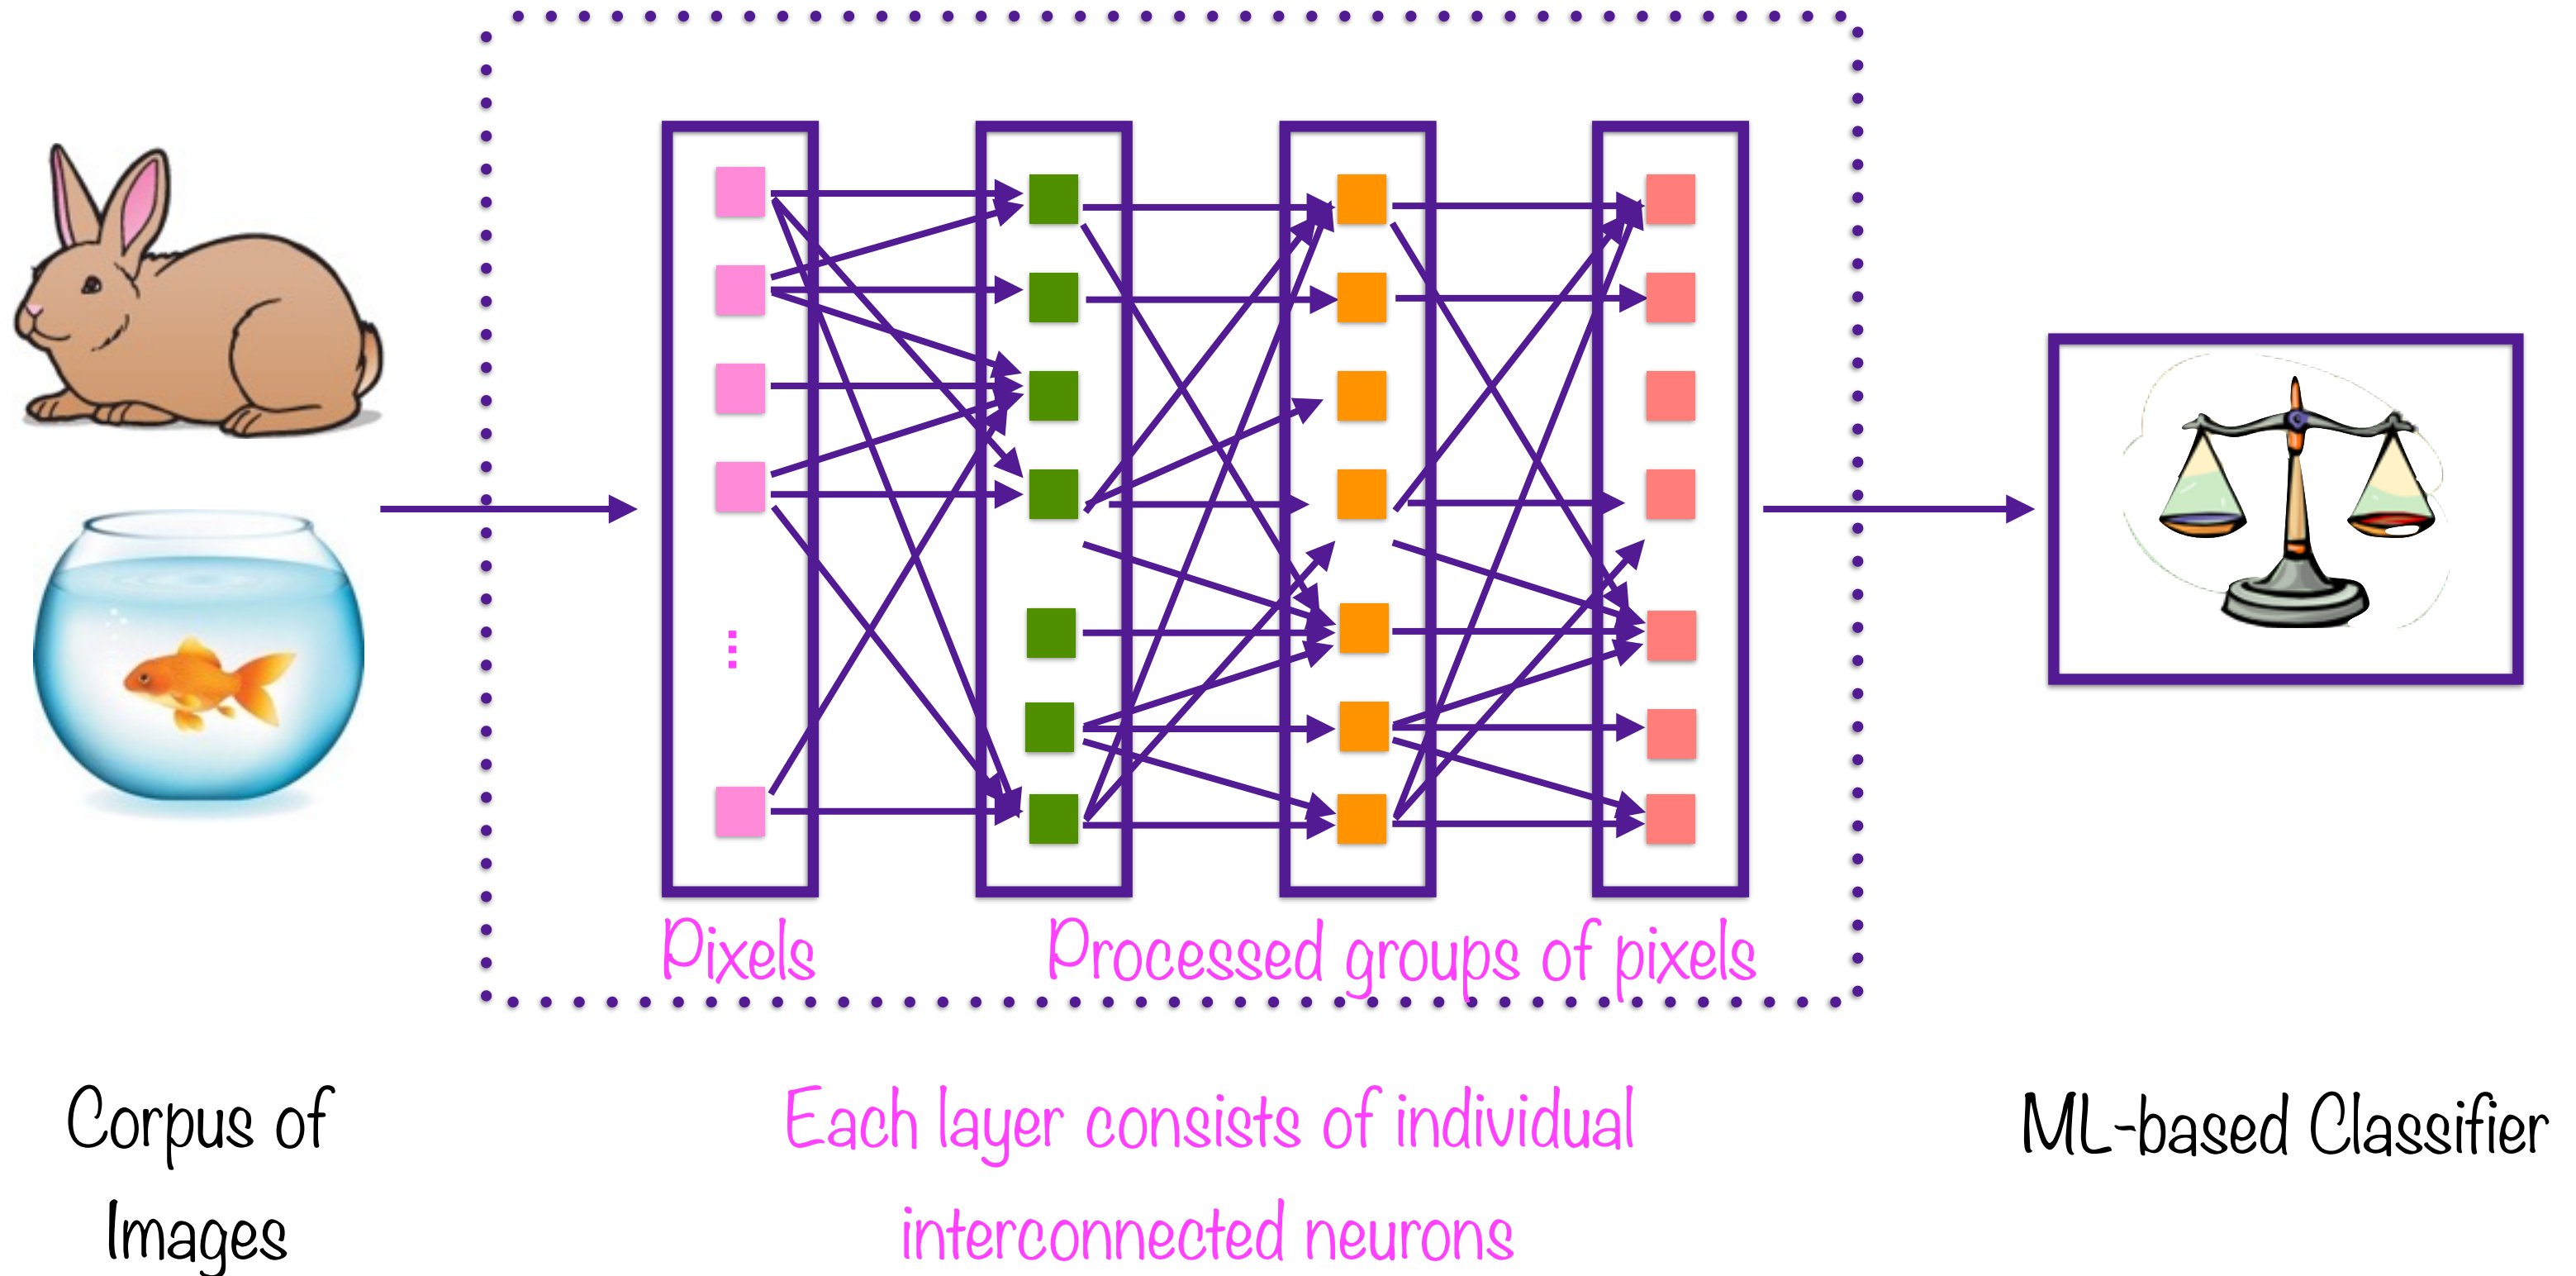
Stride size is an important hyper parameter in
CNNs

Convolutional Neural Networks

Neural Networks for Image Classification



Neural Networks for Image Classification



Parameter Explosion



Consider a 100 x 100 pixel image (10,000 pixels)

If first layer = 10,000 neurons

Interconnections $\sim O(10,000 * 10,000)$

100 million parameters to train neural network!

Parameter Explosion



Dense, fully connected neural networks can't cope

Convolutional neural networks to the rescue

CNNs Introduced



Eye perceives visual stimulus in 2D visual field

Eye sends 2D image to visual cortex

Visual cortex adds depth perception

Individual neurons in cortex focus on small field

“Local receptive field”

CNNs Introduced



CNNs perform spectacularly well at many tasks

Particularly at image recognition

Dramatically fewer parameters than DNN with similar performance

Inspirations for CNNs



Two Dimensions

Data comes in expressed in 2D



Local Receptive Fields

Neurons focus on narrow portions

CNN Layers




Convolution layers - zoom in on specific bits of input

Successive layers aggregate inputs into higher level features


Pixels >> Lines >> Contours/Edges >> Object

Convolutional Layers


Feature Maps



Image




Pixels




Feature Map

Feature Maps



Pixels




Convolutional Layer

Feature Maps

Local Receptive
Field of Neuron i

Neuron i




Pixels


Convolutional Layer

Feature Maps

Number of neurons in
receptive field = kernel
size



Kernel Size

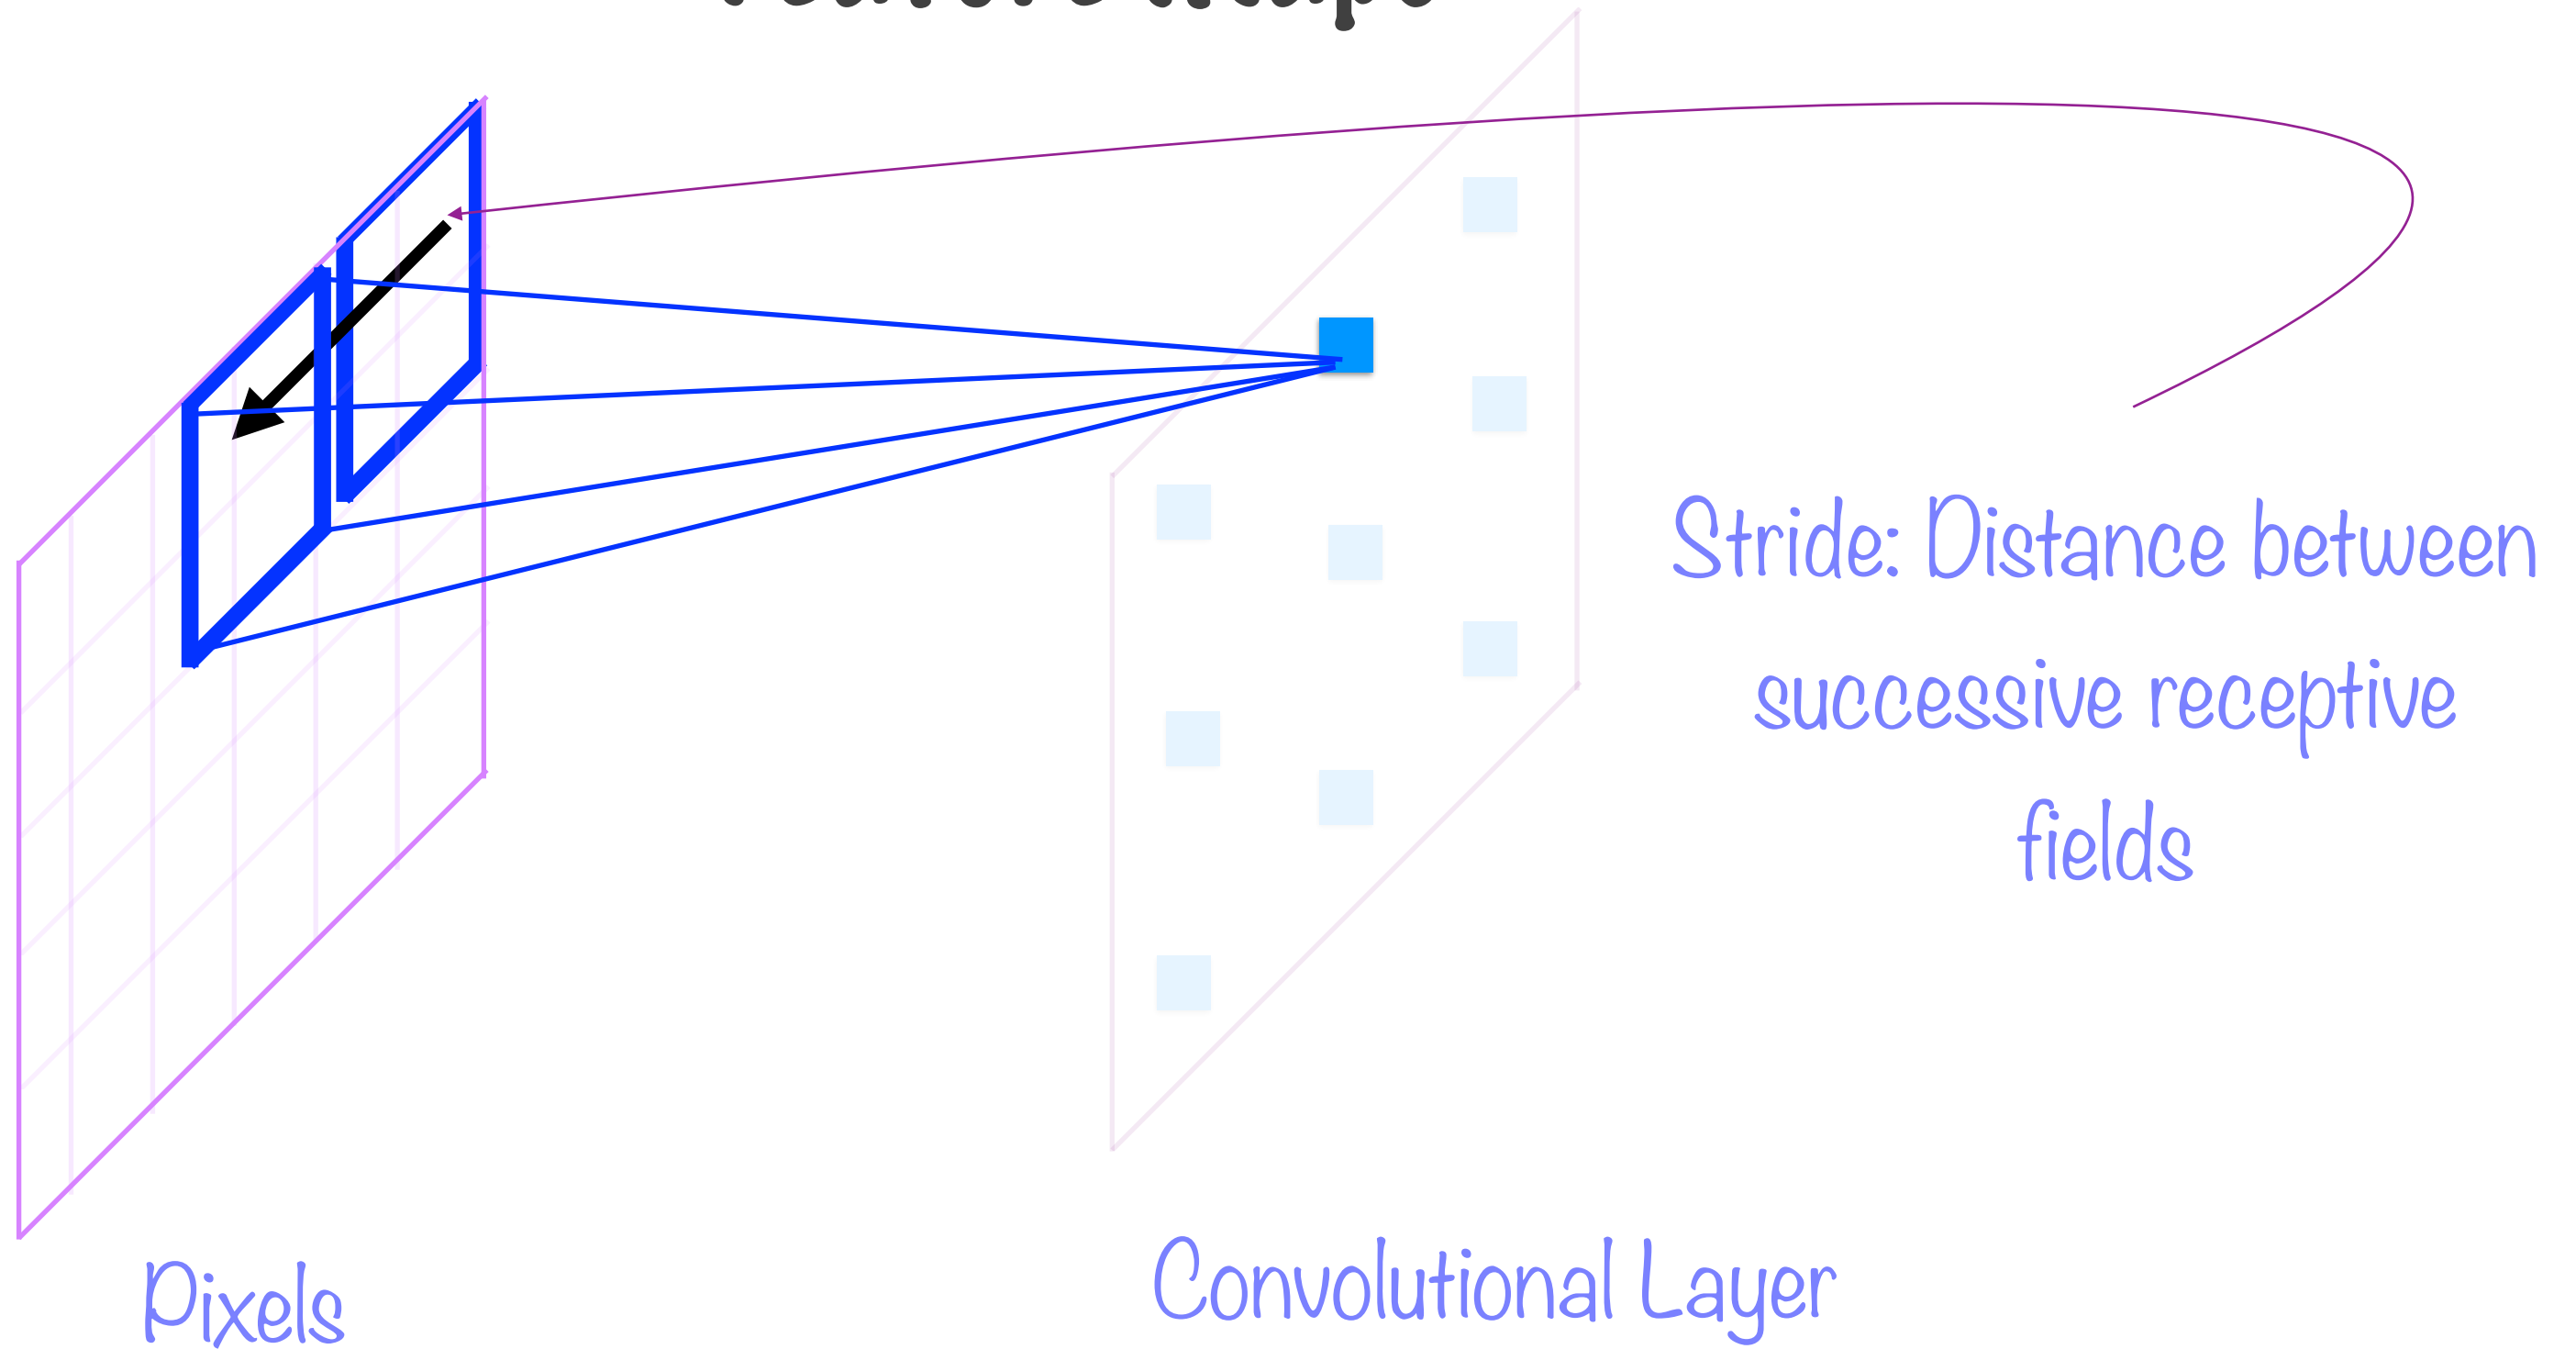


The convolutional kernel size is usually expressed in terms of width and height of receptive area


Use small convolutional kernels, more efficient

Stacking 2 3x3 kernels is preferable to 1 9x9 kernel


Feature Maps




Feature Maps




Feature Maps



Feature Maps




Feature Maps




Feature Maps

Zero padding may
be needed at the

edges



Feature Maps




All neurons in a feature map have the same weights and biases

Two big advantages over DNNs

- Dramatically **fewer** parameters to train
- CNN can recognise feature patterns **independent** of location

Feature Maps




The parameters of all neurons in a feature map are collectively called the filter

Why filter?


Because weights **highlight (filter)** specific patterns from the input pixels

Filters



Horizontal Filter


Neuron will detect horizontal
lines in input



Vertical Filter

Neuron will detect vertical lines
in input


Feature Maps



Notice also that neurons are not connected to all pixels

CNNs are *sparse* neural networks


Convolutional Layer



Each convolutional layer consists of several feature maps of equal sizes

The different feature maps have different parameters

Convolutional Layer




Each neuron's receptive field includes the feature maps of all previous layers

This is how **aggregated** features are picked up


The CNN as a whole consists of multiple convolutional (and pooling) layers

More on pooling layers in a bit


CNNs



Feature Map




Convolutional
Layer




CNN


RGB Channels



RGB




Feature Map




Convolutional
Layer

RGB Channels




Feature Map




Convolutional
Layer

RGB


Output of a Convolution Layer Neuron




Input Image



Layer 1




Layer 2




Layer L


Output of a Convolution Layer Neuron




Input Image



Layer 1




Layer 2




Layer L

Map m ,
Column c ,
Row r


Output of a Convolution Layer Neuron



Input Image



Layer 1



Layer 2



Layer L

Map m ,
Column c ,
Row r

Neuron output depends on corresponding* neurons from
each preceding layer

(*corresponding: same receptive field and feature maps,
different layers)

Pooling Layers

Two Kinds of Layers in CNNs

Convolutional

Local receptive field

Pooling

Subsampling of inputs

Convolution

6

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

6

Matrix



x1	x0	x1
x0	x1	x0
x1	x0	x1

4

1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8	2.0	1.8

4

Convolution Result

Two Kinds of Layers in CNNs


Convolutional

Local receptive field

Pooling

Subsampling of inputs

Pooling Layers



Neurons in a pooling layer have no weights or biases

A pooling neuron simply applies some aggregation function to all inputs

Max, sum, average...

Max Pooling

4

4

0.2	0.8	0.3	0.6
0.2	0.9	0.3	0.8
0.3	0.8	0.8	0.9
0	0	0.2	0.8

Matrix



Max,
2x2 filter,
stride = 2

2

2

0.9	0.8
0.8	0.9

Pooling Result

Max Pooling

4

4

0.2	0.8	0.3	0.6
0.2	0.9	0.3	0.8
0.3	0.8	0.8	0.9
0	0	0.2	0.8

Matrix



Max,
2x2 filter,
stride = 2

2

2

0.9	0.8
0.8	0.9

Pooling Result

Max Pooling

4

4

0.2	0.8	0.3	0.6
0.2	0.9	0.3	0.8
0.3	0.8	0.8	0.9
0	0	0.2	0.8

Matrix



Max,
2x2 filter,
stride = 2

2

2

0.9	0.8
0.8	0.9

Pooling Result

Max Pooling

4

4

0.2	0.8	0.3	0.6
0.2	0.9	0.3	0.8
0.3	0.8	0.8	0.9
0	0	0.2	0.8

Matrix



Max,
2x2 filter,
stride = 2

2

2

0.9	0.8
0.8	0.9

Pooling Result

Max Pooling

4

4

0.2	0.8	0.3	0.6
0.2	0.9	0.3	0.8
0.3	0.8	0.8	0.9
0	0	0.2	0.8

Matrix



Max,
2x2 filter,
stride = 2


2

2

0.9	0.8
0.8	0.9

Pooling Result


Pooling Layers



Why use them?

- greatly **reduce memory usage** during training
- **mitigate overfitting** (via subsampling)
- make NN recognise features independent of location (**location invariance**)

Pooling Layers




Pooling layers typically act on each channel independently

So, usually, output area < input area but

Output depth = Input depth


CNNs for Classification

Typical CNN Architecture




Alternating groups of convolutional and pooling layers

Typical CNN Architecture




Each group of convolutional layers usually followed by a ReLU layer

Typical CNN Architecture




The output of each layer is also an image

Typical CNN Architecture




However successive outputs are smaller and smaller (due to pooling layers)

Typical CNN Architecture



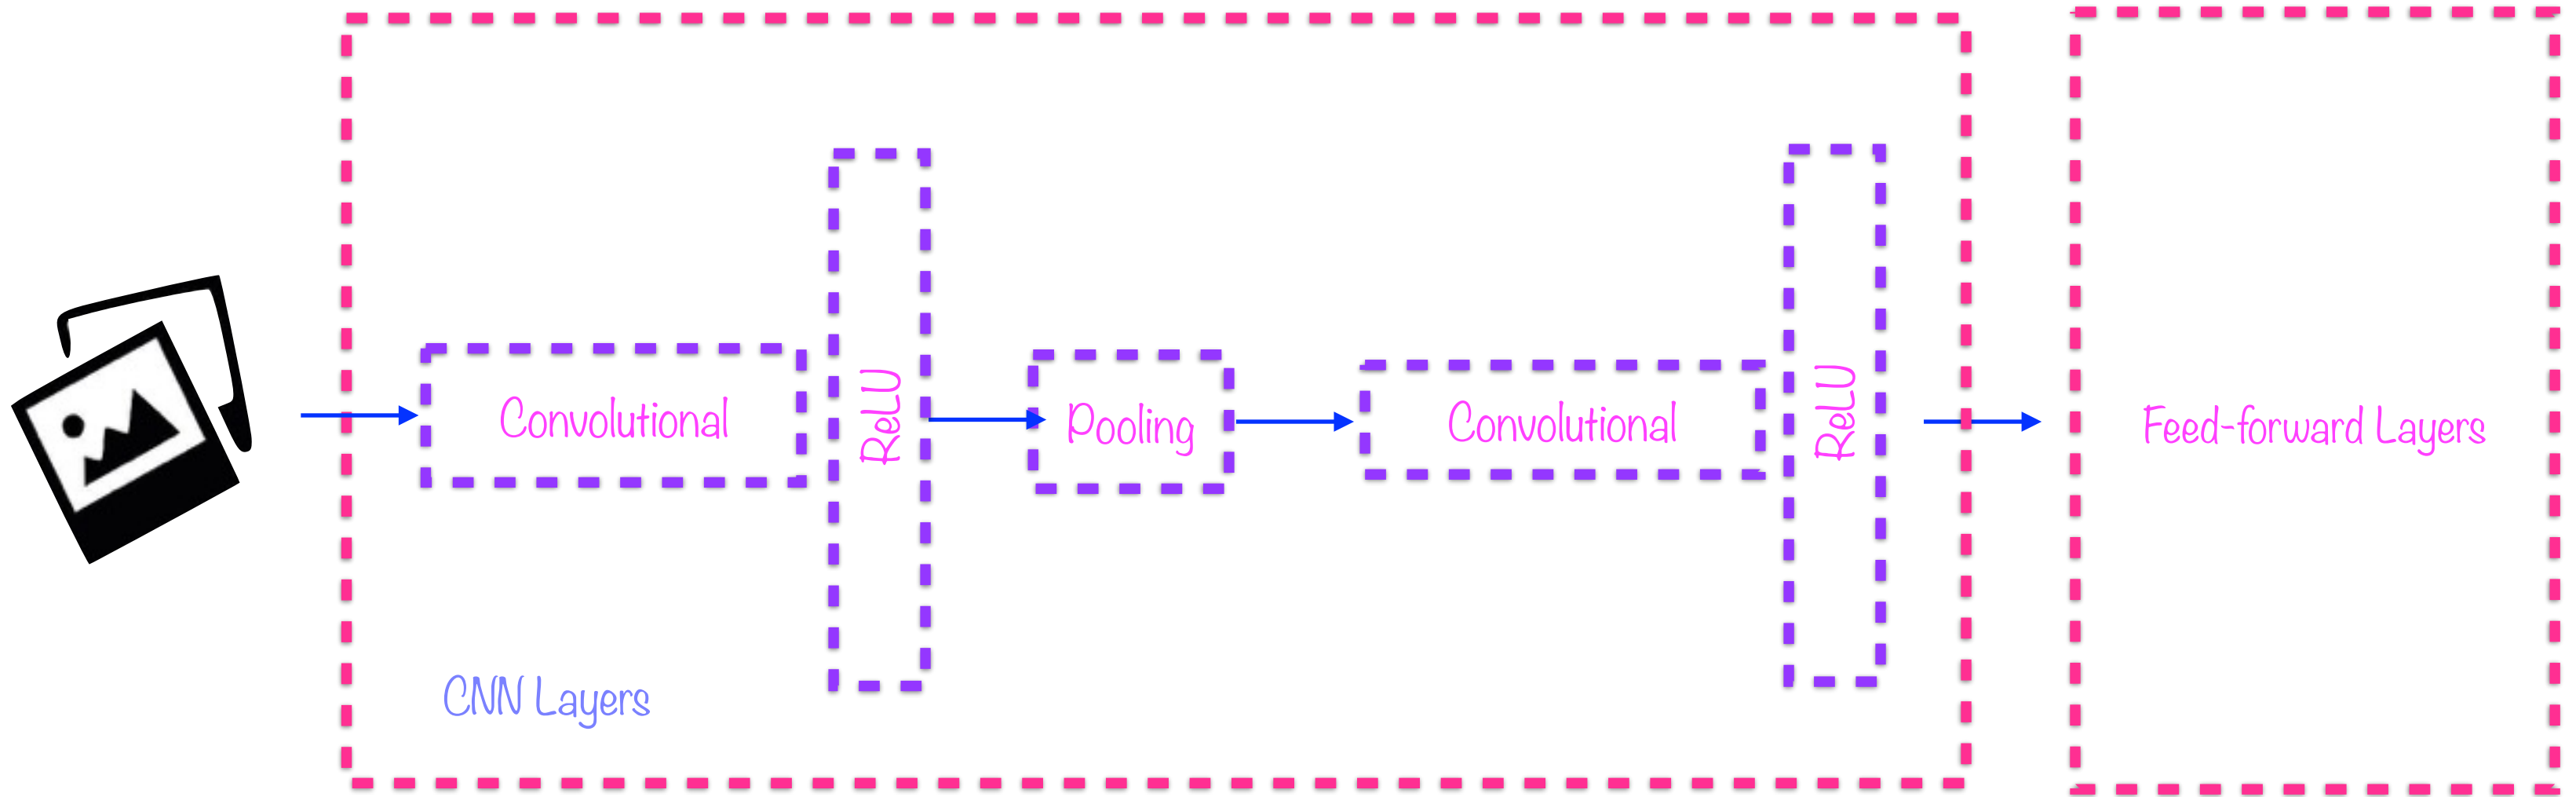
As well as deeper and deeper (due to feature maps in the convolutional layers)

Typical CNN Architecture



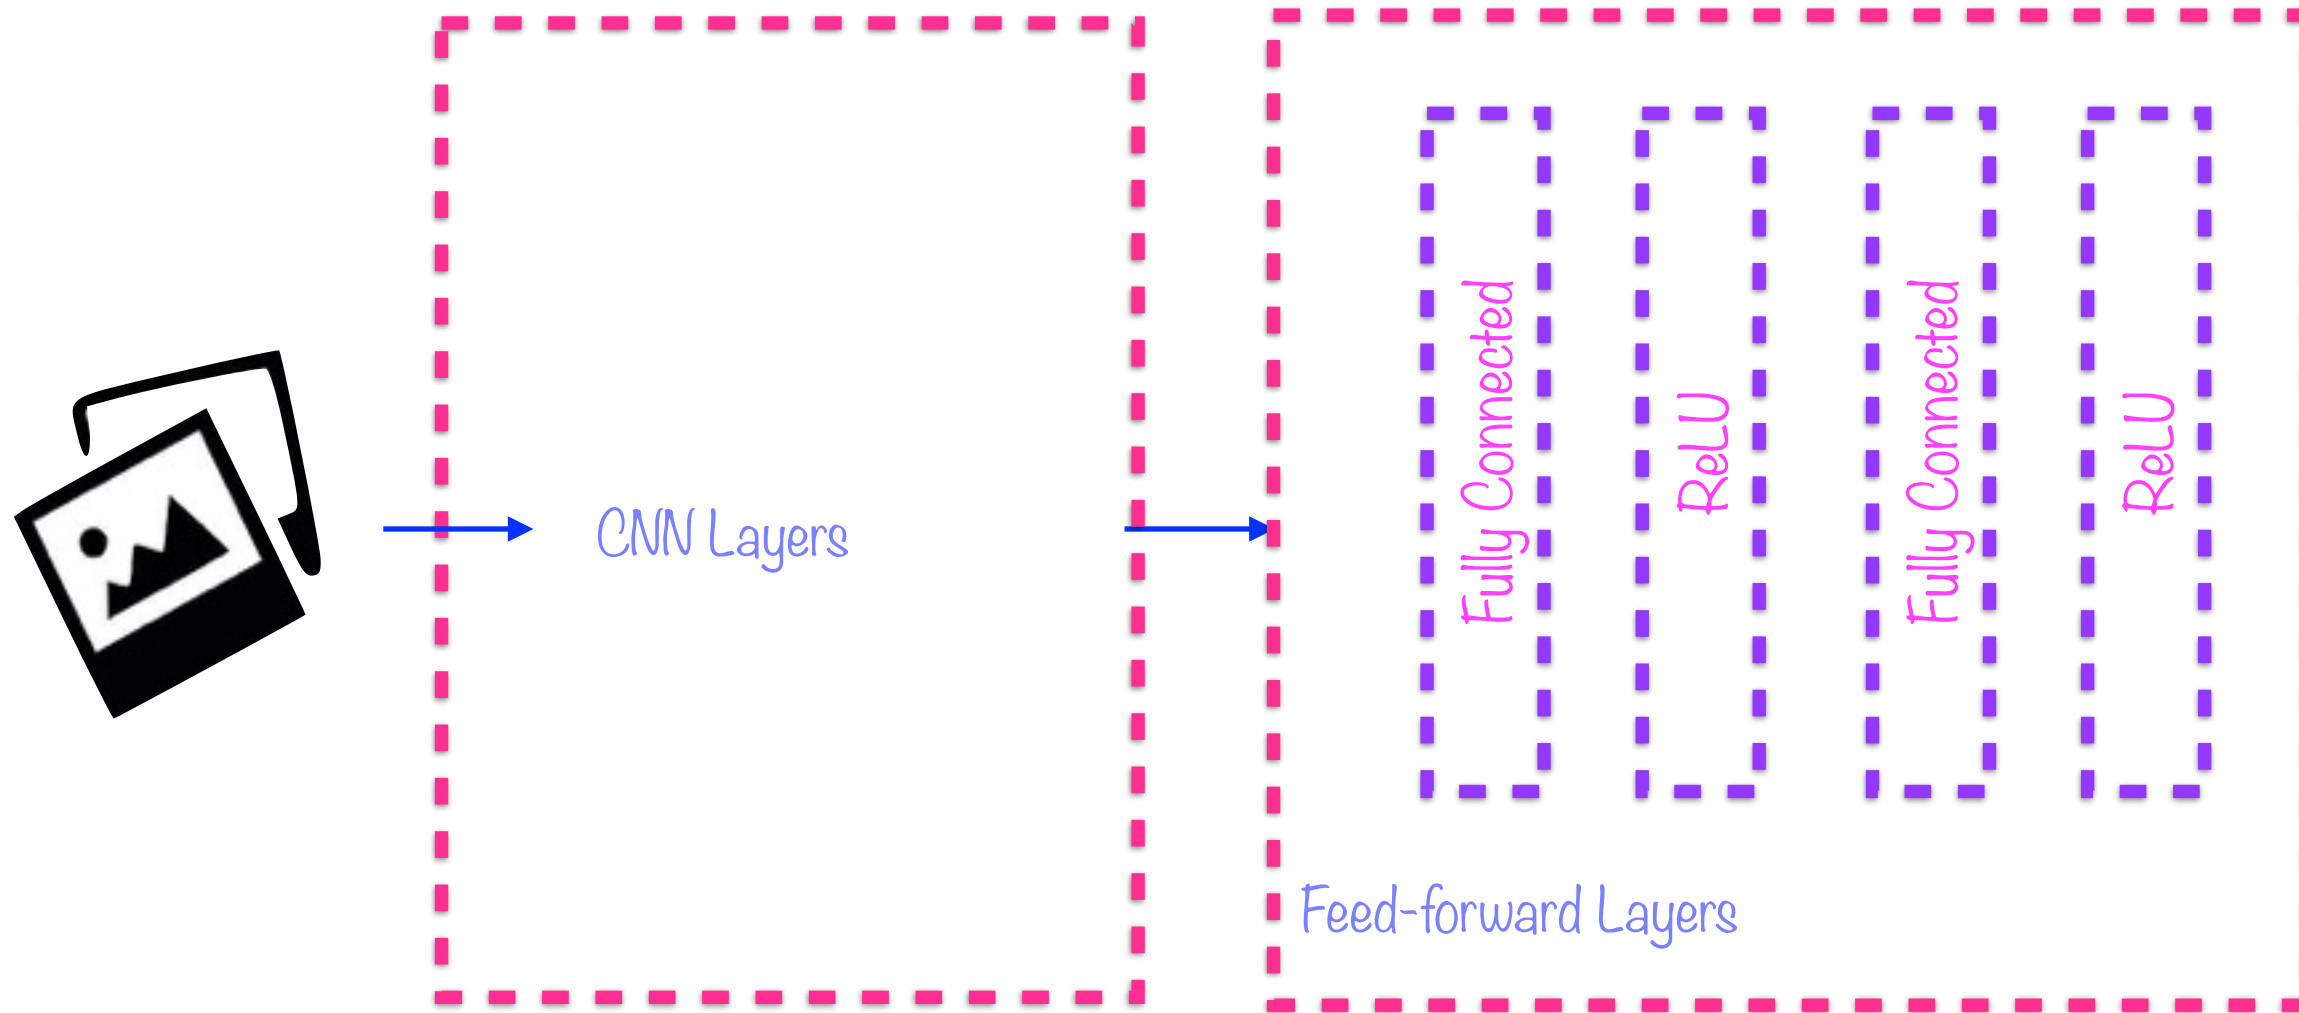
This entire set of layers is then fed into a regular, feed-forward NN

Typical CNN Architecture




This entire set of layers is then fed into a regular, feed-forward NN

Typical CNN Architecture



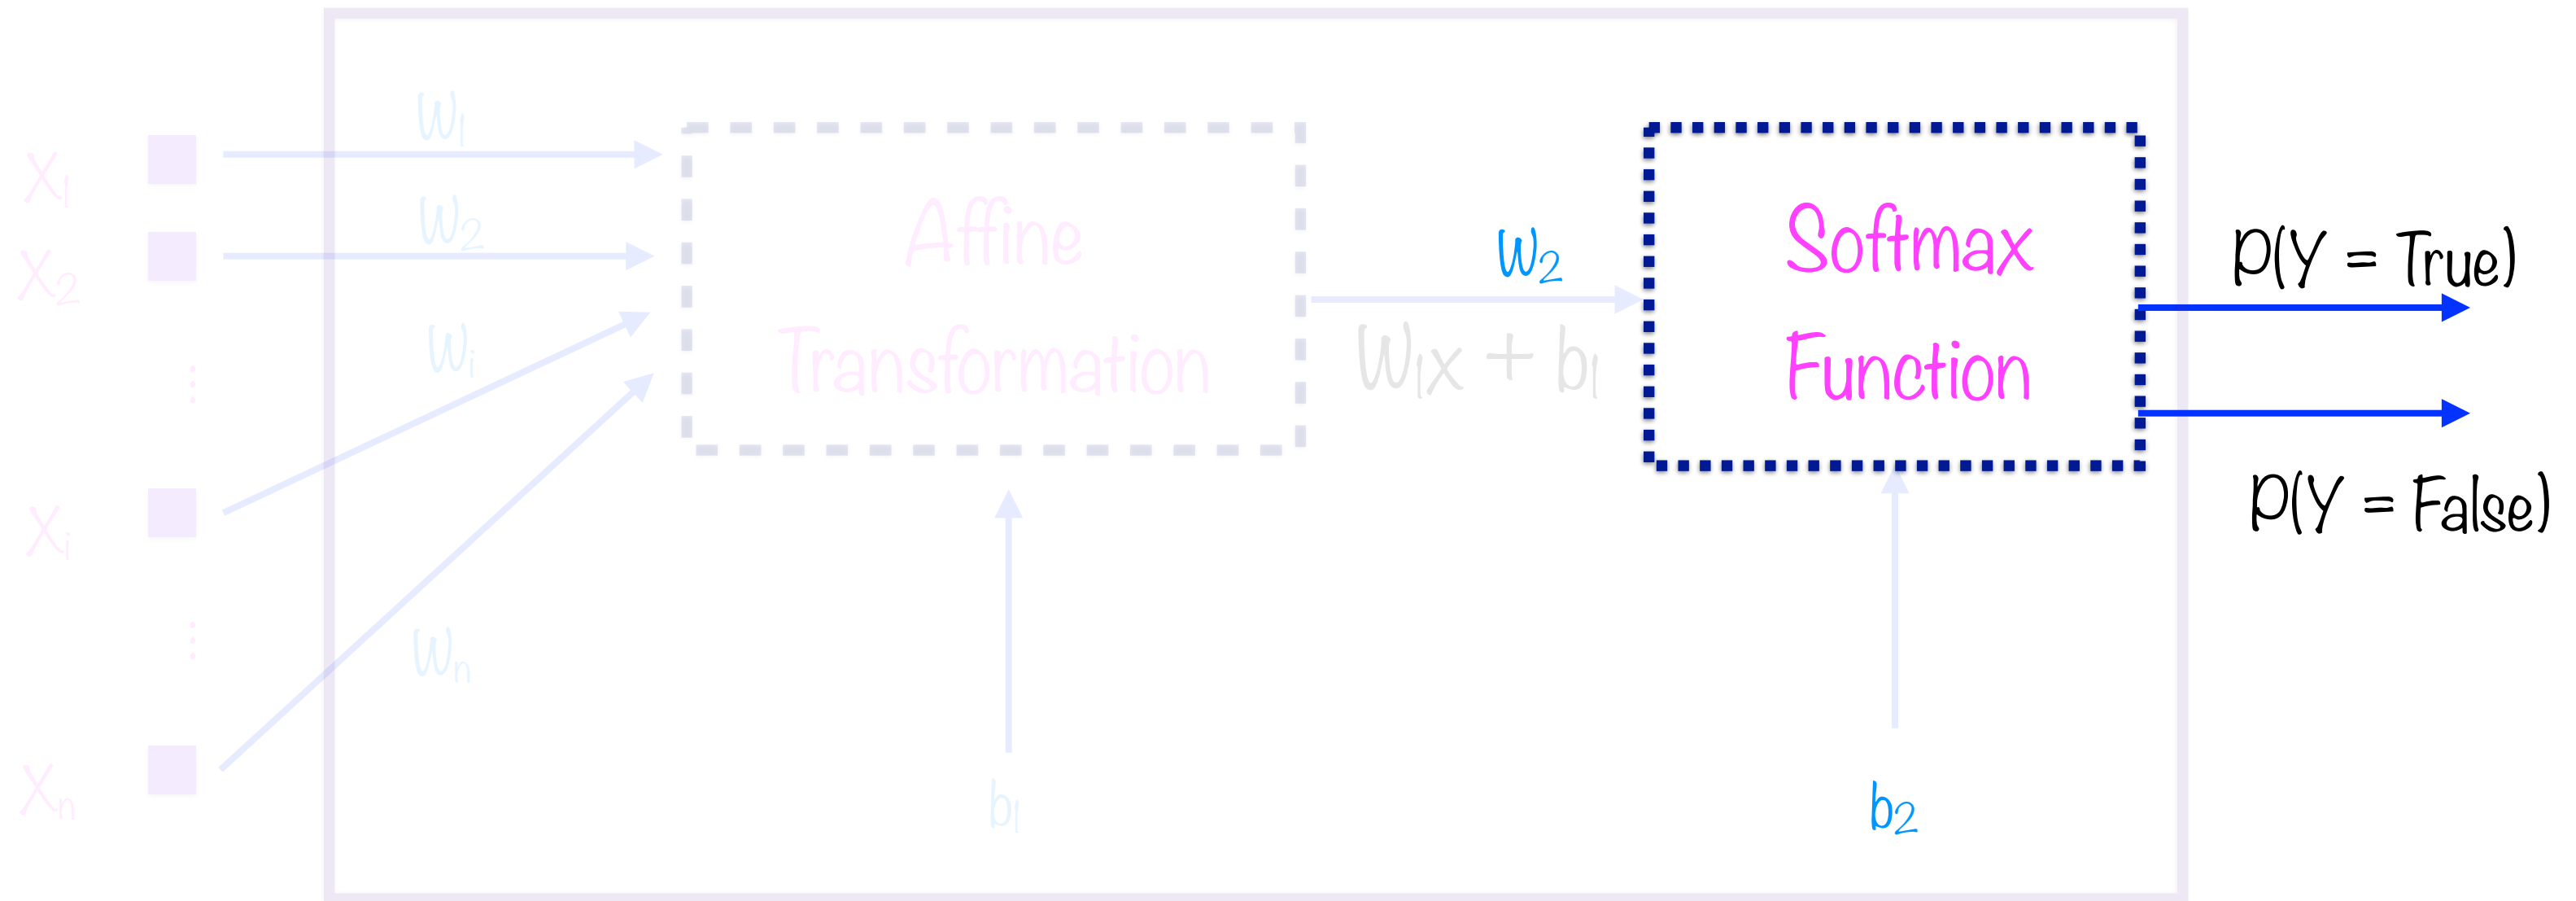
This feed-forward has a few fully connected layers with
ReLU activation

Typical CNN Architecture

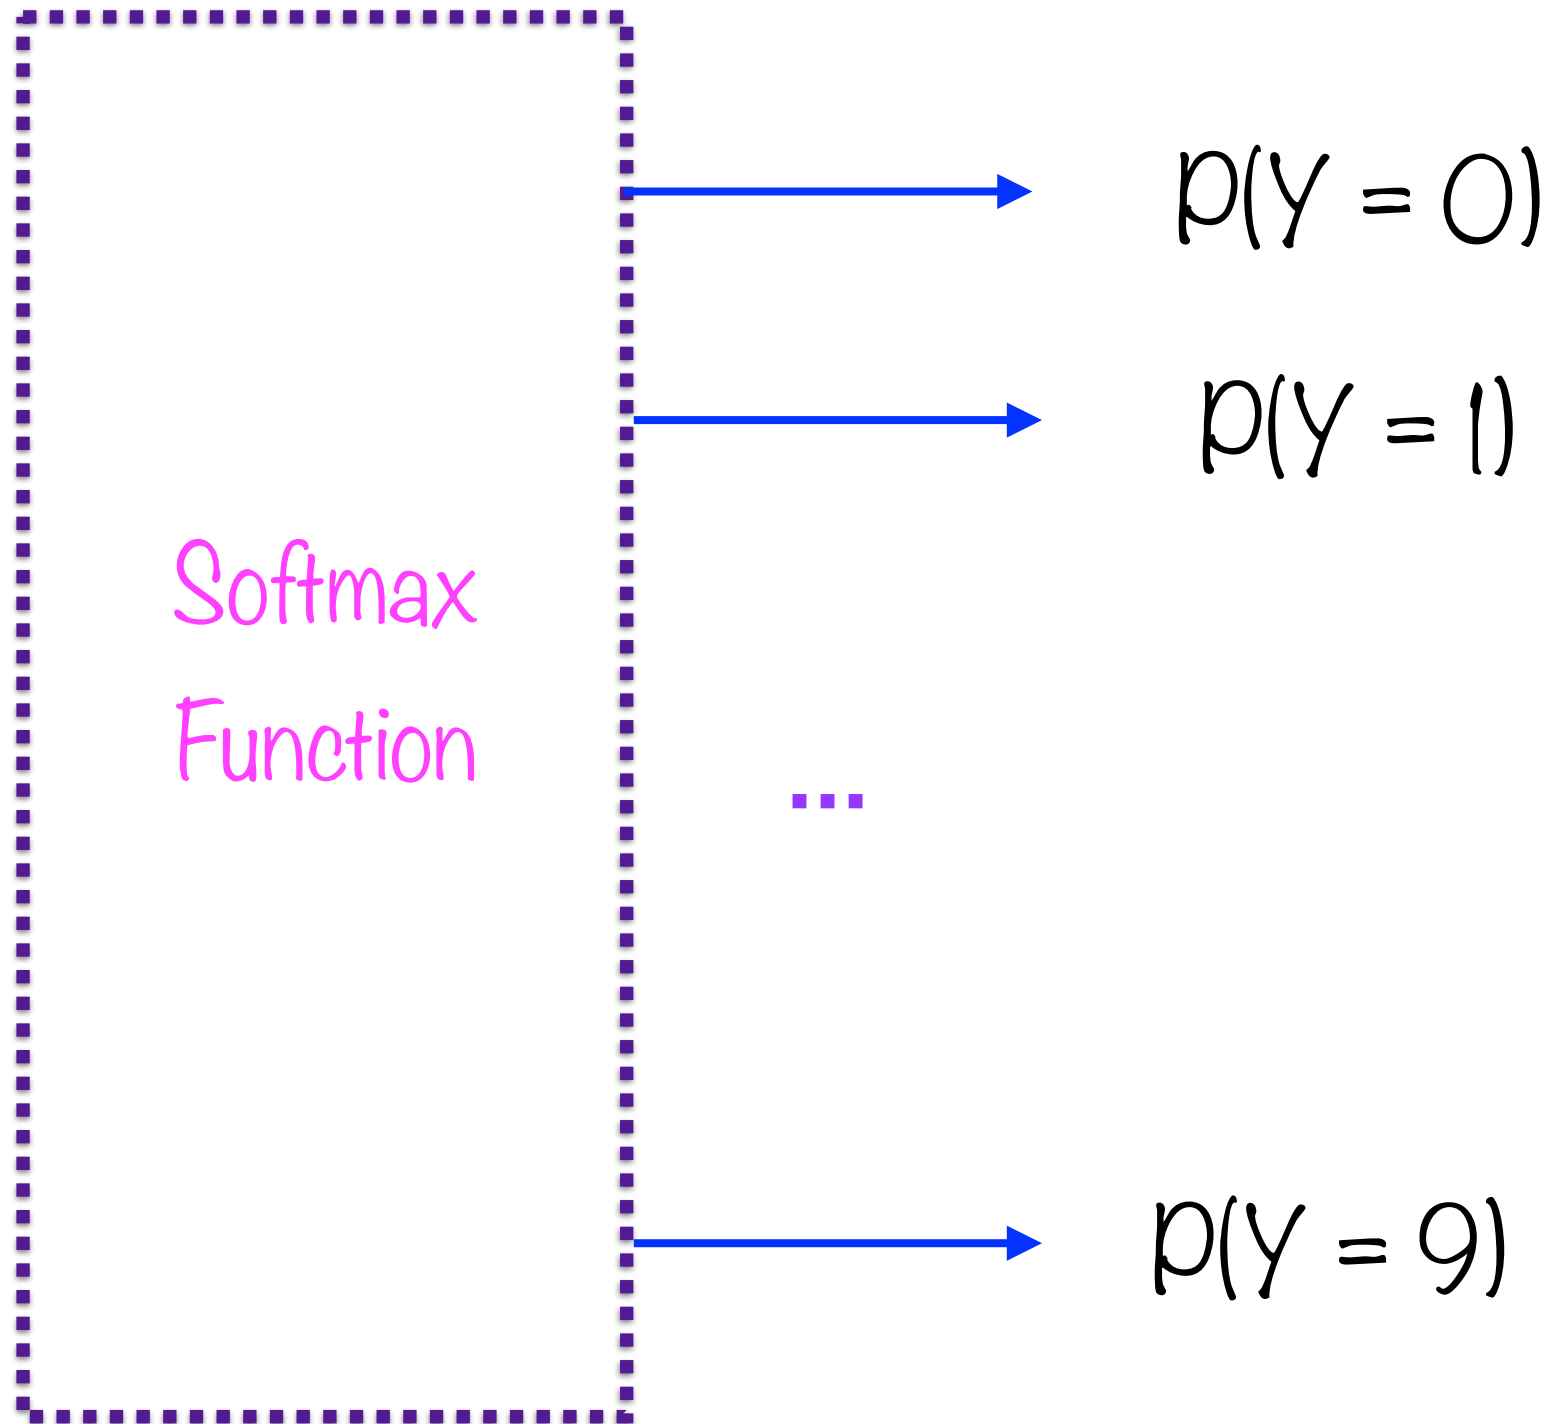


Finally a SoftMax prediction layer

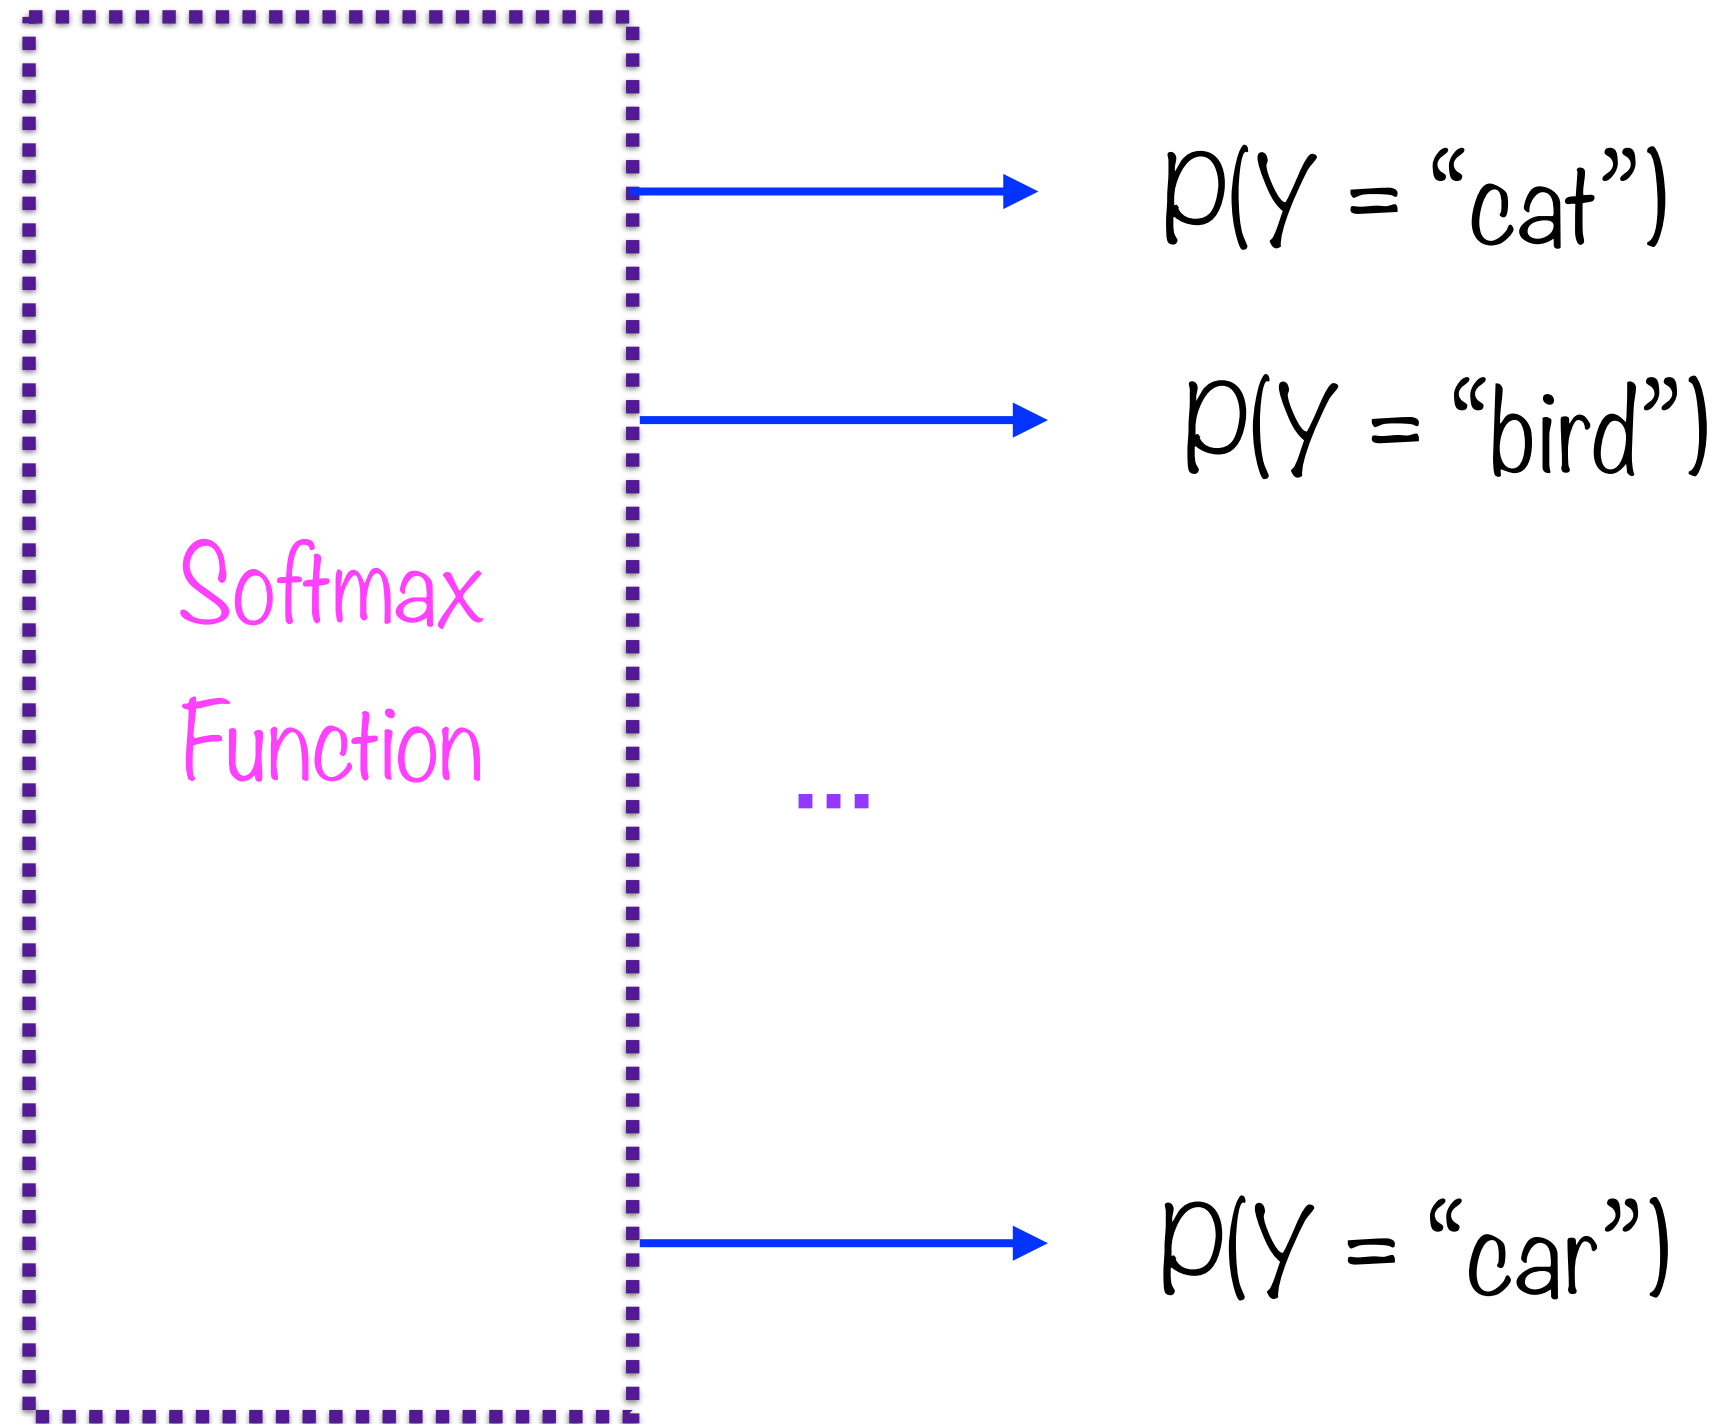
Logistic Regression with One Neuron



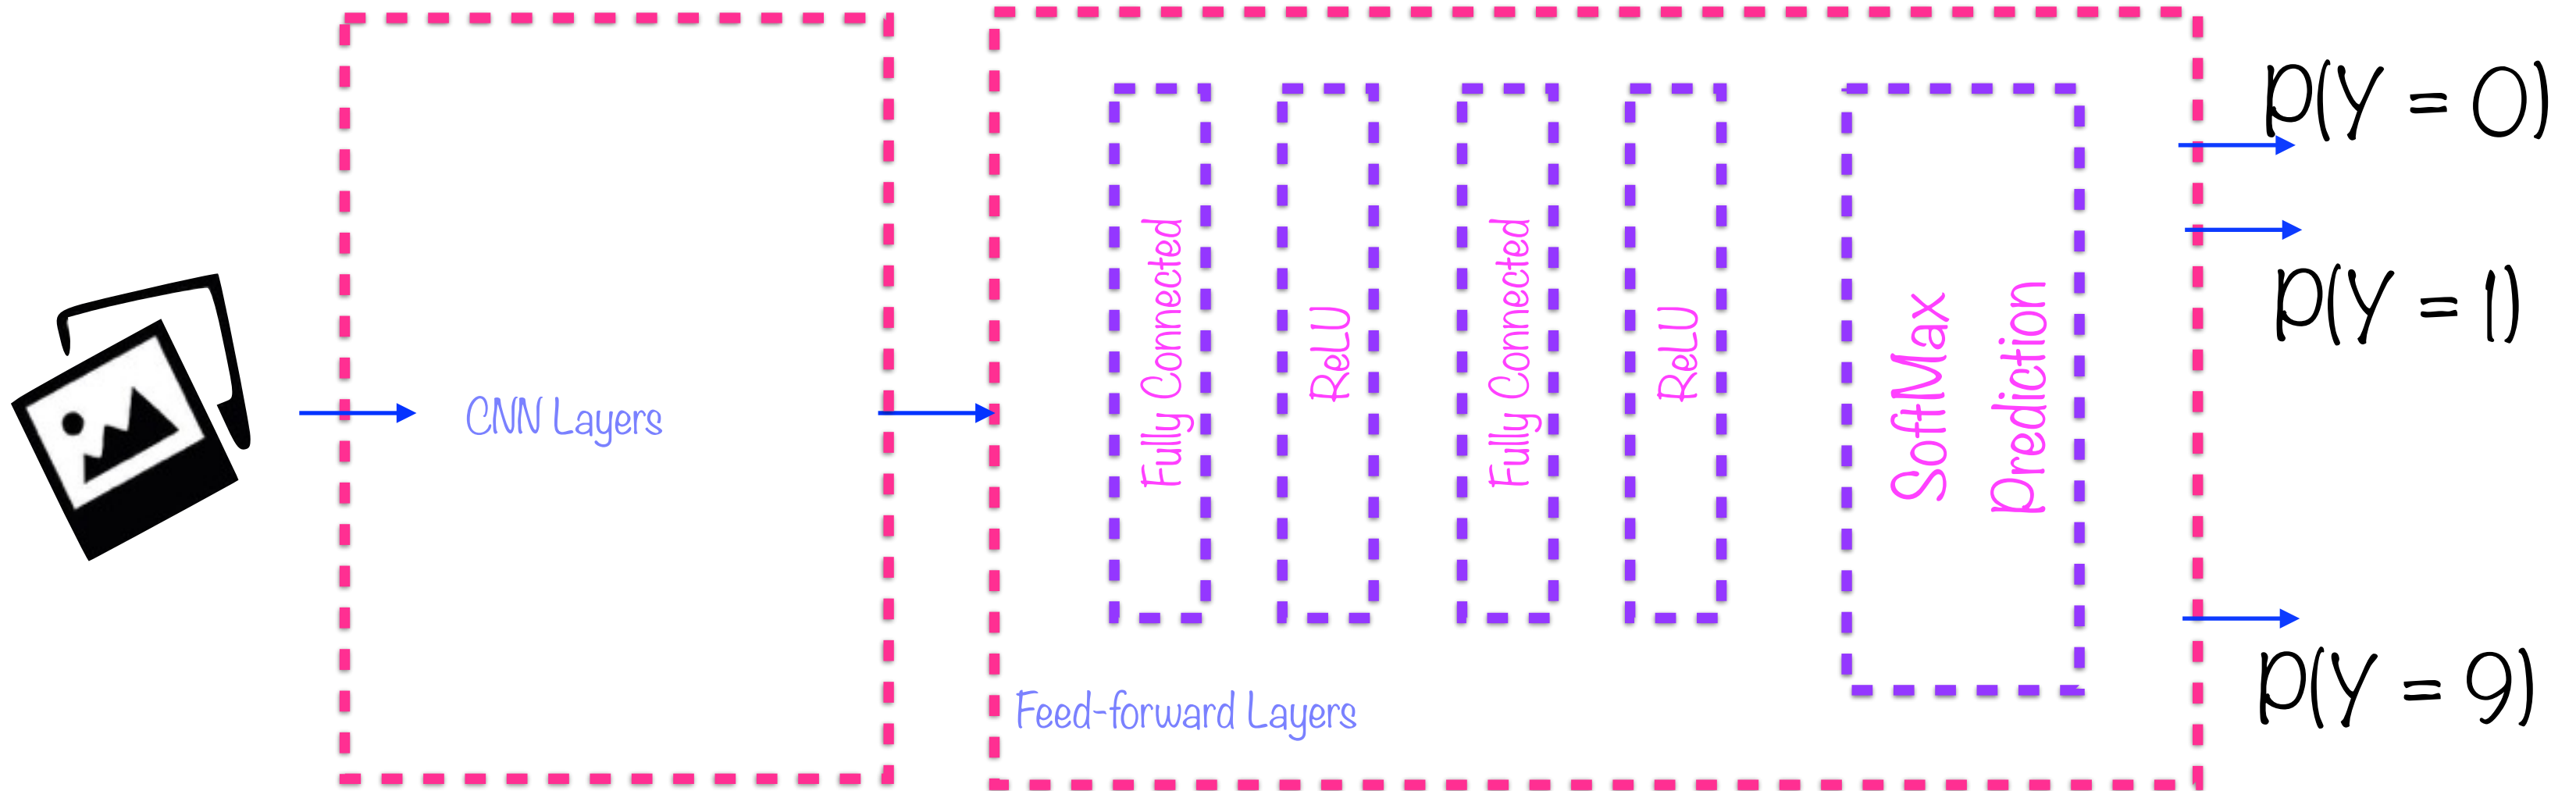
SoftMax for Digit Classification



SoftMax for Image Classification




Typical CNN Architecture



This is the output layer, emitting probabilities

Typical CNN Architectures



Alternating groups of convolutional and pooling layers

Each group of convolutional layers usually followed by a ReLU layer


Image gets smaller and smaller (due to pooling)

Also deeper and deeper (due to convolution)


Typical CNN Architectures

At output end of CNN, regular feedforward NN stacked on

- Few fully connected layers
- Input into these are small images
- ReLU activations
- Finally, a Softmax prediction layer



Typical CNN Architectures



Input is an image

Outputs are probabilities